


agent_pos eval R^2 (frame-level CV, steps 250-500)

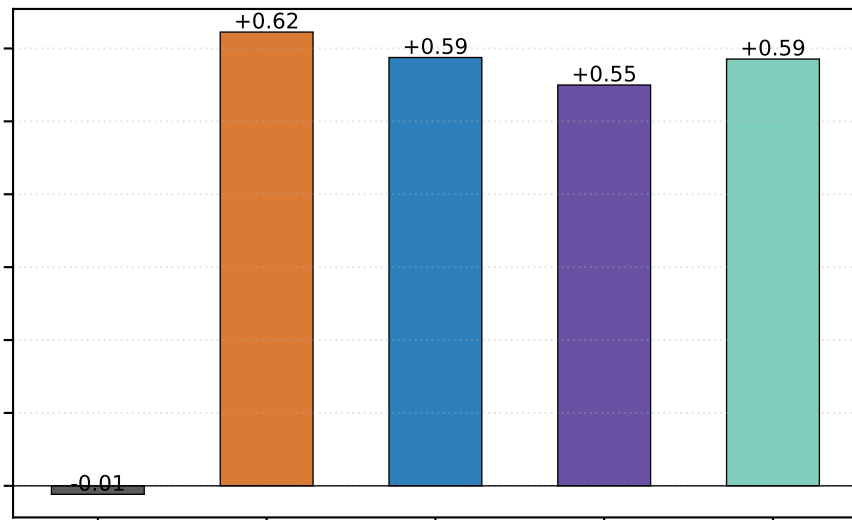
World-model probe: frozen DINOv2-B + small LSTM, Memory Maze 9×9 (architecture-agnostic test)

A — encoder bandwidth → position decode (Memory Maze)

B — non-linear gap



MLP — linear R^2 (format-shift recovery)



blind

coarse 1×1

foveated 4×4 $\sigma=4$

uniform 4×4

fov-logpolar

blind

coarse 1×1

foveated 4×4 $\sigma=4$

uniform 4×4

fov-logpolar