
Sensor structure shapes the format of cognitive maps in navigation agents

Anonymous authors
Paper under double-blind review

Abstract

The format of spatial memory in the hippocampus varies systematically with a species’ visual ecology, but whether this reflects a general sensor–memory coupling is hard to test directly: animal sensor ecology is fixed within species, and AI systems rarely co-train encoder and memory across a graded sensor design. We test this in silicon: a navigation agent where only the visual sensor varies along a graded scale of encoder spatial bandwidth, with architecture, task, and training held fixed. Across five conditions, agents reach near-identical task performance yet the hidden-state code for position differs in form. With a vision-limited encoder, position concentrates into a scene-invariant direction linearly readable from the hidden state; with a rich encoder, it distributes into non-linear directions and partly offloads to the encoder. Linear readability shifts faster than total position information — a format shift, not a total-information collapse. We frame this as an emergent *capacity allocation* between the encoder route and the memory-integration route, measured along three concomitant aspects (magnitude, format, consumption). Memory transplants are asymmetric (rich-encoder policies fail on bottleneck states; the reverse works), paralleling cross-modality remapping in hippocampal coding under hidden-state inference. Probe-readable position further dissociates from policy reliance: a bottleneck-encoder agent under-weights its readable code while a rich-encoder agent reads a non-spatial substitute. This positions sensor structure as a training-time lever on cognitive-map format, generates a falsifiable architecture-agnostic prediction for transformer navigators, and reframes bandwidth-restricting perception as potentially cognition-enabling rather than only compute-saving.

1 Introduction

Visual perception and downstream memory in artificial systems are commonly treated as a sequential pipeline, guided by the implicit assumption that scaling the encoder uniformly improves downstream representations. Recent benchmarks complicate this picture: vision–language models with strong visual recognition still fail at fine-grained spatial reasoning [Ramakrishnan et al., 2025, Yang et al., 2025]. Existing explanations have focused on the training distribution [Chen et al., 2024] or on differences in the geometric content of pretrained encoders [Tong et al., 2024, El Banani et al., 2024]. We instead question the underlying assumption: by treating perception and memory as a co-trained system, we ask how the structure of the encoder constrains what spatial information is accessible from memory through simple linear readout.

Cognitive neuroscience has long recognised that sensor structure shapes the format of cognitive maps. Across species, distinct sensory ecologies drive distinct spatial-coding strategies — foveated primates rely on view cells while near-panoramic rodents rely on place cells [Wirth et al., 2017, Rolls, 2024]; echolocating bats and congenitally blind humans show striking cross-modal remapping [Geva-Sagiv et al., 2016, Fortin et al., 2008]. Empirically isolating the causal role of sensor ecology is notoriously difficult in biology: body plans and evolutionary histories cannot be held




Figure 1: **Experimental setup and analysis pipeline.** *Row 1 (architecture):* four sensor conditions span a graded scale of encoder spatial bandwidth (*blind, coarse, foveated, uniform*; leftmost column), feeding a shared ResNet-18 encoder + 3-layer LSTM; *blind* bypasses the encoder. A log-polar foveation control (Appendix I) adds a 5th intermediate point on the same scale. *Row 2 (analysis):* three method blocks — decodability, intrinsic structure, causal intervention — read the frozen top-layer hidden state $h_2 \in \mathbb{R}^{512}$ to test three views of one capacity allocation: *magnitude, format, consumption* (§2; synthesis in §4).

constant, so cross-species effects are confounded with everything else that differs between species. Artificial systems offer a powerful alternative, yet modern deep reinforcement learning rarely co-trains perception and memory across a systematically controlled sensor design.

We bridge this gap *in a silicon setting*. Building on deep-RL navigation agents known to develop cognitive-map-like representations even in the absence of vision [Wijmans et al., 2023, Banino et al., 2018, Cueva and Wei, 2018],¹ we introduce a comparative testbed that holds the encoder backbone, recurrent memory, task, optimisation, and reward entirely fixed, and varies a single degree of freedom: the spatial bandwidth of the visual sensor across five conditions (Figure 1).

We hypothesise an emergent *capacity-allocation* principle: the visual encoder and the recurrent integration of continuous sensor history (e.g., GPS) compete for a finite hidden-state budget. Our results confirm this tradeoff. As encoder spatial bandwidth increases, the linearly-readable, world-frame position signal in the recurrent state collapses monotonically while task performance stays near-perfect (93–99% success). We establish this reallocation along three concomitant views: *magnitude* (how much linearly-accessible position information lives in the recurrent state), *format* (which subspace carries it across conditions), and *consumption* (which route the policy actually uses for behaviour). Robustness against single-method artefacts is confirmed by convergent results across linear, non-linear (MLP), and information-theoretic (MINE) probes, leave-one-scene-out validation, and Procrustes shape distance. We further ground the format shift behaviourally with asymmetric memory transplants: rich-encoder policies fail to navigate using bottleneck states, demonstrating the representational change is mechanistically intrinsic and not an artefact of the probing framework [Orozco et al., 2025].

¹“Cognitive map” in our paper is operationalised as the linearly-decodable world-frame position code at the policy-readable hidden state. Aspects of the broader concept [Epstein et al., 2017] are engaged by our LOSO scene-invariance test (relational across rooms) and predictive-horizon analysis (SR-style future-state code, Stachenfeld et al. 2017); hierarchical scene clustering and topological reachability are deferred to follow-up.

Our contributions are three folds. **(i) The capacity-allocation hypothesis:** a unifying framework explaining how visual bandwidth dictates spatial-memory format, organising five concomitant signatures (magnitude collapse, format divergence, consumption dissociation, scene-invariance, place-vs-view-cell character) into three concomitant views (magnitude, format, consumption). This matched-condition design enables an analytical avenue biology cannot easily realise: the systems-neuroscience toolkit applied to a *controlled comparative population* rather than a single recording. We treat this comparative methodology as an explicit secondary contribution, developed in full in §5. **(ii) Causal behavioural validation:** asymmetric memory transplants demonstrating that rich-encoder policies fail on bottleneck states, mechanically grounding the format shift and mirroring biological global remapping under cross-modality reorganisation. **(iii) Broader theoretical implications:** a falsifiable prediction for transformer-based navigators (§4) and an architecture-side reframing of bandwidth-restricting perception (foveation, log-polar) as *cognition-enabling* rather than only compute-saving — complementing data-side accounts [Chen et al., 2024] of VLM spatial-reasoning failures.

2 Methods

Architecture, task, and conditions (Figure 1, Row 1). PointGoal navigation in Habitat [Savva et al., 2019] trained with DD-PPO [Wijmans et al., 2020] across five conditions varying only the visual sensor: *blind* (no encoder); *coarse* (48×48 RGB, 1×1 encoder feature map after ResNet-18 downsampling); *foveated* (256×256 with eccentricity-dependent Gaussian blur $\sigma_{\max}=8$, 4×4 feature map); *uniform* (256×256 unblurred, 4×4); *foveated-logpolar* ($\sim 2 \times 2$ via log-polar resampling onto a 64×64 retinal grid, control point isolating encoder-spatial-output from blur, Appendix I). All five share an identical 3-layer LSTM (512-d) and the Wijmans non-visual sensor stack (goal-in-start-frame, GPS, compass, close-to-goal); Gibson-0+ [Xia et al., 2018] and MP3D training pool; 18 MP3D test scenes for distribution-shift probes; 250M frames each; single seed per condition (cogsci modelling convention [Wijmans et al., 2023, Banino et al., 2018, Cueva and Wei, 2018, Schaeffer et al., 2023]; rigour from convergent multi-method evidence, §4.4); blind has an independently re-trained checkpoint (Appendix A). After training, weights are frozen and per-step top-layer hidden state $\mathbf{h}_2 \in \mathbb{R}^{512}$ is collected from $n = 500$ Gibson held-out deterministic-rollout episodes per condition.

Probing methodology (Figure 1, Row 2). Three method blocks read $\mathbf{h}_2$ to test three concomitant aspects of one capacity allocation:

- *Decodability $\rightarrow$ magnitude:* linear Ridge with Hewitt–Liang selectivity [Hewitt and Liang, 2019], 2-layer MLP, MINE [Belghazi et al., 2018], Skaggs spatial information [Skaggs et al., 1996]; plus step-binned, lag- k , predictive-horizon [Stachenfeld et al., 2017, Whittington et al., 2020], and excursion-forgetting variants for memory dynamics.
- *Intrinsic structure $\rightarrow$ format:* LOSO cross-validation [Sanders et al., 2020], cross-condition probe transfer, CKA [Kornblith et al., 2019], 1-NN purity, Procrustes shape distance [Williams et al., 2021], principal angles and position-direction cosines, participation ratio and TwoNN intrinsic dimension [Ansuini et al., 2019, Facco et al., 2017], eigenspectrum power-law [Stringer et al., 2019].
- *Causal intervention $\rightarrow$ consumption:* 5×5 memory transplant and memory-init transplant; paired same-scene shortcut-discovery (Tolman 1948 in silico); GPS-sensor mid-rollout ablation.

Each method is introduced at point of use in §3; full protocols and hyperparameters in Appendix B. Convergence across the three views on the same condition-level contrast is the strongest claim the methodology supports.

3 Results

The five agents reach near-identical task performance (Table 1) despite substantially different visual inputs — the encoder–memory–policy stacks must therefore internally compensate for the input differences, holding different internal representations of the same task. We read those representations along four views matching the three method blocks of Figure 1: how much linearly-readable posi-

Table 1: Per-condition headline. All five agents trained for 250M frames (§2); blind uses an independently retrained checkpoint (Appendix A). Probe R^2 : 5-fold episode-level CV mean $\pm$ std on full-length trajectories (no per-episode step cap). Sub-chance R^2 in rich-encoder rows reflect long-tail collapse (mid-episode $R^2 \in [0.6, 0.95]$, Figure 10b); rich-encoder GPS R^2 is statistically indistinguishable from zero (one-sample t -test: foveated $p \approx 0.58$, uniform $p \approx 0.25$; $df=4$). Bold marks blind/coarse rows where GPS and compass remain decodable ($p < 0.001$) and best task performance.

Condition	SPL	Success	GPS R^2 (no-cap)	Compass R^2 (no-cap)
Blind	0.59	0.93	$+0.934 \pm 0.04$	$+0.844 \pm 0.03$
Coarse	0.853	0.989	$+0.582 \pm 0.099$	$+0.659 \pm 0.061$
Foveated	0.894	0.993	$+0.178 \pm 0.659$	$+0.503 \pm 0.142$
Foveated-logpolar	0.879	0.990	-0.029 ± 0.759	$+0.467 \pm 0.330$
Uniform	0.891	0.993	-1.785 ± 2.993	-1.244 ± 3.579

tion lives in memory (magnitude, §3.1); what shape it takes spatially (§3.2) and temporally (§3.3); and what behaviour reads from memory under causal intervention (consumption, §3.4). Each view points to the same underlying capacity-allocation tradeoff between an encoder-derived visual route and a memory-integration route. The Discussion (§4) applies the same set of analyses to a different environment and encoder family, as a step toward a comparative cognitive neuroscience of artificial navigation agents.

3.1 Magnitude: linear position code shrinks with encoder bandwidth

Picture two routes by which the agent could know “where am I”: integrating motion over time (works even without vision) or reading position-correlated visual features (only works with rich vision). Both can supply position to memory; the policy reads memory through a linear operation. As we widen the visual encoder, the agent reallocates from the first route to the second, and **the linearly-readable position signal at the policy-readable layer shrinks monotonically with encoder bandwidth** (Figure 2a). All five agents reach near-ceiling task performance (Table 1), so this is a representational dissociation, not a skill gap. The position information itself is not lost — a non-linear probe still recovers it (§3.2) — only the linear, policy-accessible component shrinks.

The bandwidth–magnitude relationship is robust against two natural alternative readings, and a sharper reading of the trigger emerges from the controls. The chance-level rich-encoder GPS R^2 is a real property of the hidden state rather than a limit of the probe: the same linear decoder reads distance-to-goal (DtG) at $R^2 \geq 0.81$ in every condition (Appendix C; Figure 18), so position-correlated structure is recoverable wherever it exists. The discriminator is not the number of encoder output cells either; FOVEATED-LOGPOLAR, whose encoder produces a $\sim 2 \times 2$ feature map close to coarse’s 1×1 , nonetheless lands in the rich-encoder regime — placing the relevant variable on per-step *feature variety*, the diversity of visual features delivered at each timestep, rather than cell count (Appendix I). Read this way, blind reproduces Wijmans et al. [2023]’s emergent-cognitive-map result for goal-driven agents without vision, and coarse extends it to a *vision-with-bottleneck* regime — a visual encoder is present, but its 1×1 output collapses per-step variety to one channel, leaving the same integration route to dominate the linear position code.

Capacity allocation between two routes. Memory has two routes to a position code — integrating the position sensor’s signal over time, or reading position-correlated features from the encoder — and policy gradient biases capacity toward whichever route the encoder makes viable. Two views in the main figure are consistent. *Across training* (Figure 2b): sighted conditions start with high linear readability; rich-encoder conditions then decay along trajectories whose timescale tracks encoder informativeness, while coarse plateaus and blind stays stable. *Predictive horizon* (Figure 2c): the successor-representation cognitive-map signature [Stachenfeld et al., 2017] — memory predicting where the agent will be many steps ahead — is sustained only in blind, where the integration route carries the position code.

Corroborating diagnostics. Three appendix tests converge on the same picture (Appendix Figure 10). (i) *Pipeline view*: position readability is comparable across conditions where the position sensor first enters memory, and only diverges at the policy-readable layer — isolating the divergence to where the policy reads. (ii) *Peaked-unit content*: the most spatially-informative units in rich-encoder agents




Figure 2: **Three views of magnitude.** (a) *Cross-condition*: top-layer linear GPS R^2 falls with encoder bandwidth. (b) *Across training*: rich-encoder conditions start high then lose readability; bottleneck conditions hold it throughout. (c) *Predictive horizon*: only bottleneck conditions linearly decode pos_{t+k} at long k . Shaded band: $R^2 < 0$, the linear probe is worse than predict-mean baseline — the “no-signal zone”. Uniform’s curve dives deep into this zone (its hidden state is scene-conditional, so 5-fold CV trains on some scenes’ position-directions and tests on incompatible ones, producing systematically negative R^2); the dive is noise, not a meaningful trend.




Figure 3: **Three views of spatial format.** (a) *Format-shift severity (probe-depth sweep)*: GPS R^2 as a function of probe depth (linear $\rightarrow$ MLP-1 $\rightarrow$ MLP-2 $\rightarrow$ MLP-4) per condition. Bottleneck conditions plateau near $R^2 \approx 0.95$ from the very first column — linear already suffices, so deeper probes recover almost nothing more. Rich-encoder conditions ramp up sharply with probe depth — position is encoded but only non-linearly, and the depth needed to recover it grows with encoder bandwidth. (b) *Format dichotomy*: bottleneck sits in the scene-invariant + transplant-robust quadrant; rich-encoder conditions cluster in the scene-conditional + transplant-brittle quadrant. (c) *Subspace divergence*: all 10 condition pairs cluster near “fully orthogonal” (90°, $\cos = 0$), corroborating (b) at the geometric level.

track trajectory-phase and visual-landmark features rather than raw position, while in blind they track raw position — a view-cell-like vs. path-integrated dissociation [Skaggs et al., 1996, Rolls, 2024]. (iii) *Perturbation robustness*: a 25-step random-action mid-episode excursion fragments blind’s position decoding most — the inverse of a passthrough memory and consistent at the high level with rodent path-integration under removed visual landmarks [Chen et al., 2016].

3.2 Spatial Format: collapsed linear readability is a format shift, not an information loss

Picture two ways the agent’s memory might organise a position code: along a single direction shared across rooms (one “where am I” axis usable in any scene), or along scene-specific directions selected per context. The policy reads memory through a linear operation, so it can decode the first format but not the second. All five conditions *preserve* the position information itself — a non-linear probe recovers it where a linear probe is at chance, and information-theoretic estimates of total information between memory and position shift far less than the linear R^2 (Figure 3a; Appendix D.1) — but **the magnitude collapse of §3.1 is a format shift, not an information loss**: encoder bandwidth shifts which of the two formats the memory uses. The remainder of §3.2 characterises this shift along two complementary views: *within* a condition (does the readout direction depend on the scene?) and *across* conditions (do different conditions even share an axis?).

Within-condition: scene-invariance dichotomy. Leave-one-scene-out cross-validation gives a clean answer (Figure 3b, x-axis): bottleneck conditions generalise across held-out scenes (a single direction works in any room), while rich-encoder conditions do not (the position-direction differs per scene, so a single linear hyperplane cannot capture it; a non-linear probe can, by switching its readout on visual context — exactly the format-shift gap in Panel a). Intuitively, a bottlenecked encoder has no scene-specific features to bind to, so memory is forced toward a universal direction; a rich encoder lets memory exploit scene-specific features, fragmenting the position-direction across rooms. The dichotomy mirrors the no-remapping vs. global-remapping spectrum of hippocampal coding under hidden-state inference [Sanders et al., 2020, Achille and Soatto, 2018].

Across-condition: non-interchangeable subspaces. The picture sharpens further across conditions: each condition’s position-direction lives in a mostly-orthogonal subspace of every other’s (Figure 3c, all 10 pairs cluster near 90° and direction-cosine ≈ 0) — the codes are not just *different*, they are nearly maximally distinct. A behavioural test confirms this in policy-execution terms (Figure 3b, y-axis): a 5×5 memory transplant shows a direction-asymmetry — a bottleneck-style memory state breaks rich-encoder policies, while a rich-encoder memory state still drives bottleneck policies, paralleling cross-modality remapping in echolocating bats [Geva-Sagiv et al., 2016]. Manifold-shape scalars (intrinsic dimension and eigenspectrum power-law α , Appendix Table 2) corroborate at the geometric level: blind separates on α (sharper variance fall-off, concentrated in a few directions) and on intrinsic dimension (lowest), with the four sighted conditions clustering on both metrics — the same bottleneck-vs-sighted contrast surfaced earlier in this section.

Synthesis. Spatial format is capacity allocation in geometric terms: bottleneck encoders push memory toward a scene-invariant single-direction code; rich encoders distribute capacity into scene-conditional directions occupying orthogonal subspaces. The first format is linearly readable by the policy; the second is not — this is the policy-side mechanism behind the magnitude collapse of §3.1. The structural separation is laid down well before convergence (Appendix Figure 15c).

3.3 Temporal Format: blind recomputes the code each step, sighted read it as stationary

Picture two ways memory might track a slowly-changing target like “where is the goal”: read the same direction throughout the trajectory (a stationary code; the value updates but the axis stays put), or rotate the readout direction step-by-step as the agent integrates new motion (a transient code; every coordinate’s meaning shifts). Bottleneck conditions take the second path; sighted-rich conditions take the first. **The static format shift of §3.2 has a temporal counterpart: the readout axis for goal direction rotates step-by-step in blind, but stays put in sighted conditions.**

Step-by-step decoder generalisation. Train a linear decoder of egocentric goal-vector at trajectory step t_{train} , evaluate it at every other step t_{test} — for each $(t_{\text{train}}, t_{\text{test}})$ pair, ask whether the decoder still works [King and Dehaene, 2014]. The five heatmaps in Figure 4 (top row) reveal three regimes: a *diagonal-only* pattern in blind (the decoder at step t does not transfer to step $t \pm 30$), a *square-block* pattern in coarse (the same decoder works across the trajectory), and graded intermediates in rich-encoder conditions. Mechanistically: blind must integrate motion from $t=0$ onward to compute its goal-vector, so the meaning of every memory coordinate shifts step-by-step; coarse derives goal-direction from a slowly-varying colour cue, so the memory→goal mapping is quasi-stationary. Our pre-registered prediction was the opposite direction (blind as the stationary block); we report what we found.

The within-sighted spectrum tracks encoder bandwidth. The bottom row of Figure 4 reads the same dichotomy from two quantitative angles using one matched episode (Almena, episode 157), where all five conditions are dropped at the same start with the same goal and all five complete it. *Left:* cumulative memory displacement $\Sigma|\Delta \mathbf{h}_2|$ up to step t (running L2 path length of $\mathbf{h}_2$ in the recurrent state space) plotted as a curve per condition. Blind’s curve ends at 1453 over 350 winding steps, while the four sighted conditions stop early at 309–467 over 79–101 direct steps, with the within-sighted ranking foveated $309 < \text{uniform } 330 < \text{fov-LP } 435 < \text{coarse } 467$. *Right:* mean decoder R^2 as a function of trajectory-step lag $|i - j|$, collapsing each top-row TGM heatmap into a single curve per condition. **The five curves fan out monotonically: blind drops to ≈ -1.2 within the first few steps (catastrophic transfer, the diagonal-only signature); coarse holds $R^2 \approx +0.18$ at small lag (the strongest stationary block); fov-LP, foveated, and uniform interpolate at $+0.13$, $+0.13$, and $+0.08$ respectively.** The sighted ranking matches the within-sighted spread of $\Sigma|\Delta \mathbf{h}_2|$ in the left panel and inverts the encoder-bandwidth ranking from §3.2: less encoder bandwidth produces a




Figure 4: **Temporal format: blind-vs-sighted dichotomy and within-sighted spectrum across three quantitative views.** *Top row:* step-by-step generalisation of a linear goal-vector decoder. Each 50×50 heatmap shows decoder R^2 when fit at step i (row) and evaluated at step j (column). A diagonal-only structure (blind) signals a step-by-step rotating basis: the decoder at step t does not transfer to step $t \pm 30$. A square-block structure (coarse) signals a stationary readout. Sighted-rich conditions interpolate. *Bottom row, left:* cumulative memory displacement $\Sigma|\Delta h_2|$ vs step-in-episode on the matched episode (Almena, episode 157; same start, same goal, all five conditions complete it), one curve per condition. Blind’s curve ends at 1453 over 350 winding steps; sighted conditions stop early at 309–467 over 79–101 direct steps (foveated $309 < \text{uniform } 330 < \text{fov-LP } 435 < \text{coarse } 467$). *Right:* mean decoder R^2 as a function of step lag $|i - j|$ — each top-row TGM heatmap collapsed to a single curve. The five curves fan out monotonically: blind drops to ≈ -1.2 (catastrophic transfer); coarse holds $R^2 \approx +0.18$ at small lag (strongest stationary block); fov-LP $+0.13$, foveated $+0.13$, uniform $+0.08$. The within-sighted ranking inverts encoder bandwidth: less bandwidth, more stationary readout. Per-unit autocorrelation τ values (Murray [Murray et al., 2014]) and a persistent-homology corroboration [Gardner et al., 2022] are reported in Appendix F.

more stationary memory readout, because memory has to do more work as a scene-anchor when the visual stream carries fewer high-frequency updates. The two views read one finding from different angles: blind’s memory state is a *slow-rotating integrator* reading from a step-by-step rotating basis; sighted units sit at a stationary readout — a *fast pendulum with a fixed axis* — whose stiffness scales inversely with encoder bandwidth. Murray-style per-unit autocorrelation timescales ($\tau \approx 12$ for blind, $\tau \approx 7$ –8 for sighted) and a topological corroboration via persistent homology [Gardner et al., 2022] are reported in Appendix F.

3.4 Consumption: causal intervention dissociates probe-readability from policy reliance

A linear probe tells us *where in memory the position code lives*, but not whether the policy actually uses it. The two could come apart in either direction: a condition might carry a probe-readable position code that the policy ignores (“readable but unused”), or might lack one yet behave as if it had one (“unused but readable”). Both cases would warn against reading probe results as direct policy mechanism. We test policy reliance with a paired same-scene different-goal shortcut protocol [Tolman, 1948]: re-init the recurrent memory between episodes vs. persist it; the SPL drop measures how much the policy depends on persistent memory. **Probe-readability and policy reliance dissociate in two off-diagonal cases.**

Coarse carries a linear position code yet has a small shortcut SPL drop (“readable but unused” — the policy under-weights its own readable code); uniform has no linear position code yet has the largest shortcut SPL drop (“unused but readable via a substitute” — the policy reads something else from memory) (Figure 5a). A direct readout test names the substitute: a scene-id classifier separates scenes substantially better in uniform than in blind, while within-scene linear position decoding goes the other way — uniform’s memory state appears to encode scene identity more




Figure 5: **Probe-readable position is not what drives behaviour.** (a) *Off-diagonal cases*: per-condition shortcut SPL drop (y -axis = policy reliance on persistent memory) vs. linear position decoding from memory (x -axis = probe-readability). Coarse sits in “readable but unused”, uniform in “unused but readable via a substitute”. (b) *The substitute uniform reads*: end-trajectory margin (distance from persistent end to OLD goal — distance to NEW goal) on same-floor failure cases. Only uniform ends closer to the previous episode’s goal — consistent with reading a scene-history anchor instead of the absent position code. Foveated-logpolar is excluded from (b): it has the most robust persistent-memory behaviour of the five conditions and therefore yields too few paired-failure cases ($n=2$) to estimate a margin.

strongly than position, the inverse of blind. Trajectory-level analysis (Figures 5b, 25) gives the candidate mechanism: only uniform’s persistent agent ends up closer to the *previous* episode’s goal than the new one — a visual / scene-history anchor (“head to where you were going last episode”) is one consistent reading.

Boundary checks: upstream and downstream causal interventions. Two further causal stress-tests characterise the consumption mechanism without separating conditions (Figure 6). *Memory-init transplant* (re-initialise the agent’s memory from the previous-episode end-state rather than from zero): all five conditions show a comparable SPL drop (Figure 6a). The cross-condition uniformity is itself the finding — every policy at $t=0$ expects a zero memory state, regardless of encoder. **Memory in this architecture is a dynamics-only activity-slot store, not a synaptic store [Dorrell et al., 2024]**: persistent memory is bound to the trajectory dynamics, not to the underlying weights. *GPS-sensor ablation* (zero GPS+compass from step 0) collapses all five conditions to near-zero success (Figure 6b): the position sensor is causally critical across the spectrum. The absolute drops cluster tightly across the four sighted conditions and do not cleanly track the magnitude vs. consumption split; we read this as a check on the upstream sensor’s causal role rather than independent ranking-level evidence for the integration-vs-substitute split established above. Boundary tests, classifier protocol, and the failed mid-rollout visual ablation are detailed in Appendix B.




Figure 6: **Causal interventions: memory is dynamics-bound, and the position sensor is causally critical across all conditions.** (a) *Memory-init transplant*: re-initialising memory from the previous episode’s end-state (instead of zero) drops SPL by a comparable amount across all five conditions. The cross-condition uniformity is the finding — every agent’s policy expects a zero memory state at $t=0$, so memory acts as an activity-slot store bound to the trajectory dynamics. (b) *GPS+compass ablation*: zeroing GPS and compass from step 0 collapses success to near-zero across all conditions. The drops cluster tightly across the four sighted conditions; blind drops less because its baseline is already lower. The position sensor is the upstream causal source whose integration the memory encodes.




Figure 7: **Synthesis: three views of one capacity allocation.** *X-axis*: top-layer GPS R^2 on h_2 (magnitude). *Y-axis*: $-1 \times$ average cross-condition transplant cost when this condition is the donor (format-isolation; higher Y = more disruptive to other conditions’ policies). *Marker size*: second-episode shortcut SPL drop % (consumption; larger = more policy dependence on persistent recurrent memory). Within each axis, multiple convergent measurements — e.g. information-allocation, scene-invariance, predictive-horizon, place-cell, and excursion-fragility findings (Figures 3, 10) — sharpen the *interpretation* but are not separate views of variation.

4 Discussion

4.1 Synthesis: three views of one capacity allocation

A reader could imagine three independent things going on across our five conditions — they differ in how much position lives in memory, in how that position is encoded, and in whether the policy uses it. We argue the three are views of one allocation: the visual encoder and the integration of

the position sensor compete for finite memory capacity, and the encoder’s bandwidth tilts the balance. **Bottleneck conditions retain a linear position code in memory, rich-encoder conditions reallocate to the visual route, and the policy reads from whichever route the memory makes available.** Each condition occupies a distinct (X, Y, size) coordinate in Figure 7: magnitude on X (linear position decode, §3.1), format-isolation on Y (cross-condition transplant cost, §3.4), policy reliance as marker size (shortcut SPL drop, §3.4). Foveation occupies a hybrid position: it shares uniform’s rich-encoder magnitude but separates on every other measure — most strikingly, foveated agents recover a held-out-MP3D position code that uniform agents do not, consistent with the visual and integration routes being independent capacity sinks rather than two ends of one knob. Coarse and uniform are both on the same magnitude side of the diagram for different reasons: coarse’s encoder cannot supply spatially-resolved features (so capacity defaults to integration), while uniform’s encoder supplies rich features (so capacity is reallocated to the visual route, leaving integration unattended). The two are different regimes of one allocation, not the same mechanism.

4.2 Biological precedent and predictions

The hypothesis “sensor niche shapes spatial memory” has long-standing comparative-cognition precedent — rodent dark-rearing [Chen et al., 2016], bat cross-modality remapping [Geva-Sagiv et al., 2016], congenitally blind humans with larger anterior hippocampal volume [Fortin et al., 2008]. We do not claim that our silicon results *demonstrate* the bio-AI parallel; we offer them as patterns from controlled within-architecture variation that one would expect if the hypothesis held in animals. Loosely mapped: blind $\leftrightarrow$ subterranean rodents [Kimchi et al., 2004]; coarse $\leftrightarrow$ single-channel sensory summaries (echolocation, infrared pit-detection); foveated $\leftrightarrow$ primate / raptor (gaze-fixed); uniform $\leftrightarrow$ insect compound eye / zebrafish near-isotropic sampling [Geva-Sagiv et al., 2015, Toledo et al., 2020] — interpretive scaffolding, not mechanistic identity claims.

Two specific bio-AI bridges land cleanly under the framework of Sanders et al. [2020], in which place fields remap when sensory evidence supports distinct hidden states. (i) Our LOSO catastrophe in rich-encoder conditions matches global-remapping (rich per-step visual evidence supports scene-specific hidden states; the linear position direction reorganises across rooms); bottleneck conditions exhibit no remapping (a single scene-invariant code). (ii) Blind’s transient-code temporal-generalisation signature and longer intrinsic timescale [King and Dehaene, 2014, Murray et al., 2014] reproduce the cortical-hierarchy slow-integrator pole. Both are “the same metric one would run on monkey or human cortex, applied unchanged to a recurrent navigator”. Three falsifiable predictions follow: (i) within-species manipulations of effective per-step visual variety (gaze restriction, dark-rearing, prosthetic compound vision) should shift hippocampal coding along the no-/global-remapping spectrum; (ii) gaze-fixed primates should show coding closer to rodent place cells than to free-gaze primate view cells, with Wirth et al. [2017] the closest existing dataset; (iii) the Skaggs, LOSO, predictive-horizon, MINE, and subspace-divergence findings, catalogued separately in cognitive map theory [Tolman, 1948, O’Keefe and Nadel, 1978, Stachenfeld et al., 2017, Whittington et al., 2020], may all be views of a single mechanism. Whether biological hippocampi obey a finite-capacity allocation principle remains an open empirical question.

4.3 Implications and scope

The principle has a simple geometric reading. A finite memory has a fixed number of dimensions to spend; how it spends them depends on what its input already provides. When the encoder delivers rich, near-complete spatial information at every step, the memory has no reason to maintain a separate linearly-readable position code, and the policy is free to commit those dimensions to other content. When the encoder delivers almost nothing, the memory must reconstruct position by integrating the action stream over time, and the resulting code is forced into a form the linear action head can consume. **This is the empirical pattern of §3.1–§3.2, and it is what the invariance–disentanglement reading of the information bottleneck [Tishby et al., 1999, Achille and Soatto, 2018] predicts:** bandwidth-restricted memory is the architectural mechanism that drives the trained representation toward scene-invariant rather than scene-conditional spatial encoding (Figure 3b). The framing has a quantitative skeleton: a hidden state with effective rank d allocates dimensions among (i) directly carrying encoder-derived features (capacity d_e), (ii) integrating sensor history (capacity d_i), and (iii) carrying task-relevant non-position content (capacity d_o), subject to $d_e + d_i + d_o \lesssim d$. Bottleneck encoders force $d_e \rightarrow 0$, leaving d_i to dominate the linear position code; rich encoders enlarge d_e , so policy gradient minimises d_i down to whatever residual the

integration route still needs. Our Procrustes / participation-ratio / eigenspectrum- α measurements collectively probe d (effective rank ≈ 4 for blind, lower for coarse where the encoder collapse shrinks d overall, Appendix E.3); the magnitude collapse follows from d_i shrinking as d_e grows. We use IB descriptively rather than mechanistically; Saxe et al. [2019]’s critique of IB-deep-learning theory cautions against absolute MI-magnitude claims, but the cross-condition *ordering* reported here (linear collapse, non-linear-probe recovery, mutual-information rank in Appendix D.1) is robust to those concerns.

The principle is also architecture-agnostic: a content-level allocation should appear in any system where a finite-capacity memory reads a pretrained encoder. Replacing recurrent gating with attention-over-history should not change the headline ordering. In transformer-based navigators (e.g. Zeng et al. 2024; cf. Whittington et al. 2022 for the transformer–hippocampus correspondence), where “memory” lives in the KV cache or attention over rollout history, the principle predicts the same monotonic anti-correlation reported in Figure 2a: **attention-readable position codes under bandwidth-restricted encoders, attention-internal position codes only non-linearly decodable under rich encoders**. The cleanest falsification target is direct — train a transformer navigator under our five sensor conditions; the opposite pattern would falsify architecture-independence and bound the principle to recurrent-state systems. The same logic re-frames recent VLM spatial-reasoning failures [Ramakrishnan et al., 2025]: rich pretrained encoders may divert capacity *away from* the downstream substrate’s spatial code, rather than merely fail to teach it, and bandwidth-restricting perception (foveation, log-polar) may be *cognition-enabling* rather than only compute-saving. The relevant variable is encoder spatial-feature *variety* per step rather than raw input bandwidth (Appendix I); a finer multi-point sweep within recurrent agents is left to future work.

As a partial test of architecture-independence, we ported the encoder-bandwidth design to a different environment, encoder family, and training regime (Figure 8). Frozen DINOv2-Base [Oquab et al., 2024] patch features at five resolutions — mirroring our Habitat chassis (BLIND, COARSE, FOVEATED, UNIFORM, FOVEATED-LOGPOLAR) — feed a small recurrent memory trained on Memory Maze [Pasukonis et al., 2023] 9×9 offline trajectories with next-feature prediction. We then probe `agent_pos` from the memory state (Ridge linear and 4-layer non-linear probe, frame-level CV with 80/20 random split, eval window steps 250–500). Three pre-registered hypotheses land cleanly. (1) *Linear R^2 rises monotonically with encoder bandwidth*: BLIND 0.01 $\rightarrow$ COARSE 0.38 $\rightarrow$ FOVEATED 0.41 $\rightarrow$ FOVEATED-LOGPOLAR 0.41 $\rightarrow$ UNIFORM 0.45. (2) *Non-linear R^2 saturates near 0.998 for all sighted conditions* and falls to chance for BLIND (−0.0003). (3) *The non-linear gap (non-linear − linear) decreases monotonically with bandwidth*: COARSE +0.62, FOVEATED +0.59, FOVEATED-LOGPOLAR +0.59, UNIFORM +0.55. **Format-shift recovery is largest where linear readability is smallest — the same dichotomy our Habitat result reports, in a different environment, encoder family, and policy regime.** *Direction relative to Habitat.* The Memory-Maze setup deliberately omits the position-sensor channel our Habitat agents receive; without that explicit positional anchor, the encoder is the only sensory route to absolute position, and BLIND cannot decode it (no signal, no integration anchor). The bandwidth-allocation *tradeoff* (where blind’s integration route compensates) thus requires a competing route; the format-shift dynamic (linear-readable code emerges from a non-linear substrate as bandwidth grows) generalises across both regimes. The two together place the broad “encoder structure shapes memory format” claim on a different environment, encoder family, and policy regime, while clarifying that the specific direction (high-bandwidth-rich vs. bottleneck-rich) is mediated by which sensory channels carry positional information.

4.4 Limitations

(i) *Single seed per condition*. Following the comparative-cognition modelling convention [Wijmans et al., 2023, Banino et al., 2018, Cueva and Wei, 2018, Schaeffer et al., 2023]. Cross-condition rigour comes from convergent multi-method evidence: linear and non-linear probes, CKA, Procrustes, subspace divergence, LOSO, lag- k , 5×5 memory transplant, Skaggs spatial information, and excursion-forgetting all rank the conditions in the same direction — unlikely under single-seed noise alone — supplemented by an independent-retraining replication of the blind condition (Appendix A). Per-condition seed-variance error bars over a multi-seed grid are left for follow-up.

(ii) *Method-family scope*. Methods we do not report and that could change a finding’s strength: dynamical-systems analysis (fixed-point / slow-point structure [Sussillo and Barak, 2013]; Dynam-




Figure 8: **Architecture-agnostic confirmation: format shift reproduces on Memory Maze with a frozen DINOv2-Base encoder feeding a small LSTM.** (A) Linear (Ridge, best of $\alpha \in \{0.1, 1, 10, 100, 1000\}$, blue) and 4-layer MLP (red) eval R^2 for agent_pos from the LSTM hidden state, frame-level random CV (80/20), eval window steps 250–500 of $T=500$ Memory-Maze trajectories. Linear R^2 is monotone-increasing with encoder bandwidth (BLIND 0.01 → UNIFORM 0.45); MLP R^2 saturates at ≈ 0.998 for all sighted conditions and is at chance for BLIND. (B) The non-linear gap (MLP–linear) is monotone-decreasing with bandwidth: format-shift recovery is largest where linear readability is smallest (COARSE +0.62 vs. UNIFORM +0.55), the same dichotomy our Habitat agent shows in Figure 3a. The opposite linear- R^2 direction relative to Habitat reflects the absence of GPS/compass channels in the Memory-Maze setup — without an alternative positional anchor, the encoder is the only route, and BLIND carries no signal.

ical Similarity Analysis [Ostrow et al., 2023] as a dynamics-aware Procrustes alternative) would speak to format at the level of computational primitives. We piloted DSA on our memory-state time-series (10-d PCA pre-reduction, 50-step trajectories) and found the cross-condition signal sensitive to hyperparameter choices (n_{delays} , rank), rollout protocol (excursion-warmup vs. deterministic-rollout), and episode-phase alignment (blind has median episode length $\sim 2.5\times$ sighted, so first- T -step samples occupy different behavioural phases). Under the cleanest protocol (deterministic-rollout, all conditions, $T=50$), sighted clustered at floor ~ 0.10 angular distance while blind sat further out (~ 0.27); but blind’s within-condition split-half DSA was also ~ 0.23 , leaving the signal-to-noise ratio insufficient to publish a definitive ranking. Information-theoretic estimators of $I(h_t; \text{pos}_{t+k})$ would replace decodability R^2 with a nat-scale predictive-horizon measure; gaze-placement controls and a σ_{max} strength sweep within the foveated condition would test the rich-encoder content. All four are left for follow-up.

(iii) *Architecture, algorithm, and task scope.* We use a recurrent (LSTM) backbone with DD-PPO, PointGoal, and a ResNet-18 visual encoder, matching the canonical pipeline of [Wijmans et al., 2023, Banino et al., 2018, Cueva and Wei, 2018] for direct comparability; the architecture choice is replication-driven rather than central to the claim, and the principle is stated at the encoder–memory interface (architecture-agnostic). A reasonable concern is whether the encoder–memory race is a recurrent-gating artefact (e.g. vanishing-gradient limitations specific to LSTMs); two within-paper checks argue against it: (a) non-linear (2-layer MLP) probes recover position in rich-encoder conditions ($R^2 \in [0.35, 0.79]$), so the memory *does not* discard position — only the linear-readable form tracks encoder bandwidth, a property of the linear policy head; (b) the cross-condition memory transplant (§3.4) is a behavioural test outside any probe framework and reproduces the same ordering. The encoder–memory race is in principle architecture-agnostic; the cleanest falsification target is a transformer-based navigator under the same conditions, identified in §4 but not verified here. The magnitude and format patterns could also shift under different RL algorithms, different navigation tasks (ObjectGoal, ImageGoal, exploration), or different visual encoders.

5 Conclusion

A five-condition visual-sensor ablation in a recurrent PointNav agent shows sensor structure as a representational-content variable whose effects cohere under a single capacity-allocation principle measured as three views: *magnitude* (bandwidth–allocation tradeoff between encoder and recurrent-integration routes), *format* (within-condition scene-invariance, across-condition non-interchangeable subspaces with asymmetric memory transplant, and a King–Dehaene-style temporal-format dissociation in which the blind code rotates over time while the coarse code stabilises), and *consumption* (probe-readable position dissociates from policy reliance). The bandwidth–allocation ordering is preserved on held-out MP3D, blind units integrate over a $2.1\text{--}2.5\times$ longer Murray-cortical-hierarchy timescale than every sighted condition, and the pattern is at high level consistent with independent animal-navigation findings on hippocampal compensation under sensory deprivation, cross-modality remapping, and the no-/global-remapping axis under hidden-state inference [Sanders et al., 2020] — framed as an open bio–AI hypothesis rather than a demonstrated parallel.

Methodologically, we view the paper as a step toward a *comparative cognitive neuroscience of artificial navigation agents*: a research programme in which matched-condition populations of trained agents are analysed with the same set of metrics — population geometry, intrinsic timescales, dynamical fixed points, predictive-coding residuals, persistent-homology of the state manifold, mixed-selectivity / splitter analyses, communication subspaces, transfer entropy — that is used to map biological cortex and hippocampus. We piloted this analytic suite on the present sensor axis (Figures 4, 4, 8): temporal-axis methods (TGM, Murray timescales) and topological methods (persistent homology) detect strong cross-condition signal aligned with capacity allocation, and the format-shift dynamic itself reproduces in a frozen encoder feeding a small LSTM on a different environment. Beyond this paper, the same suite can be applied to architecture-axis populations (recurrent vs. transformer world models), training-curriculum axes (sparse vs. dense reward), or biological-vs-silicon axes (LSTMs trained on egocentric video versus rodent CA1 recordings under matched stimuli). The goal is not to claim that artificial agents *are* brains, but that the analytical methods used to ask “how does this neural system represent its world?” transfer to artificial systems where the relevant variables can be held fixed — and that the cross-condition contrasts produce theoretically informative signatures (capacity allocation, format shifts, dynamical regimes) of the same kind comparative neuroscience uses to understand sensor-niche adaptation in animals. The encoder–memory race is stated as an interface-level claim and should in principle generalise to transformer-based navigators (the cleanest falsification target, §4). Cross-condition rigour rests on multi-method convergence and an independent-retraining blind replication (Appendix A); a multi-seed grid, finer encoder-bandwidth sweep, and gaze-placement controls are reported as future work.

References

- Alessandro Achille and Stefano Soatto. Emergence of invariance and disentanglement in deep representations. *Journal of Machine Learning Research*, 19(50):1–34, 2018.
- Alessio Ansuini, Alessandro Laio, Jakob H. Macke, and Davide Zoccolan. Intrinsic dimension of data representations in deep neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- Andrea Banino, Caswell Barry, Benigno Uria, Charles Blundell, Timothy Lillicrap, Piotr Mirowski, Alexander Pritzel, Martin J Chadwick, Thomas Degris, Joseph Modayil, et al. Vector-based navigation using grid-like representations in artificial agents. *Nature*, 557(7705):429–433, 2018.
- Mohamed Ishmael Belghazi, Aristide Baratin, Sai Rajeswar, Sherjil Ozair, Yoshua Bengio, Aaron Courville, and R Devon Hjelm. Mutual information neural estimation. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*, pages 531–540, 2018.
- Boyuan Chen, Zhuo Xu, Sean Kirmani, Danny Driess, Pete Florence, Brian Ichter, Dorsa Sadigh, Leonidas Guibas, and Fei Xia. SpatialVLM: Endowing vision-language models with spatial reasoning capabilities. In *Computer Vision and Pattern Recognition (CVPR)*, 2024.
- Guifen Chen, John A. King, Neil Burgess, and John O’Keefe. Absence of visual input results in the disruption of grid cell firing in the mouse. *Current Biology*, 26(17):2335–2342, 2016.

- Christopher J Cueva and Xue-Xin Wei. Emergence of grid-like representations by training recurrent neural networks to perform spatial localization. In *International Conference on Learning Representations (ICLR)*, 2018.
- William Dorrell, Soraya El-Gaby, Timothy E. J. Behrens, and James C. R. Whittington. A tale of two algorithms: Structured slots explain prefrontal sequence memory and are unified with hippocampal cognitive maps. *Neuron*, 2024. doi: 10.1016/j.neuron.2024.10.022.
- Mohamed El Banani, Amit Raj, Kevis-Kokitsi Maninis, Abhishek Kar, Yuanzhen Li, Michael Rubinstein, Deqing Sun, Leonidas Guibas, Justin Johnson, and Varun Jampani. Probing the 3D awareness of visual foundation models. In *Computer Vision and Pattern Recognition (CVPR)*, 2024.
- Russell A. Epstein, Eva Zita Patai, Joshua B. Julian, and Hugo J. Spiers. The cognitive map in humans: spatial navigation and beyond. *Nature Neuroscience*, 20(11):1504–1513, 2017.
- Elena Facco, Maria d’Errico, Alex Rodriguez, and Alessandro Laio. Estimating the intrinsic dimension of datasets by a minimal neighborhood information. *Scientific Reports*, 7(12140), 2017.
- Madeleine Fortin, Patrice Voss, Catherine Lord, Maryse Lassonde, Jens Pruessner, Dave Saint-Amour, Constant Rainville, and Franco Lepore. Wayfinding in the blind: larger hippocampal volume and supranormal spatial navigation. *Brain*, 131(11):2995–3005, 2008.
- Richard J. Gardner, Erik Hermansen, Marius Pachitariu, Yoram Burak, Nils A. Baas, Benjamin A. Dunn, May-Britt Moser, and Edvard I. Moser. Toroidal topology of population activity in grid cells. *Nature*, 602(7895):123–128, 2022. doi: 10.1038/s41586-021-04268-7.
- Maya Geva-Sagiv, Liora Las, Yossi Yovel, and Nachum Ulanovsky. Spatial cognition in bats and rats: from sensory acquisition to multiscale maps and navigation. *Nature Reviews Neuroscience*, 16(2):94–108, 2015.
- Maya Geva-Sagiv, Sandro Romani, Liora Las, and Nachum Ulanovsky. Hippocampal global remapping for different sensory modalities in flying bats. *Nature Neuroscience*, 19(7):952–958, 2016.
- Stefan Heimersheim and Neel Nanda. How to use and interpret activation patching. *arXiv preprint arXiv:2404.15255*, 2024.
- John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2019.
- Tali Kimchi, Ariane S Etienne, and Joseph Terkel. A subterranean mammal uses the magnetic compass for path integration. *Proceedings of the National Academy of Sciences*, 101(4):1105–1109, 2004.
- Jean-Rémi King and Stanislas Dehaene. Characterizing the dynamics of mental representations: the temporal generalization method. *Trends in Cognitive Sciences*, 18(4):203–210, 2014. doi: 10.1016/j.tics.2014.01.002.
- Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *International Conference on Machine Learning (ICML)*, 2019.
- John D. Murray, Alberto Bernacchia, David J. Freedman, Ranulfo Romo, Jonathan D. Wallis, Xinying Cai, Camillo Padoa-Schioppa, Tatiana Pasternak, Hyojung Seo, Daeyeol Lee, and Xiao-Jing Wang. A hierarchy of intrinsic timescales across primate cortex. *Nature Neuroscience*, 17(12):1661–1663, 2014. doi: 10.1038/nn.3862.
- John O’Keefe and Lynn Nadel. *The Hippocampus as a Cognitive Map*. Oxford University Press, 1978.
- Maxime Oquab, Timothée Darcet, Théo Moutakanni, Huy Vo, Marc Szafraniec, Vasil Khalidov, Pierre Fernandez, Daniel Haziza, Francisco Massa, Alaaeldin El-Nouby, et al. Dinov2: Learning robust visual features without supervision. <https://arxiv.org/abs/2304.07193>, 2024.

- Saurav Orozco et al. Emergent world representations in OpenVLA. *arXiv preprint arXiv:2509.24559*, 2025.
- Mitchell Ostrow, Adam Eisen, Leo Kozachkov, and Ila Fiete. Beyond geometry: Comparing the temporal structure of computation in neural circuits with dynamical similarity analysis. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Jurgis Pasukonis, Timothy Lillicrap, and Danijar Hafner. Evaluating long-term memory in 3d mazes. In *International Conference on Learning Representations (ICLR)*, 2023.
- Santhosh Kumar Ramakrishnan, Erik Wijmans, Philipp Krähenbühl, and Vladlen Koltun. Does spatial cognition emerge in frontier models? In *International Conference on Learning Representations (ICLR)*, 2025.
- Edmund T. Rolls. Why do primates have view cells instead of place cells? *Trends in Cognitive Sciences*, 28(9):862–874, 2024.
- Ruth Rosenholtz. Capabilities and limitations of peripheral vision. *Annual Review of Vision Science*, 2:437–457, 2016.
- Honi Sanders, Matthew A Wilson, and Samuel J Gershman. Hippocampal remapping as hidden state inference. *eLife*, 9:e68480, 2020.
- Manolis Savva, Abhishek Kadian, Oleksandr Maksymets, Yili Zhao, Erik Wijmans, Bhavana Jain, Julian Straub, Jia Liu, Vladlen Koltun, Jitendra Malik, Dhruv Batra, and Devi Parikh. Habitat: A platform for embodied AI research. In *International Conference on Computer Vision (ICCV)*, 2019.
- Andrew M Saxe, Yamini Bansal, Joel Dapello, Madhu Advani, Artemy Kolchinsky, Brendan D Tracey, and David D Cox. On the information bottleneck theory of deep learning. *Journal of Statistical Mechanics: Theory and Experiment*, 2019(12):124020, 2019. ICLR 2018.
- Rylan Schaeffer, Mikail Khona, Tzuhsuan Ma, Cristobal Eyzaguirre, Sanmi Koyejo, and Ila Fiete. Self-supervised learning of representations for space generates multi-modular grid cells. *arXiv preprint arXiv:2311.02316*, 2023.
- William E Skaggs, Bruce L McNaughton, Matthew A Wilson, and Carol A Barnes. Theta phase precession in hippocampal neuronal populations and the compression of temporal sequences. *Hippocampus*, 6(2):149–172, 1996.
- Kimberly L Stachenfeld, Matthew M Botvinick, and Samuel J Gershman. The hippocampus as a predictive map. *Nature Neuroscience*, 20(11):1643–1653, 2017.
- Hans Strasburger, Ingo Rentschler, and Martin Jüttner. Peripheral vision and pattern recognition: A review. *Journal of Vision*, 11(5):13, 2011.
- Carsen Stringer, Marius Pachitariu, Nicholas Steinmetz, Matteo Carandini, and Kenneth D Harris. High-dimensional geometry of population responses in visual cortex. *Nature*, 571(7765):361–365, 2019.
- David Sussillo and Omri Barak. Opening the black box: Low-dimensional dynamics in high-dimensional recurrent neural networks. *Neural Computation*, 25(3):626–649, 2013.
- Naftali Tishby, Fernando C. Pereira, and William Bialek. The information bottleneck method. In *Proceedings of the 37th Annual Allerton Conference on Communication, Control, and Computing*, 1999.
- Sivan Toledo, David Shohami, Ingo Schiffner, Emmanuel Lourie, Yotam Orchan, Yoav Bartan, and Ran Nathan. Cognitive map-based navigation in wild bats revealed by a new high-throughput tracking system. *Science*, 369(6500):188–193, 2020.
- Edward C Tolman. Cognitive maps in rats and men. *Psychological Review*, 55(4):189–208, 1948.

- Shengbang Tong, Zhuang Liu, Yuexiang Zhai, Yi Ma, Yann LeCun, and Saining Xie. Eyes wide shut? exploring the visual shortcomings of multimodal LLMs. In *Computer Vision and Pattern Recognition (CVPR)*, 2024.
- James C. R. Whittington, Joseph Warren, and Timothy E. J. Behrens. Relating transformers to models and neural representations of the hippocampal formation. In *International Conference on Learning Representations (ICLR)*, 2022.
- James CR Whittington, Timothy H Muller, Shirley Mark, Guifen Chen, Caswell Barry, Neil Burgess, and Timothy EJ Behrens. The tolman–eichenbaum machine: unifying space and relational memory through generalization in the hippocampal formation. *Cell*, 183(5):1249–1263, 2020.
- Erik Wijmans, Abhishek Kadian, Ari Morcos, Stefan Lee, Irfan Essa, Devi Parikh, Dhruv Batra, and Manolis Savva. DD-PPO: Learning near-perfect pointgoal navigators from 2.5 billion frames. In *International Conference on Learning Representations (ICLR)*, 2020.
- Erik Wijmans, Manolis Savva, Irfan Essa, Stefan Lee, Ari Morcos, and Dhruv Batra. Emergence of maps in the memories of blind navigation agents. In *International Conference on Learning Representations (ICLR)*, 2023.
- Alex H. Williams, Erin Kunz, Simon Kornblith, and Scott W. Linderman. Generalized shape metrics on neural representations. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- Sylvia Wirth, Pierre Baraduc, Angélique Planté, Serge Pinede, and Jean-René Duhamel. Gaze-informed, task-situated representation of space in primate hippocampus during virtual navigation. *PLOS Biology*, 15(2):e2001045, 2017.
- Fei Xia, Amir R Zamir, Zhiyang He, Alexander Sax, Jitendra Malik, and Silvio Savarese. Gibson Env: Real-world perception for embodied agents. In *Computer Vision and Pattern Recognition (CVPR)*, 2018.
- Jihan Yang, Shusheng Yang, Anjali W. Cai, Boyi Zhao, Wenwen Zhao, Yifan Hu, Joy Hsu, Ellis L. Brown, Shengbang Tong, and Lerrel Yang. Thinking in space: How multimodal large language models see, remember, and recall spaces. In *Computer Vision and Pattern Recognition (CVPR)*, 2025.
- Kuo-Hao Zeng, Zichen Zhang, Kiana Ehsani, Rose Hendrix, Jordi Salvador, Alvaro Herrasti, Ranjay Krishna, Aniruddha Kembhavi, and Roozbeh Mottaghi. PoliFormer: Scaling on-policy RL with transformers results in masterful navigators. *arXiv preprint arXiv:2406.20083*, 2024.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [Yes]

Justification: The abstract and §1 state the three hypotheses H1/H2/H3 with explicit signatures (what pattern supports each), and preview the three findings with specific magnitudes (path-history $R_s=0.56$, CKA <0.004 , compass $R_s=0.94$ vs GPS $R_s=0.03$). These are expanded in §4 and §5 with no stronger claims than the data support.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: §5.4 explicitly lists six limitations (ivi) covering checkpoint staleness, single-seed H3, missing proposed analyses, foveation-model simplification, linear-probes-only, and LSTM-architecture scope.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper reports empirical results; it contains no theoretical claims or proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: §3.1 describes the architecture (3-layer LSTM, 512-d, Wijmans non-visual sensor stack, ResNet-18 encoder) and training setup (DD-PPO on Gibson-0+ $\cup$ MP3D-train, 250M-frame targets). §3.2§3.6 describe all probing and cross-condition protocols. Appendix A details the gaze-diversity ablation. Code will be released upon publication.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: §3.7 commits to open-source release of the training and probing code. Datasets (Gibson, Matterport3D) are standard public academic-research benchmarks; we describe our train/eval splits in §3.1.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: §3.1 specifies training pool and eval set; §3.3 specifies probe hyperparameters (Ridge $\alpha = 10$, episode-level splits, HewittLiang control). Appendix A gives the ablation’s loss formulation and coefficient.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: We report single-seed numbers per condition. This is documented as Limitation (ii) in §5.4. Multi-seed replication for tighter confidence intervals is in the immediate follow-up.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: §3.7 reports per-condition training budget (57 days on 20EV100 for 250M frames) and probing budget (<1h per condition on one V100).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research uses public academic-research benchmarks, publicly-released open-source RL frameworks, and involves no human subjects or deployment. We also report and remediate a numerical-stability bug upstream (Appendix B).

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: §5.3 discusses implications sensor design as a tunable axis for interpretable embodied agents. Negative societal impacts are minimal: the work is fundamental research on simulated indoor navigation with no direct deployment pathway.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: We release neither a high-risk dataset nor a generative model. Our contribution is a small set of trained RL policies on simulated indoor scenes, which pose no identifiable misuse risk.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Gibson Database and Matterport3D are used under their released academic-research licenses; Habitat-Sim and habitat-lab are MIT-licensed. Our use complies with their terms; citations §3.1.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: Code will be released with documentation upon publication. Figures are novel; they are released under the paper’s license alongside the code. The gaze-diversity auxiliary loss and NaN-sanitisation patch are small standalone modules with docstrings.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper involves no crowdsourcing and no human-subjects data.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: No human-subjects research; no IRB/equivalent review was required.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs were used only for writing assistance on the manuscript. They are not part of the core research methodology; all empirical results come from standard RL training and linear probing.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.




Figure 9: **Behavioural convergence (DD-PPO training curves).** TensorBoard event-file scalars across the 250M-frame training budget. All four sighted conditions (COARSE, FOVEATED, FOVEATED-LOGPOLAR, UNIFORM) converge to SPL 0.82–0.85 and success > 97% by ~ 100 M frames; BLIND converges to SPL ~ 0.50 over a longer horizon. Probing windows fall well after convergence for all conditions, so the encoded-content gap reported in §3 cannot be attributed to differential under-training.

A Reproducibility: hyperparameters and implementation notes

Training hyperparameters (DD-PPO, all five conditions identical) 4-GPU torchrun; num_steps=256, learning rate lr=2.5e-4 with lr_decay=False; num_environments=16/worker; total budget 250M frames per condition. All other hyperparameters follow the Habitat-baselines DD-PPO defaults from Wijmans et al. [2020].

Architectural details (sensor encoding) Each non-visual sensor (goal-in-start-frame, GPS, compass, close-to-goal scalar) is projected to 32-d via a per-sensor linear layer; the previous-action embedding is also 32-d (nn.Embedding with action-space cardinality + 1 for the no-prev-action token). The four sensor projections are concatenated with the previous-action embedding and the (flattened) visual encoder output to form the LSTM Layer-0 input vector. This matches the Wijmans et al. [2020] Wijmans-default sensor stack.

Implementation notes (foveated condition) The foveated condition disables the running-mean-and-variance input normaliser (a per-channel running normalisation layer applied to the visual stream before the ResNet-18 backbone) to avoid an in-place autograd conflict between the running-buffer update and the foveation-blur backward pass. Behavioural success and SPL are unaffected by this choice; an ablation that re-enables the normaliser (with the autograd conflict patched) is available for future work that wants to revisit it.

Implementation notes (numerical stability) Two safeguards protect against rare numerical instabilities arising from off-navmesh episodes: (i) a pre-update gradient sanitisation pass replaces NaN/Inf gradients with zero before each PPO step; (ii) a forward-looking geodesic-distance clamp prevents reward NaNs when the agent steps off the navmesh. Both are no-ops on healthy trajectories and apply identically to all five conditions.

Log-polar foveation transform The log-polar control (Appendix I) applies a differentiable log-polar resampling step before the ResNet-18 backbone: the 128×128 avg-pooled image is mapped onto a 64×64 grid with $n_r=64$ log-spaced radial samples and $n_\theta=64$ angular samples, mimicking the cortical magnification of mammalian primary visual cortex. Visual front-end only; LSTM, sensors, and training pipeline are identical to the foveated condition.

B Probing protocol details

Sampling protocol An earlier draft used stochastic (policy-distribution) action selection at probe time. Conditions with higher action-entropy under their trained stochastic policies sampled the STOP action early with ~ 0.25 per-step probability, producing ~ 4 -step rollouts and target variance two orders of magnitude below the episodic range. The resulting probe R^2 values were trivially high across all conditions (any predictor fits a near-constant target). This confounded cross-condition comparisons; we switched to the deterministic (argmax) protocol used here and elsewhere in our pipeline. All paper results use the deterministic protocol.

Length-matching ablation Under deterministic rollouts, episode lengths span ~ 100 –400 mean steps depending on condition. We use 5-fold episode-level CV to keep splits balanced per scene

rather than length-matching across conditions. A controlled length-matching ablation (truncate every condition to a common length) is left for follow-up; we expect the magnitude ordering to be preserved because bottleneck conditions decode at high R^2 even at the smallest sub-sample, but have not verified.

Cross-heading probe The proposal included a cross-heading generalisation probe (train on heading A , test on opposite heading). Our implementation returned “insufficient samples” sentinels at our probe sizes; not in our results.

LSTM Layer 0 sensor branch The non-visual sensor stack (goal-in-start-frame, GPS, compass, close-to-goal scalar, prev-action) projects each sensor to 32-d, concatenates them with the visual encoder’s flattened output, and feeds the result as the LSTM’s Layer 0 input. The GPS embedding therefore enters the LSTM via this concatenation; this explains why Figure 2c shows linearly-decodable GPS at Layer 0 even for conditions whose visual encoder discards GPS entirely. The full per-layer breakdown (linear R^2 at $\mathbf{h}_0, \mathbf{h}_1, \mathbf{h}_2$ and at the cell states $\mathbf{c}_0, \mathbf{c}_1, \mathbf{c}_2$) is in Figure 11: Layer 0 carries GPS for every condition; Layer 1 dips for everyone; Layer 2 separates blind/coarse from foveated/uniform; cell states do not carry the GPS code as cleanly as hidden states. The headline “top-layer linearly-decodable GPS” claim is therefore specifically about $\mathbf{h}_2$, not the cell pathway.

Non-linear (MLP) probe details For each condition we fit a 2-layer MLP (hidden 256–128, ReLU, L_2 weight decay 10^{-4} , sklearn `MLPRegressor` with early-stopping) on the same top-layer hidden states and GPS labels as the linear Ridge probe, with the same 5-fold episode-level CV. The full probe-depth sweep (linear $\rightarrow$ MLP-1 $\rightarrow$ MLP-2 $\rightarrow$ MLP-4) shown in Figure 3a establishes a depth-graded recovery: bottleneck conditions plateau near $R^2 \approx 0.95$ from depth-0 (linear suffices), while rich-encoder conditions ramp up sharply with probe depth. Top-layer MLP-2 recovery (Ridge linear $\rightarrow$ MLP-2): blind 0.92 $\rightarrow$ 0.98, coarse 0.51 $\rightarrow$ 0.75, FOVEATED 0.06 $\rightarrow$ 0.79, FOVEATED-LOGPOLAR $-0.07 \rightarrow 0.63$, UNIFORM $-2.28 \rightarrow 0.35$. The MLP probe alone cannot disentangle “non-linear position code” from “MLP regressing through position-correlated features (wall distance, scene identifiers, episode-phase markers)”; the log-polar foveation control (Appendix I) targets that question at one intermediate point.

C Alternative-account rebuttals (full)

Task-competence differences. All conditions solve PointGoal at 93–99% success (Table 1); training curves (Figure 9) confirm that all five conditions converge behaviourally before the probing window. The encoded-content gap cannot be attributed to navigation skill.

Probe-capacity / probe-fitting artefacts. A direct test: distance-to-goal (DtG) decodes robustly across all four conditions ($R^2 \geq 0.81$; Figure 18) while the goal-direction component is at chance everywhere (PointGoal provides the goal vector as a per-step input, so the policy has no incentive to re-encode it). The linear probe is therefore capable of fitting position-correlated structure when it exists; the chance-level rich-encoder GPS R^2 is a hidden-state property, not a probe artefact. The real-vs-permuted-label gap (Hewitt–Liang selectivity) likewise tracks GPS R^2 .

Sample-size / trajectory-distribution artefacts. CV σ spans come from the same 500-episode pool with ~ 100 –400-step rollouts under identical deterministic evaluation; the protocol confound that haunted an earlier draft is documented and retired in Appendix B.

Linear-probe artefacts. Three complementary checks. (a) *Behavioural*: the memory-transplant intervention goes outside the probe framework and reproduces the bottleneck-vs-rich-encoder ordering; the 5×5 matrix’s asymmetry adds structure no linear-similarity measure can read. (b) *Non-linear (MLP) probe*: a 2-layer MLP recovers $R^2 \in [0.35, 0.79]$ from rich-encoder top-layer states (Appendix B). (c) *Non-linear similarity (t-SNE)*. Beyond linear similarity (CKA $< 10^{-4}$ off-diagonal, Figure 16), a non-linear t-SNE embedding (Figure 19) also separates the four conditions with 1-NN purity 1.000 vs. chance 0.20 — the format divergence is not an artefact of linear-similarity measures. H1 is therefore specifically about linear top-layer readability — what the linear DD-PPO action head can consume — not GPS absence.




Figure 10: **Magnitude: corroborating diagnostics (referenced from §3.1 and Table 1).** (a) Top-layer GPS R^2 on Gibson and held-out MP3D (5-fold CV, deterministic rollouts, 300 ep/cond; lighter bars = MP3D; the MP3D shift is single-seed and awaits replication). (b) GPS R^2 binned by step-in-episode (7 log-spaced bins): rich-encoder readability holds in the mid-episode window then collapses in the long tail. (c) Pipeline view: GPS R^2 along Encoder $\rightarrow$ L0 (LSTM input) $\rightarrow$ L1 $\rightarrow$ L2 (single-fit; the 5-fold CV version is condensed in main Figure 2). Readability is comparable across conditions at the LSTM input and only diverges at L2; the shaded band at L2 is the 2-layer MLP probe recovery range. (d) Per-unit Skaggs spatial information [Skaggs et al., 1996] is condition-flat at the headline metric; the per-unit regression spatial analysis in §3.1 (trajectory-phase vs. raw-position R^2) is the more diagnostic reading. (e) Mid-episode excursion forgetting (25-step random-action perturbation; per-segment scale-invariant MAE): blind fragments most, consistent with the strongest reliance on the integrated action stream.

D Magnitude: corroborating diagnostics

The main-text Figure 2 (§3.1) summarises the magnitude story across three views: cross-condition R^2 , training-trajectory dynamics, and predictive horizon. This appendix collects the corroborating diagnostics referenced from §3.1: the composite of cross-environment generalisation / step-binned readability / pipeline view / Skaggs information / excursion forgetting (Figure 10); the per-LSTM-layer breakdown of where the GPS code lives (Figure 11); a single-fit companion check (Figure 12); the GPS+compass version of the substitution-dynamics curves (Figure 13); and the nat-scale information-theoretic capacity estimate (Figure 14, §D.1).

D.1 Mutual-information capacity accounting (MINE)

To anchor the information-allocation finding in §3.1 on a nat-scale rather than the bounded, model-dependent R^2 scale of probes, we estimate $I(\mathbf{h}_2; \text{pos})$ per condition using the Mutual Information Neural Estimator [Belghazi et al., 2018]. MINE exploits the Donsker–Varadhan dual representation of the KL divergence:

$$I(X; Z) \geq \sup_T \mathbb{E}_{P_{XZ}}[T(x, z)] - \log \mathbb{E}_{P_X \otimes P_Z}[\exp T(x, z)],$$

with $T_\theta : \mathcal{X} \times \mathcal{Z} \rightarrow \mathbb{R}$ a neural network optimised by gradient ascent on the right-hand side. Joint samples (x_i, z_i) come from paired $(\mathbf{h}_2, \text{pos})$ along episodes; marginal samples are obtained by within-batch shuffling of z . We use a 3-layer MLP with 256 hidden units, ELU activations, batch size 512, learning rate 10^{-4} , 4,000 training steps per condition, and the bias-corrected gradient (eq. 12 of Belghazi et al. 2018) to mitigate the SGD bias of the log-denominator estimate.





Figure 11: **Where GPS lives along the LSTM stack: per-layer linear probe R^2 .** 5-fold CV linear probe at each hidden state h_ℓ and cell state c_ℓ , $\ell \in \{0, 1, 2\}$, for all 4 conditions. *Left:* GPS. *Right:* compass. Three patterns. (i) At h_0 all conditions decode GPS at $R^2 \sim 0.95$ — the GPS sensor embedding enters the LSTM concat at Layer 0 regardless of vision condition (Appendix B). (ii) h_1 dips for every condition including blind ($R^2 = +0.22$); information passes through the middle layer with reduced linear decodability for everyone, then partly recovers at h_2 for the bottleneck conditions. (iii) Cell states c_ℓ do not carry the GPS code as cleanly as h_ℓ : c_1 is at chance/negative for every condition; c_2 is positive only for blind ($R^2 = +0.57$). The headline “top-layer linearly-decodable GPS” result in §3.1 is therefore specifically about h_2 (the policy readout), not the cell-state pathway. Same single-seed limitation as the main paper.



(a) Per-LSTM-layer GPS / compass probe R^2 across all three layers, all four conditions (single train/test fit; the 5-fold CV mean reported in Table 1 matches qualitatively but with different magnitudes for rich-encoder conditions). Compass tells the same story as GPS at smaller magnitude.

(b) Companion to Figure 10a (which uses 5-fold CV mean): single train/test fit on the same checkpoints; the qualitative bottleneck-vs-rich-encoder partition is the same.

Figure 12: **Single-fit per-layer and MP3D companion checks.** Companion to the headline 5-fold CV results: a single train/test split reproduces the bottleneck-vs-rich-encoder partition reported in Table 1, ruling out CV-fold-induced artefacts. The MP3D panel re-evaluates the same checkpoints on a held-out scene distribution; the partition still holds.

E Spatial Format: corroborating diagnostics

The main-text Figure 3 (§3.2) shows three convergent measures of format divergence. This appendix collects: the full transplant matrix together with cross-condition probe transfer and subspace evolution (Figure 15); a linear-similarity null and a non-linear similarity check (Figures 16, 19); the geometric structure of the position code within h_2 (Section E.1); a population-coding analysis (Figure 17); the eigenspectrum power-law per condition with associated geometric scalars (Section E.3); and a causal subspace ablation (Section E.2).

E.1 Capacity-allocation: additional geometric tests

The figures below provide independent geometric tests of the capacity-allocation principle, beyond the headline information-allocation (Fig. 3a) and LOSO scene-invariance (Fig. 3b) results in §3.1. (The cross-condition probe-transfer matrix is shown as panel (b) of Figure 15; we do not duplicate it here.)




Figure 13: **Cross-training substitution dynamics: GPS and compass (companion to main Figure 2b).** Probe R^2 across training checkpoints, truncated at 250M frames. (a) Top-layer GPS R^2 (the source panel for Figure 2b). (b) Top-layer compass R^2 tells the same story at smaller magnitude — substitution applies to spatial codes broadly, not just position.




Figure 14: **MINE capacity accounting.** (a) MINE training curves: $I_\theta(\mathbf{h}_2; \text{pos})$ in nats vs. optimisation step, per condition. All curves converge by $\sim 3,000$ steps. (b) Final $I(\mathbf{h}_2; \text{pos})$ in bits per condition (mean $\pm$ std across 3 MINE seeds): Blind 6.01 ± 0.06 , Coarse 4.55 ± 0.03 , Foveated 4.30 ± 0.04 , Foveated-logpolar 4.66 ± 0.04 , Uniform 4.47 ± 0.06 . Total mutual information decreases with encoder bandwidth, but the decrease is gentle ($\approx 1.3 \times$ blind/uniform) compared to the linear-probe R^2 collapse (effectively factor-of-zero for uniform). The within-sighted ordering (foveated $<$ uniform $<$ coarse $<$ foveated-logpolar) is reproducible across MINE seeds (per-cond std ≤ 0.06 bits, much smaller than the within-sighted gap ≈ 0.36); the dominant signal is the bottleneck-vs-rich-encoder gap (Blind 6.01 vs sighted mean 4.49). The interpretation: rich-encoder conditions retain substantial position information in $\mathbf{h}_2$, but it is encoded non-linearly and partly offloaded to the encoder route, paralleling the LOSO scene-invariance result (Figure 3b).

E.2 Single-direction β -scrubbing reveals redundant population coding

The Ridge probe identifies a 2-d direction β in $\mathbf{h}_2$ that linearly predicts (x, y) . We test whether this direction is *causally responsible* for the recoverable position code by projecting $\mathbf{h}_2$ onto the null-space of β (rank-510 subspace) and re-evaluating both Ridge and MLP probes (Heimersheim & Nanda 2024 noising-style patching: tests whether β was *necessary* for the encoded behavior [Heimersheim and Nanda, 2024]).




Figure 15: **Spatial Format: corroborating diagnostics (referenced from §3.2).** (a) The full 5×5 memory-transplant matrix at midpoint $t=30$ (the per-recipient mean enters main Figure 3b as the “transplant cost” axis); the matrix shows the same direction-asymmetry condensed there. (b) Cross-condition probe transfer matrix (train on row, test on column): off-diagonal R^2 is catastrophically negative (-3.5×10^2 to -3.5×10^4 unclipped; clipped to $[-1, 1]$ for visualisation) — the position-axes encoded across conditions are jointly incompatible at the readout level, even when the same scenes are visited and the same task is solved. (c) Subspace divergence is established early: principal angles are already near-orthogonal at the earliest probed checkpoint (~ 50 M frames) and stay so throughout, indicating the structural separation is laid down well before convergence.

	Linear orig	Linear scrub	MLP orig	MLP scrub
Blind	+0.95	+0.93	+0.94	+0.95
Coarse	+0.72	+0.62	+0.79	+0.84
Uniform	−1.08	−1.46	+0.42	+0.31

The MLP probe is essentially unaffected by scrubbing across all three conditions ($\Delta R^2 \in [-0.05, +0.11]$). The Ridge probe shows small drops for Blind (−0.02) and Coarse (−0.10), and a more-negative shift for Uniform (the absolute value increases due to the variance of negative




Figure 16: Unaligned linear CKA [Kornblith et al., 2019] between top-layer LSTM hidden states ($n=95,640$ steps per condition under deterministic-rollout probes); every off-diagonal is below 10^{-4} . Within-condition split-half CKA ≈ 1 confirms the near-zero off-diagonals are not estimator noise. Pairs that look identical here differ markedly under transplant and probe-transfer (Figure 15); CKA is the least informative of the three format-divergence measures because it cannot read the asymmetric structure that transplant and probe-transfer show.




Figure 17: **Population-coding analysis: per-unit spatial information vs. population readout (referenced from §3.2).** (a) Per-unit spatial-information distribution: each curve is one condition’s units sorted by spatial information (descending), x-axis log scale. Counter-intuitively, rich-encoder conditions (uniform, foveated) have a few more peaked spatially-tuned units (max ≈ 2 bits) than bottleneck conditions; blind has the flattest distribution (max ≈ 1 bit, no unit reaches the dashed 1-bit threshold). (b) Sparse-vs-distributed GPS decoding: probe R^2 when only the top- k spatial-info units are used as features. Blind reaches $R^2 \approx 1.0$ at $k = 512$; coarse climbs to $R^2 \approx 0.5$; rich-encoder networks fail to decode GPS at any k even though their top units carry higher per-unit spatial information. The combined picture: *higher per-unit spatial information $\neq$ position readout*. Rich-encoder peaked units appear to encode features correlated with position (e.g. visual landmarks, trajectory phase) rather than position itself.

R^2 estimates). Two interpretations: (i) the rank-2 subspace identified by Ridge is one of many redundant directions encoding position — when removed, Ridge re-fits to a different direction in the remaining 510 dims; (ii) MLP, having access to non-linear combinations of remaining dims, trivially recovers the position code from this redundant population. This is the population-coding-redundancy property familiar from biological place-cell systems [Stringer et al., 2019]: single-cell or low-rank ablation does not break the code because of high-dimensional distributed representation. Methodologically, single-direction scrubbing is too weak to causally localise position information in h_2 ; higher-rank ablation, attribution-based pruning, or path-patching [Heimersheim and Nanda, 2024] would be next steps. Foveated is deferred to the 250M re-probe.




Figure 18: **Goal-vector probe R^2** (referenced from Appendix C, “probe-capacity rebuttal”). Goal distance (DtG) decodes strongly across all conditions ($R^2 \in [0.81, 0.90]$); goal direction and the full ego-frame goal vector are at chance or negative everywhere. PointGoal provides the goal vector as a per-step input, so the policy has no incentive to re-encode it.




Figure 19: **Non-linear similarity: t-SNE of top-layer hidden states** (referenced from §3.2). *Left*: joint t-SNE of 5,000 samples (1,000 per condition) in 512-d hidden space (perplexity 30). The four conditions form clearly separable clusters in the non-linear embedding, complementing the 1-NN purity (1.000 vs. chance 0.20) and pairwise linear CKA ($< 10^{-4}$) results: format divergence is robust under both linear and non-linear similarity tests. *Right*: per-condition t-SNE coloured by distance-to-goal ($n=2,000$ per panel). Within-condition structure varies: blind shows a clearer DtG gradient than the rich-encoder conditions, consistent with §3.1: blind’s hidden state carries position information more directly while rich-encoder conditions encode it less linearly.

E.3 Eigenspectrum power-law per condition

We adapt the eigenspectrum analysis of Stringer et al. [2019] to the LSTM hidden state, fitting a power law $\lambda_n \propto n^{-\alpha}$ to the eigenvalues of $\mathbf{h}_2$ per condition (PCs $n \in [5, 200]$ to skip finite-sample bias and machine-precision tail). All five conditions show a tight power-law eigenspectrum (Figure 23, $R^2 \geq 0.99$). The fitted exponent shows a sharp bottleneck-vs-rich-encoder split:

Condition	α	R^2	Var. top 10 PCs
Blind	3.85 ± 0.03	0.99	95%
Coarse	1.80 ± 0.01	1.00	82%
Foveated	1.72 ± 0.01	1.00	76%
Foveated-logpolar	1.80 ± 0.01	1.00	95%
Uniform	1.90 ± 0.01	1.00	91%




Figure 20: **Position-axis lives in low-variance directions of h_2 .** Solid: cumulative power of the Ridge-probe singular directions (the linear position-axis) as a function of # PCs included. Dashed: cumulative explained variance (the manifold’s own variance distribution). For all four conditions, the manifold’s variance is concentrated in $\sim 7\text{--}9$ top PCs (90% by PC 9), but the position-axis distributes its power across 200+ PCs — it does not reach 50% within the first 200 PCs. Position information is encoded in the *low-variance tail* of h_2 across all conditions; the discriminator across conditions is therefore not where (in a high-variance sense) position lives but rather the cross-scene stability of the readout direction (Figure 3b).




Figure 21: **Linear readout vs. # PCs.** (a) Ridge-probe GPS R^2 as a function of the number of top PCs of h_2 used as the probe input. Per-cond pattern: blind plateaus quickly even with ≤ 10 PCs (linear axis aligned with low-variance dirs); rich-encoder conditions barely climb across 50 PCs (information distributed). (b) Participation ratio $(\Sigma\lambda)^2/\Sigma\lambda^2$ (number of effectively active dimensions, ignoring tail) across conds. Variance distribution is similarly compact across conditions; what differs is how position information is laid out within these dimensions, supporting the position-axis analysis (Figure 20).

Two observations. (i) Blind sits well above the critical exponent $1 + 2/d_{\text{task}} \approx 1.67$ that demarcates differentiable from fractal manifolds for a $d_{\text{task}} = 3$ stimulus (current position + heading), and far above Stringer’s V1 value ($\alpha \approx 1.04$). With $\alpha = 3.85$, blind’s eigenspectrum decays sharply: variance is concentrated in the top few PCs (top-10 already capture 95%). We do not claim that Stringer’s specific theoretical bound ($\alpha = 1 + 2/d$, derived for smooth tuning curves over a d -dim stimulus manifold) directly applies to our recurrent setting; the observation is descriptive. (ii) The four sighted conditions cluster tightly around $\alpha \approx 1.7\text{--}1.9$, all near the $1 + 2/d_{\text{task}} \approx 1.67$ critical exponent — the smoothness/efficiency border. Blind is the outlier above, in a few-dim con-




Figure 22: β -subspace scrubbing. For each cond (foveated deferred), Ridge β identifies a 2-d direction in h_2 that linearly predicts (x, y) ; scrubbed $h_2 = (I - \beta^+ \beta) h_2$ removes this 2-d subspace. MLP probe is essentially unaffected by scrubbing across all three conditions: position information is redundantly distributed across many h_2 dimensions, not concentrated in the 2-d β subspace.

Table 2: **Geometric scalars: per-condition manifold shape (companion to Figure 23; referenced from §3.2).** Participation ratio (PR), intrinsic dimension (ID-MLE / ID-CDF), and eigenspectrum power-law exponent α on h_2 . All values from post-retrain hp-consistent agents (3-seed mean over 30,000-step subsamples; std < 0.05). Blind separates on α (sharper tail) and on ID (lowest); the four sighted conditions cluster on α and ID (1.73–1.92, 2.04–2.41). PR varies more across sighted conditions and does not track encoder bandwidth monotonically; we read it as different geometric distributions of the visual route’s residual variance rather than a clean ordering.

Condition	PR	ID-MLE	ID-CDF	α
Blind	4.41	1.72	1.54	3.63
Coarse	5.25	2.41	2.20	1.81
Foveated-logpolar	1.70	2.25	2.04	1.88
Foveated	9.11	2.41	2.30	1.73
Uniform	3.09	2.32	2.17	1.92

centrated regime, while the four sighted conditions distribute variance across more PCs. The split bottleneck-vs-rich-encoder reading: blind’s heavy reliance on the integrated GPS code crystallises into a low-dimensional eigenspectrum, while sighted conditions (including coarse, whose 1×1 encoder still feeds a ResNet-18 stack with diverse channels) distribute hidden-state variance across more directions to support encoder-route information.

F Temporal Format: corroborating diagnostics

The main-text §3.3 reports the temporal-generalisation matrix and the phase-portrait + intrinsic-timescale corroboration. This appendix collects two further diagnostics: a predictive-horizon test grounded in successor-representation theory (Section F.2), and a topological-distinguishability test via persistent homology (Section F.1, referenced from §3.3).

F.1 Topological distinguishability via persistent homology

A parameter-free pipeline used for toroidal grid-cell structure in mouse cortex [Gardner et al., 2022] confirms that conditions are topologically distinguishable: cross-condition Wasserstein-2 distance between H1 persistence diagrams is roughly twice the within-condition seed noise (Figure 24b). Persistent-loop counts order monotonically with encoder bandwidth (blind highest, uniform lowest; Figure 24a). We treat this as suggestive corroboration rather than independent mechanistic evidence. Our pre-registered prediction was the opposite direction (we expected uniform highest, blind lowest); we report what we found.




Figure 23: **Eigenspectrum power-law per condition.** (a) Eigenvalue λ_n vs. PC index n for $\mathbf{h}_2$, log-log; reference power laws $\alpha = 1$ (Stringer V1) and $\alpha = 1 + 2/3$ (theoretical critical exponent for $d_{\text{task}} = 3$) shown for comparison. (b) Fitted power-law exponent α per condition (linear fit on $\log-\lambda_n$ vs. $\log-n$, PCs 5–200); error bars are the std of the slope estimate. The monotonic ordering is consistent with the capacity-allocation principle: rich-encoder conditions distribute variance across more PCs and obtain lower α than bottleneck conditions.




Figure 24: **Topological corroboration of the temporal-format dichotomy.** (a) Persistent-loop counts (number of H_1 features with persistence > 0.1) per condition. (b) Wasserstein-2 distance between H_1 persistence diagrams: the cross-condition gap exceeds within-condition (seed) noise across all five conditions, so the $\mathbf{h}_t$ manifolds are topologically distinguishable.

F.2 Predictive horizon: does $\mathbf{h}_t$ encode future positions?

Successor-representation theory [Stachenfeld et al., 2017, Whittington et al., 2020] predicts that hippocampal place cells encode not just current position but a discounted future-occupancy code: place-cell firing at s_t corresponds to expected visits to nearby states s_{t+k} . We test the analog claim for our recurrent state: does a linear probe from $\mathbf{h}_t$ recover *future* positions $(x, y)_{t+k}$ for $k \in \{0, 1, 5, 10, 20, 50\}$?

	$k = 0$	$k = 1$	$k = 5$	$k = 10$	$k = 20$	$k = 50$
<i>Linear (Ridge $\alpha = 10$) R^2, all-episode pooling, episode-level 5-fold CV</i>						
Blind	+0.94	+0.94	+0.94	+0.92	+0.90	+0.79
Coarse	+0.57	+0.57	+0.57	+0.56	+0.50	+0.37
Foveated	+0.06	+0.06	+0.08	+0.13	+0.19	+0.06
Foveated-logpolar	+0.17	+0.18	+0.17	+0.11	−0.14	−0.50
Uniform	−4.31	−4.32	−4.32	−4.26	−4.39	−10.14

Three observations. (i) *Blind exhibits a long predictive horizon.* Linear R^2 from $\mathbf{h}_t$ to $(x, y)_{t+k}$ remains ≥ 0.90 from $k = 0$ to $k = 20$ (only 0.04-point drop), and 0.79 at $k = 50$. The same linear projection that decodes *current* position from $\mathbf{h}_t$ at $R^2 = 0.94$ also decodes positions 20 steps in the future at $R^2 = 0.90$. This is a strong form of the predictive-map signature predicted by SR theory: $\mathbf{h}_t$ encodes future-state occupancy, not just current state. (ii) *Coarse follows the same pattern at lower magnitude* (linear $R^2 \in [0.50, 0.57]$ across $k = 0$ –20, dropping to +0.37 at $k = 50$). (iii) *Rich-encoder conditions have no clear predictive horizon.* Foveated stays in $[0.06, 0.19]$ across $k = 0$ –50 with no monotonic decay — the small positive values reflect probe-level noise rather than a sustained future-state code. Foveated-logpolar starts at +0.17 but turns negative by $k = 20$. Uniform is already strongly negative at $k = 0$ ($R^2 = -4.31$) and grows more negative. Rich-encoder $\mathbf{h}_t$ encodes neither current nor future position linearly — consistent with the LOSO scene-invariance reading in §3.2, but adding the temporal aspect: even with non-linear flexibility, future positions are not encoded. The signature is asymmetric: bottleneck conditions encode $\mathbf{h}_t$ as a smoothly-varying integrated GPS code that linearly predicts upcoming positions; rich-encoder conditions encode $\mathbf{h}_t$ reactively to the current visual scene without temporal extrapolation.

G Consumption: corroborating diagnostics

The main-text §3.4 reports two complementary causal-intervention tests: the persistent-memory failure-mode test (which discriminates conditions) and the memory-init transplant + GPS-sensor ablation boundary checks (which characterise the consumption mechanism without separating conditions). This appendix shows: the canonical paired-episode trajectory per condition (Figure 25); the full 4×4 failure-mode catalog spanning the per-condition margin distribution (Figure 26); and a memory-transplant midpoint sweep that defends the $t=30$ choice in the main-text 5×5 matrix (Figure 27).

G.1 Persistent-memory failure-mode catalog

The main-text Figure 25 shows one canonical paired-episode case per condition; the full 4×4 catalog below (Figure 26) shows four cases per condition spanning the per-condition margin range (most-tries-new $\rightarrow$ most-locks-onto-old). All cases share the selection criterion: reset SPL > 0.5 , persistent SPL < 0.2 , persistent steps ≥ 30 , $|\Delta y_{\text{goal}}| < 0.5$ m (same-floor old vs. new goal). Within each row, panels are ordered by margin (column 1 = most-negative = closer to NEW goal direction; column 4 = most-positive = closer to OLD goal direction).

G.2 Memory-transplant midpoint sweep: defending the $t=30$ choice

The main-text 5×5 memory-transplant matrix is reported at a fixed midpoint $t=30$ (the full matrix appears as panel (a) of Figure 15). A natural concern is whether the asymmetry pattern is specific to that midpoint. The midpoint sweep below addresses this: across midpoints $\{0, 15, 30, 60, 90, 200, 400, 800\}$ time-steps, the bottleneck-vs-rich-encoder gap appears for $t \in$




Figure 25: **Persistent-memory failure modes (referenced from §3.4): canonical paired-episode case per condition.** Extended catalog: Figure 26. Aggregate per-condition margin = avg dist(persistent end $\rightarrow$ old goal) – avg dist(persistent end $\rightarrow$ new goal) over same-floor failures (reset SPL > 0.5, persistent SPL < 0.2, persistent steps ≥ 30 , $|\Delta y_{\text{goal}}| < 0.5\text{m}$): Blind -0.38m ($n=27$), Coarse -0.57m ($n=35$), Foveated -0.59m ($n=16$), **Uniform** $+1.83\text{m}$ ($n=46$, **closer to OLD goal**). Single seed.

Table 3: Foveation experiment status (submission time).

Experiment	Tests	Status
F1–F4 strength sweep ($\sigma_{\text{max}} \in \{2, 4, 12, 20\}$)	encoder–memory race as continuous lever	future work
F3 log-polar foveation	spatial-sampling vs. blur bottleneck	in training (RCP probe-4, unified hp)
Encoder feature-map probe	where position information is held	landed; see §4

[30, 90] and decays toward zero by $t \geq 400$, with a protocol-noise floor at $t=0$ where the donor’s hidden state has had no time to integrate.

H Foveation experiments: status and design

This appendix details the four controlled foveation experiments referenced from the main paper. All four are submitted.

Foveation strength sweep (F1–F4) — planned for follow-up, not in this submission. The protocol would train four additional foveated agents at $\sigma_{\text{max}} \in \{2, 4, 12, 20\}$ alongside the existing $\sigma_{\text{max}} = 8$, holding gaze location, encoder stack, training budget, and training pool fixed. Expected signature: at low σ_{max} the foveated agent should look like uniform; at high σ_{max} it should approach coarse. The R^2 vs. σ_{max} curve would directly test the encoder–memory race as a continuous lever rather than a binary one.

Spatial sampling vs. blur (F3 log-polar). Gaussian blur removes high-frequency content but preserves spatial layout: a $\sigma_{\text{max}} = 8$ foveated image still feeds the encoder 256×256 pixels, so the ResNet-18 stack still produces an 4×4 spatial feature map. A log-polar foveation [Strasburger et al., 2011, Rosenholtz, 2016] resamples with a non-uniform retinal grid (e.g. 64×64): the encoder spatial output drops to $\sim 2 \times 2$, much closer to coarse’s 1×1 than to uniform’s 4×4 . This isolates the variable our sharpened H1 mechanism identifies as the trigger: *spatial-feature variety in the encoder output*. The preregistered prediction (written before result was available) was that log-polar should produce LSTM GPS $R^2 \geq 0.3$, between coarse’s $+0.58$ and uniform’s ≈ 0 , with corresponding behavioural shifts; if log-polar matched uniform, the encoder-spatial-output-dimensionality mechanism would be wrong. The result (next paragraph) falsifies that simple reading.

Why disclose this rather than wait. Each experiment has a falsifiable signature; we disclose the designs openly so (i) the submission is honest about what foveation-specific evidence already exists vs. what is left for follow-up, and (ii) the reader can identify which incoming result would change the paper’s interpretation.




Figure 26: **Persistent-memory failure-mode catalog: 4 cases per condition spanning the per-condition margin distribution.** Each panel: reset trajectory (solid, condition colour) reaches the new goal; persistent trajectory (dashed) demonstrates the failure. The cross-condition pattern: bottleneck conditions (Blind, Coarse) cluster at “tries-new” (most cases attempt new goal direction but fail to reach); Uniform clusters at “locks-onto-old” (most cases stay near previous-episode goal); Foveated is mixed with no dominant mode. Single seed per condition (cogsci modelling convention; see §4.4).

I Encoder-capacity scaling: log-polar foveation control

A direct test of the H1 mechanism (encoder spatial-feature variety per step drives top-layer LSTM spatial encoding) is to vary that variable continuously across our existing 4-condition spectrum. The most diagnostic single point is the *log-polar foveation* control: log-polar resampling onto a 64×64 retinal grid produces an encoder spatial output of $\sim 2 \times 2$, intermediate between coarse’s 1×1 and uniform’s 4×4 (verified by direct forward pass through the trained encoder). Only the visual front-end differs from the foveated condition; LSTM, sensors, and training pipeline are identical. **Preregistered prediction:** if the H1 mechanism is the simple “encoder spatial-output dimensionality” reading, log-polar should produce LSTM GPS R^2 in the bottleneck regime (i.e. closer to coarse’s $+0.58$ than to uniform’s ≈ 0); if log-polar matches uniform on H1 instead ($R^2 \approx 0$), the encoder-spatial-output mechanism story needs reframing toward channel information, encoder spatial-feature *variety* per step, or another aspect. **Result (converged 250M frames).** The log-polar agent at convergence (ckpt-49, ~ 245 M frames) yields top-layer GPS $R^2 = -0.029 \pm 0.759$ (5-fold CV, $n = 500$ episodes) — *not* in the bottleneck regime, but indistinguishable from foveated




Figure 27: **Memory-transplant midpoint sweep (referenced from §3.4).** Two representative donor→recipient pairs (foveated→uniform, foveated→blind) across midpoints $\{0, 15, 30, 60, 90, 200, 400, 800\}$ time-steps ($n=100\text{--}150$ episodes/cell). The full 5×5 matrix in main Figure 3 is the $t=30$ slice. (a): full SPL per condition — baseline / self / cross. The $t=0$ self/cross collapse to a protocol-noise floor (the diagonal in the main-text matrix is correctly read as rollout-divergence noise, not condition mismatch); SPL decays for all conditions at $t \geq 200$ because by then most successful episodes have already terminated, leaving only long atypical episodes. (b): cross-minus-self isolates condition-mismatch alone. The gap peaks in the $t \in [30, 90]$ window — where the agent has integrated history but is not yet committed to a trajectory — and decays toward 0 at $t \geq 400$. The asymmetry pattern reported at $t=30$ is therefore neither a $t=30$ -specific artefact nor a late-midpoint inflation.

($R^2 = +0.178 \pm 0.659$) and clearly far from coarse ($R^2 = +0.582 \pm 0.099$). The simple “encoder spatial-output dimensionality” framing is therefore *not* sufficient to explain the H1 phenomenon: log-polar’s 2×2 output does not place it in the bottleneck regime. We read the result as supporting the alternative “encoder spatial-feature *variety* per step” framing: log-polar’s radial sampling preserves rich gaze-centred feature variety per timestep even at low spatial output, so the policy can still source position from the visual route and the LSTM does not commit capacity to the integrated code. Substitution-dynamics analysis (Figure 2b) corroborates this reading: log-polar starts at $R^2 = +0.92$ at ~ 50 M frames (matching all 4 main conditions’ early high R^2), then decays along the rich-encoder trajectory rather than plateauing in the bottleneck range. A finer multi-point input-resolution sweep ($K \in \{32, 64, 96, 128, 192\}$) is reported as future work; its training cost (multiple from-scratch DD-PPO runs at ~ 250 M frames each) was outside this submission’s compute window.