
How Foveated Vision Shapes Cognitive Maps in Navigation Agents

Anonymous authors

Paper under double-blind review

Abstract

What determines the content of a recurrent navigation agent’s learned internal state—the task, the architecture, or the structure of the visual input? Prior work shows that blind PointNav agents develop cognitive-map-like representations in their recurrent state as a byproduct of task pressure [Wijmans et al., 2023]. We generalise this observation. Training DD-PPO agents in Habitat under five visual conditions that differ only in the *information capacity of the visual encoder* (blind, a 1×1 -collapsed encoder, foveated, uniform, and foveated with a learned-gaze module), we show that LSTM spatial encoding is driven by encoder bottleneck rather than by the presence or absence of vision specifically. World-frame GPS and compass decode robustly from the LSTM when the visual encoder cannot resolve them ($R^2 = 0.95, 0.78$ for blind, matched-compute) and decode at chance when the encoder can ($R^2 \in [-2.4, +0.06]$ for full-RGB conditions). The GPS code in bottleneck conditions is persistent across 10+ time-steps of lag, the integration signature of a cognitive map. Representational-format divergence (near-zero pairwise CKA, probe-transfer failure, cross-condition memory-transplant baseline SPL drops of 0.19–0.21; isolation metric up to 0.17) persists across all five conditions: sensor structure shapes not only whether spatial memory arises, but its linear-subspace geometry. Whether gaze *location* within the rich-encoder regime acts as a second representational-content axis is tested by a foveated-shifted causal control (gaze hardcoded to a different static image location, same foveation transform) currently in training; results will populate §4.4. Preliminary MP3D probes show the same H1 ordering across scene sets (deterministic re-collection in progress; §4.5); the findings have direct analogs in the animal-navigation literature (enhanced hippocampal coupling in blind humans; grid-cell persistence in darkness). The encoder-bottleneck axis gives a quantitative prediction testable by sweeping encoder capacity, which we begin here.

1 Introduction

Recurrent navigation agents trained with deep RL build internal states that support memory, planning, and generalisation, and the content of those internal states is at least as interesting as the task metrics they drive. A natural question is what *determines* that content: architecture, objective, training budget, or visual input? Wijmans et al. [Wijmans et al., 2023] showed that map-like spatial representations emerge from task pressure alone, even in blind agents that receive only goal displacement and proprioception; Dwivedi et al. [Dwivedi et al., 2022] probed sighted agents for related concepts; Banino et al. [Banino et al., 2018] and Cueva & Wei [Cueva and Wei, 2018] demonstrated grid-cell-like structure in path-integration networks. Together these results establish that *some* spatial structure is learned without being designed in. They leave open the complementary question that we take up here: within agents that all achieve high task success, how much of *what* recurrent memory encodes is determined by the structure of the visual input rather than by the task? In biology the answer is clear — primate hippocampal cells code joint place-view-action state because primates sample the scene




Figure 1: The five visual conditions differ only in what reaches the agent’s encoder. (a) Blind: no visual input; only the non-visual sensor stack (goal-in-start-frame, GPS, compass, close-to-goal scalar) — the Wijmans [Wijmans et al., 2023] configuration. (b) Uniform: 256×256 full-resolution RGB. (c) Foveated (fixed gaze): eccentricity-dependent Gaussian blur centred on the image middle ($\sigma_{\max} = 8$, quadratic falloff). The two additional conditions not shown are *matched-compute* (a down-sampled uniform view with the same total pixel budget as foveated) and *foveated (learned gaze)*, which adds a learned (x, y) gaze offset to (c). (d) All conditions solve PointGoal navigation in the same Habitat Matterport3D / Gibson scenes; the top-down view is shown for reference only and is not available to the agent.

with a foveated sensor [Wirth et al., 2017, Rolls, 2024], and the same bat produces non-overlapping hippocampal populations in vision versus echolocation [Geva-Sagiv et al., 2016]. We investigate the analogous question in a controlled agent setting where the only free variables are the visual input structure and the gaze that directs it.

Our five-condition design isolates three degrees of freedom: sensor type (blind / uniform / foveated), information volume (uniform vs. matched-compute), and gaze location (fixed vs. learned). Every condition uses the same DD-PPO learner, the same Gibson-0+MP3D training pool, the same LSTM backbone, and the same non-visual sensor stack (Figure 1). We adopt the LSTM architecture from Wijmans et al. [Wijmans et al., 2023] deliberately: the emergent-map literature we build on is recurrent, and matching that architecture is what lets our findings be interpretable as statements about *the structure of the task-shaped memory* rather than about the recurrent-versus-transformer design choice. We make our claims specifically about recurrent PointNav agents, and discuss in §5 whether the logic plausibly extends to transformer-backed or non-navigation systems (an empirical question we do not settle here). Within this setup, two accounts make divergent predictions for what a foveated agent’s memory should do. A *passive buffer* stores what it currently perceives, so decodable content should fall when current input is degraded. A *compensatory* memory invests in internal state when current input is unreliable, so decodable content should persist over longer horizons than for an agent whose current input is always reliable. Which pattern the agent converges to, and whether gaze location can act as a separate representational-content lever even when task performance is held fixed, are the questions this paper asks empirically.

We test three hypotheses. **H1 (visual-encoder bottleneck $\rightarrow$ LSTM spatial compensation)**. Agents whose visual encoder cannot resolve world-frame position or heading on its own should develop these quantities in their recurrent state; agents with a rich visual encoder should not. Signature: current-state GPS/compass probe R^2 tracks visual-encoder information content monotonically across five conditions (blind, matched-compute at 1×1 feature map, uniform, foveated, foveated-with-learned-gaze), with a persistent-integration signature (decoding remains accurate at lags of 5–10 time steps). **H2 (sensor structure $\rightarrow$ format-level divergence)**. Different sensor conditions should produce

internal representations in distinct linear subspaces, not scaled versions of a common code. Signature: pairwise CKA near zero, probe weights fail to transfer across conditions (while within-condition probes succeed), and cross-condition memory transplants degrade task performance beyond same-condition self-transplants. **H3 (gaze location $\rightarrow$ orthogonal content lever)**. Where a foveated sensor is centred (*gaze location*) should act as a representational-content axis independent of the sensor transform. We test this with a foveated-shifted control in which gaze is hardcoded to a different image location from the default central gaze; same foveation transform, same training budget, different gaze target. The behavioural effect of shifting gaze location on the LSTM’s probe signature is what we report (§4.4). Our minimal end-to-end learned-gaze module (a deterministic sigmoid MLP with slow-gaze updates, §4.4 and Appendix B) collapsed to a static output; we use that collapsed location as the foveated-shifted control rather than as evidence about active gaze. Each hypothesis has a direct analog in the animal-navigation literature [Kupers and Ptito, 2014, Chen et al., 2016, Geva-Sagiv et al., 2016, Jutras et al., 2013, Najemnik and Geisler, 2005], which we return to in Discussion as evidence that the effects reflect general sensor–memory coupling principles rather than architectural artefacts.

Findings H1 is supported by a clean monotonic effect: LSTM world-frame GPS encoding tracks visual-encoder information bottleneck from $R^2 = 0.95 \pm 0.02$ (blind, no encoder) and $R^2 = 0.78 \pm 0.10$ (matched-compute, 1×1 encoder) down to chance ($R^2 \in [-2.4, +0.06]$ with CV $\sigma \in [0.7, 4.0]$) for the three full-resolution-encoder conditions (uniform, foveated, foveated-learned). H2 (representational-format divergence) is supported independently by near-zero cross-condition CKA and by mid-episode memory-transplant SPL drops of 0.19–0.21. H3 asks whether gaze location is a second content axis within the rich-encoder regime; a clean causal test (foveated-shifted control with gaze hardcoded to the learned-gaze module’s collapsed position) is in training at the time of submission and will populate §4.4.

2 Related Work

Cognitive maps in RL navigation agents Wijmans et al. [Wijmans et al., 2023] demonstrated that map-like representations emerge in the recurrent memories of DD-PPO agents even without visual input; Dwivedi et al. [Dwivedi et al., 2022] probed sighted agents for related spatial concepts; Banino et al. [Banino et al., 2018] and Cueva & Wei [Cueva and Wei, 2018] found grid-cell-like units in recurrent networks trained on path integration; Gornet & Thomson [Gornet and Thomson, 2024] extended this via visual predictive coding. Hennig et al. [Hennig et al., 2023] showed that RL under partial observability can induce belief-state-like representations without explicit probabilistic objectives. Ramakrishnan et al. [Ramakrishnan et al., 2025] examined spatial cognition in frontier models. All these studies use either absent or spatially uniform perception; none vary input quality spatially, and none test whether different sensors produce different representational formats.

Foveated vision and gaze policies in deep learning Foveated rendering [Strasburger et al., 2011, Rosenholtz, 2016] exploits eccentricity-dependent acuity falloff for efficiency. Learned glimpse policies appear in recurrent attention models [Mnih et al., 2014]; Atanov et al. [Atanov et al., 2024] showed that sensor layout changes downstream task performance. Pourrahimi & Bashivan [Pourrahimi and Bashivan, 2025] probed a foveated visual-search model for neuroscience-predicted features; Shang & Ryoo [Shang and Ryoo, 2023] jointly learned gaze and motor policies following the Sprague & Ballard [Sprague and Ballard, 2003] framework. These address visual search or task performance; none asks how sensor structure reshapes the learned spatial memory of a downstream navigation policy.

Probing neural representations Belinkov [Belinkov, 2022] surveys probing classifiers; Hewitt & Liang [Hewitt and Liang, 2019] introduced control tasks to distinguish genuine structure from probe memorisation; Kornblith et al. [Kornblith et al., 2019] introduced linear CKA for cross-model similarity. We adopt linear probing, the Hewitt–Liang selectivity metric, and unaligned linear CKA, and complement them with two behavioural interventions directly on the recurrent state (memory transplant and shortcut discovery) to validate the probe-based findings outside the linear-probe framework.

Biological precedent Our hypotheses and framing draw on three threads of animal-navigation research. (i) Primate hippocampal cells code joint place–view–action state [Wirth et al., 2017, Rolls, 2023]; Rolls [Rolls, 2024] argues that primate *spatial-view cells* and rodent *place cells*

reflect the different sampling statistics of foveated versus near-panoramic vision. (ii) The same bat in the same environment produces non-overlapping hippocampal populations when it switches between vision and echolocation [Geva-Sagiv et al., 2016]—format-level divergence across sensor modalities within a single animal. (iii) In primates, each saccade resets hippocampal 3–8 Hz rhythms and the reliability of the reset predicts subsequent memory [Jutras et al., 2013]; only overt (not covert) attention automatically writes to visual working memory [Tas et al., 2016]. Gaze targets themselves are selected in an information-gain-maximising fashion consistent with ideal-observer predictions [Najemnik and Geisler, 2005, Gottlieb et al., 2013, Hayhoe and Matthis, 2018]. We use these results only to frame our ML predictions and interpret our findings; our contribution is the controlled ablation in a learned agent that cross-species and cross-modality biological comparisons cannot cleanly supply.

3 Methods

3.1 Agent Architecture and Training

All agents are trained on PointGoal navigation in Habitat [Savva et al., 2019] using DD-PPO [Wijmans et al., 2020] on a merged pool of Gibson-0+ [Xia et al., 2018] (411 scenes) and Matterport3D train (61 scenes), totalling 472 training scenes. Evaluation uses 18 held-out MP3D test scenes (1008 episodes).

All agents share an identical 3-layer LSTM (512-d) backbone and the Wijmans non-visual sensor stack: goal-in-start-frame, GPS, compass, and a close-to-goal scalar, each projected to 32-d and concatenated with a 32-d previous-action embedding. All sighted conditions are trained for a target of 250 M environment frames; the blind baseline is trained to 342 M frames to allow slower convergence from proprioception alone.

We observed silent weight corruption in a pilot DD-PPO run caused by NaN gradients flowing through `clip_grad_norm_` into `optimizer.step()`; we mitigate this with a gradient-sanitisation patch to `PPO.before_step`, described in Appendix C.

3.2 Visual Conditions

Five conditions vary only the visual input (Figure 1):

1. *Blind*: no visual input; non-visual sensor stack only, replicating Wijmans et al. [Wijmans et al., 2023].
2. *Uniform*: full-resolution 256×256 RGB encoded by ResNet-18.
3. *Foveated (fixed)*: eccentricity-dependent Gaussian blur ($\sigma_{\max} = 8$, quadratic falloff) with gaze locked at image centre; same ResNet-18 encoder.
4. *Matched-compute*: 48×48 uniform RGB. At this resolution, ResNet-18’s 32-fold spatial downsampling collapses the feature map to 1×1 , discarding all spatial structure from the encoder’s output.
5. *Foveated (learned)*: same blur model, but gaze centre predicted by a lightweight MLP trained end-to-end with the navigation policy.

The matched-compute condition is central to H1: its 1×1 encoder feature map creates a clear bottleneck that mirrors the blind agent’s inability to resolve world-frame position from the visual stream, while still providing a visual observation. This separates the “no vision” account from the encoder-bottleneck account.

Implementation note The foveated conditions (fixed, shifted, learned) disable the running-mean-and-variance input normaliser used by uniform and matched-compute to avoid an in-place autograd conflict along the gaze-decoder gradient path. All reported results were obtained with this setting; the qualitative H1 ordering (pass-through for all three rich-encoder conditions) is robust to this difference.

3.3 Probing Methodology

After training, we freeze all weights and collect per-step LSTM hidden states (top-layer $\mathbf{h}$, 512-d) paired with ground-truth spatial labels from $n=500$ Gibson held-out validation episodes (the split not

used in training). Rollouts use *deterministic* (argmax) action selection, matching the eval protocol of our shortcut-discovery, memory-transplant, and standard habitat-baselines evaluators. Under this protocol all five conditions produce comparable rollout lengths (mean $\sim 100\text{--}400$ steps, dominated by policy competence on the task), so no cross-condition length-matching is needed. We train Ridge regression probes ($\alpha = 10$) with episode-level 5-fold cross-validation, reporting mean $\pm \sigma$ across folds, and verify probe specificity with Hewitt–Liang label-permutation selectivity [Hewitt and Liang, 2019].

Probed variables: GPS position (episodic), compass heading (episodic), distance-to-goal, lag- k path history (target value at $t - k$ decoded from $\mathbf{h}_t$, $k \in \{0, \dots, 20\}$), visited-region occupancy, full occupancy grid (20×20), per-unit spatial information (place-cell analysis), and the ego-frame goal vector $\mathbf{R}(-\theta)(\mathbf{g} - \mathbf{p})$ decomposed into scalar goal-distance and 2-D direction.

Sampling protocol An earlier draft of this work used stochastic (policy-distribution) action selection for probe rollouts. Under stochastic sampling, conditions with higher action-entropy (specifically the rich-encoder conditions foveated, uniform, matched, blind under their trained stochastic policies) sampled the STOP action early with probability ~ 0.25 per step, producing mean episode lengths of ~ 4 steps and target variance two orders of magnitude below the episodic range. The resulting probe R^2 values were trivially high across all conditions (any predictor fits a near-constant target). This confounded cross-condition comparisons, and we switched to the deterministic protocol used here and elsewhere in the probe/eval pipeline. All results in this paper use the deterministic protocol.

For foveated-fixed, all analyses use checkpoint ckpt.36 at $\sim 174\text{M}$ frames, the last clean intermediate before the silent-corruption window (Appendix C). A small number of our earliest probe tables were collected from a slightly earlier intermediate ($\sim 154\text{M}$) on the same pre-corruption training trajectory; the two are within the training curve’s noise band, and we are re-running the probes at ckpt.36 for full consistency.

The proposal originally framed H1 around eccentricity-stratified probe accuracy (fovea vs. periphery). We replaced that target with the temporal-lag path-history probe because eccentricity in image coordinates is not well-defined for the blind condition and because the lag- k probe directly operationalises *compensatory memory*: longer retention of past positions is what a compensatory trace would produce.

3.4 Cross-Condition Analysis

To test H2, we employ two complementary methods. *Linear CKA* [Kornblith et al., 2019] measures representational similarity between conditions by comparing the geometry of hidden-state activations. Values near 1 indicate identical geometry; near 0 indicates unrelated representations. *Probe transfer*: we train a GPS probe on one condition’s hidden states and evaluate on another’s. If spatial information is encoded in a shared linear subspace, cross-condition transfer should succeed. Catastrophic failure indicates fundamentally different encoding formats. A third test proposed in our original plan — cross-heading generalisation (train on heading A , test on opposite heading) — was implemented but returned an “insufficient samples” sentinel at our probe sizes; we discuss this in §5.5.

3.5 H3 Ablation: Gaze-Diversity Auxiliary Loss

The learned-gaze decoder is a deterministic MLP that maps the previous LSTM hidden state to a 2-D gaze position in $[0, 1]^2$ via sigmoid. Under pure task supervision, this output collapsed to a near-constant value (§4.4). To test whether H3 can be recovered by explicitly discouraging collapse, we add a per-batch auxiliary loss: $\mathcal{L}_{\text{div}} = \lambda \cdot \text{relu}(\sigma_{\text{target}}^2 - \text{Var}_{\text{batch}}[g])$. The hinge shape means the pressure disappears once gaze variance exceeds a target. Pilot failure modes are reported in Appendix B.

3.6 Behavioural Probes

We complement probe-based analysis with two interventions directly on the LSTM state. *Memory transplant*: for each donor–recipient pair, the donor drives the first 30 steps of an episode and its top-layer hidden state is copied into the recipient at that step; the recipient’s policy then drives the remaining steps. We compare *baseline* (recipient alone, full episode), *self-transplant* (recipient as its own donor, controls for transplant mechanics), and *cross-transplant* (different donor). Midpoint 30

Table 1: Per-condition summary. Ridge $\alpha = 10$ probes on hidden states collected under deterministic argmax-action rollouts ($n=500$ held-out Gibson episodes per condition). Probe R^2 columns report mean $\pm$ std across 5-fold episode-level CV. *GPS* and *Compass* are world-frame; *DtG* is ego-relative distance-to-goal. Lag-5 GPS R^2 reports stability of the GPS code across 5 time-steps (persistent-integration signature). Bold entries mark the top within the H1-relevant ordering (blind $>$ matched $\gg$ others).

Condition	Training			Current-state probes (R^2)			H1	H3	r
	Frames	SPL	Succ.	GPS	Compass	DtG	lag-5 GPS	goal-dir R^2	
Blind	342M	0.59	0.93	+0.95$\pm$0.02	+0.81$\pm$0.08	+0.90 $\pm$ 0.03	+0.95$\pm$0.02	−0.06	
Matched-compute	250M	0.67	0.98	+0.78$\pm$0.10	+0.64$\pm$0.10	+0.85 $\pm$ 0.12	+0.78$\pm$0.10	−0.03	
Uniform	250M	0.85	0.99	−0.31 $\pm$ 0.86	+0.36 $\pm$ 0.23	+0.86 $\pm$ 0.09	−0.27 $\pm$ 0.85	−0.03	
Foveated (fix)	174M	[TODO: 0.83]	[TODO: 0.98]	+0.06 $\pm$ 0.88	+0.07 $\pm$ 0.69	+0.82 $\pm$ 0.09	+0.04 $\pm$ 0.93	−0.08	
Fov-learned	250M	0.82	0.98	−2.43 $\pm$ 3.98	−1.34 $\pm$ 3.14	+0.81 $\pm$ 0.09	−3.60 $\pm$ 6.20	[TODO: H3]	

was chosen because Gibson PointNav episodes average ~ 100 – 120 steps; a later midpoint left too few steps after transplant. We use 150 episodes per pair per condition, sampled uniformly from the held-out Gibson validation split. *Shortcut discovery*: for each condition, we run episodes in pairs of two using the same scene but different goals, sampling 20 scenes with 10 paired episodes per scene (200 paired runs per condition per setting). Per-condition SPL is reported as the mean of the 20 scene-level mean SPLs (prevents high-episode-count scenes from dominating). We compare *reset* (LSTM zeroed between the two episodes) vs. *persistent* (LSTM state carried over) on the second-episode SPL. A reduction under the persistent setting measures how tightly the stored representation is locked to the previous goal. For the foveated-fixed condition, both behavioural probes use the ckpt.36 checkpoint (~ 174 M frames, the last clean intermediate before the silent-corruption window, §4); the linear-probing results in §4.2–4.5 use a slightly earlier ~ 154 M intermediate from the same pre-corruption window.

4 Results

4.1 Summary of Conditions and Key Metrics

Table 1 collects the per-condition numbers this section returns to. All sighted conditions solve PointGoal at 97–99% success; blind reaches 93% at a longer budget, reproducing Wijmans et al. [Wijmans et al., 2023]. Differences in what the hidden state encodes therefore cannot be attributed to differences in what the agents *do*. The probe machinery itself is calibrated: Hewitt–Liang label-permutation selectivity exceeds 0.76 for every sighted condition, and distance-to-goal decodes robustly across all five conditions ($R^2 \in [0.81, 0.90]$). The headline pattern previewed in the table is monotonic in visual-encoder information: conditions with no encoder (blind) or a 1×1 -collapsed encoder (matched) robustly encode world-frame GPS and compass in LSTM hidden state (GPS $R^2 = 0.95, 0.78$; compass $R^2 = 0.81, 0.64$); conditions with a full-resolution encoder (uniform, foveated, foveated-learned) do not (GPS $R^2 \in [-2.4, +0.06]$ with CV $\sigma \in [0.7, 4.0]$, i.e. chance). We develop this as H1. Two training-stability caveats are disclosed in full in Appendix C: foveated-fixed uses ckpt.36 at ~ 174 M frames (last clean checkpoint before a silent NaN-gradient corruption at ≈ 182 M), and matched-compute’s 1×1 feature-map collapse was a training quirk that turned out to be the mechanism, not a caveat.

4.2 H1: Visual Encoder Bottleneck $\rightarrow$ LSTM Spatial Compensation

Finding Visual-encoder capacity inversely determines LSTM spatial-encoding content. Blind agents encode current GPS in the recurrent state at $R^2 = 0.95 \pm 0.02$; matched-compute agents (48×48 RGB input, 1×1 encoder feature map) at $R^2 = 0.78 \pm 0.10$; all three full-resolution-encoder conditions (uniform, foveated, foveated-learned) decode GPS at chance ($R^2 \in [-2.4, +0.06]$ with CV $\sigma \in [0.7, 4.0]$). Compass shows the same monotonic ordering at a smaller effect size ($0.81 \rightarrow 0.64 \rightarrow 0.36 \rightarrow 0.07 \rightarrow -1.3$). DtG decodes robustly everywhere ($R^2 \in [0.81, 0.90]$). Figure 2 visualises the pattern.




Figure 2: Current-state probe R^2 (5-fold CV, episode-level splits, Ridge $\alpha = 10$ on deterministic-rollout hidden states) across five conditions. *GPS* (world-frame position) and *compass* (world-frame heading) show a clean monotonic ordering tracking visual-encoder information content: blind (no visual input) and matched-compute (48×48 RGB $\rightarrow 1 \times 1$ feature map) develop rich world-frame representations in their recurrent state; uniform, foveated, and foveated-learned (full 256×256 RGB) do not. *DtG* (distance-to-goal, an ego-relative task quantity) decodes robustly in *all* conditions because every agent must track it to navigate. Error bars: $\pm 1\sigma$ across 5 episode-level CV folds.

Integration signature at lag k The bottleneck conditions’ GPS encoding is *persistent*, not an instantaneous readout. Decoding GPS_{t-k} from $\mathbf{h}_t$ yields $R^2 = 0.95$ at $k=0, 2, 5, 10$ for blind and $R^2 = 0.78$ at every k up to 10 for matched. The rich-encoder conditions decode at chance at every k . Compass decays steadily with k in both bottleneck conditions ($0.81 \rightarrow 0.19$ over 10 steps for blind; $0.64 \rightarrow -0.16$ for matched) but never becomes available in rich-encoder conditions. This is the integration signature of a compensatory cognitive-map code: the LSTM maintains a stable world-frame position representation across many time steps, precisely because the visual encoder cannot do so on its own.

Why The mechanism is direct. In PointNav, GPS and compass are observations that the policy reads each step; the recurrent state only needs to re-encode them if it cannot rely on the visual stream to do so. A full-resolution RGB encoder re-derives position and heading at every step, removing that need. An impoverished encoder cannot, and the policy pressure to act consistently across the full trajectory forces the LSTM to take over. Blind and matched-compute bracket this principle: blind has *no* visual encoder (pure compensation); matched has a *collapsed* 1×1 encoder (partial compensation); foveated, uniform, and foveated-learned all have full encoders (pass-through). Task success is tightly bunched (0.59–0.85 SPL), so the encoding differences are not a skill-gap artefact.

Relation to Wijmans et al. Wijmans et al. [Wijmans et al., 2023] showed that blind PointNav agents develop spatial representations in LSTM memory. Our blind result ($R^2 = 0.95$) replicates this; matched-compute extends it beyond the blind special case. The trigger is encoder-level information bottleneck, not the absence of vision specifically. Appendix D reports an encoder-resolution sweep (32×32 to 256×256 ; in progress) quantifying the bottleneck–compensation relationship as a scaling curve.

Bio parallel The pattern maps cleanly onto the animal-navigation literature cited in our Introduction: congenitally blind humans develop enhanced hippocampal coupling and outperform sighted controls on non-visual spatial tasks [Kupers and Ptito, 2014]; rodent grid cells maintain their periodicity in darkness despite intact head-direction inputs [Chen et al., 2016]. In both cases, absence or impoverishment of visual input shifts spatial bookkeeping into a recurrent substrate — the same qualitative shift we measure in LSTM hidden states.

4.3 H2: Sensor Structure $\rightarrow$ Format-Level Divergence

Finding Five agents that solve the same task at the same success rate do not share a spatial representation. Pairwise linear similarity (CKA) sits near zero for every condition pair (Figure 3); a GPS probe trained on one condition’s hidden states and evaluated on another’s does not generalise at all (Table 2; diagonal 0.90–0.97, off-diagonals all $R^2 \leq -2.7$). In plain terms: what each agent calls “position” is written in a different linear subspace. Visual input type shapes *format*, not just strength.




Figure 3: Unaligned linear CKA [Kornblith et al., 2019] between top-layer LSTM hidden states ($n = 1,907$). Every off-diagonal is below 0.004; the most-similar pair is the two foveated variants (fixed vs. learned gaze) at 0.0035.

Table 2: Probe-transfer R^2 on GPS: train on row condition, test on column. Self-accuracy diagonal 0.90–0.97; every cross-condition transfer is catastrophically negative.

Train \ Test	Blind	Uniform	Fov-fix	Fov-lrn	Matched-compute
Blind	0.97	−9.7	−2.7	−613	−15.7
Uniform	−7.3	0.97	−13.5	−410	−11.5
Fov-fix	−11.3	−5.1	0.96	−232	−10.7
Fov-lrn	−482	−430	−456	0.90	−855
Matched-compute	−5.7	−11.7	−15.5	−390	0.96

Ruling out alternatives Self-probes succeed ($R^2 > 0.90$) for every condition, so probe capacity is not the bottleneck. CKA is computed on a common-sample truncation, ruling out sample-size asymmetry. Task competence is tightly bunched (97–99% success), so the divergence is not a skill difference. And transfer fails symmetrically (both directions), which rules out the possibility that one condition’s code simply embeds another’s.

A t-SNE projection of pooled hidden states confirms the CKA result visually (Appendix A): the five conditions occupy non-overlapping regions of a 2-D manifold.

Boundaries Our similarity tests are linear (CKA) or near-linear (t-SNE default); non-linear similarity measures could reveal higher-order shared structure. Our five sensors are diverse but not exhaustive; hybrid sensors (coarse uniform periphery + sharp fovea, depth-only inputs) might produce intermediate geometries.

Memory transplant Probe-based similarity measures could still miss *functional* equivalence: two representations with low CKA might still be interchangeable if the downstream policy reads them the same way. We test this with memory transplant — paste a donor’s mid-episode hidden state into the recipient’s policy at step t , let the recipient drive the rest of the episode. The isolated condition-mismatch cost (cross-transplant SPL − self-transplant SPL, which subtracts both a small $t=0$ protocol-noise floor and the rollout-divergence cost of re-initialising a recurrent state mid-episode) is stable across midpoints $t \geq 15$ and carries a clear ranking (Figure 4b): foveated-learned→foveated incurs the largest drop (−0.17 SPL), foveated→uniform a smaller one (−0.10), and foveated→blind is essentially neutral (−0.02). (The raw cross-transplant drop vs. baseline is correspondingly 0.19–0.21 for the first two pairs; self-transplant contributes a 0.02–0.04 SPL overhead that accounts for the difference.) The striking piece is the foveated-learned→foveated drop: this is the pair with the *highest* pairwise CKA (0.0035) of any pair we tested, yet it is the most behaviourally incompatible. Higher probe-based similarity does not imply higher functional interchangeability — the representations of agents that share the same sensor transform but differ in gaze location are more compatible to look at than to use. (CKA and probe transfer use common-sample truncation for fair cross-condition




Figure 4: Memory transplant across a midpoint sweep $\{0, 15, 30, 60, 90\}$ steps ($n = 100\text{--}150$ episodes/cell). (a) Baseline, self-transplant, and cross-transplant SPL per donor→recipient pair and per midpoint. At $t = 0$ the donor has accumulated no information yet, so self- and cross-transplant reduce to a protocol-noise floor (self ≈ -0.03 , cross ≈ -0.06 SPL vs. baseline). Self-transplant SPL drops monotonically with increasing midpoint: the recipient’s own past recurrent state genuinely diverges from the state it would have occupied had it driven the rollout from step 0. (b) Cross minus self (the isolated condition-mismatch component, with both protocol noise and rollout divergence subtracted) is roughly stable across midpoints $t \geq 15$: foveated-learned→foveated is the most behaviourally incompatible pair at every midpoint tested (-0.17 SPL on average), *despite having the highest pairwise CKA* (0.0035) of any condition pair—higher probe-based similarity does not predict higher functional interchangeability.

comparison; transplant uses full on-policy rollouts because the task is not truncatable. The two measurement scopes are deliberate, and their conclusions are consistent.) Together with CKA, probe transfer, and t-SNE, transplant provides a fourth convergent signal that the five representations are format-level distinct.

4.4 H3: Gaze Location as an Orthogonal Content Lever

H1 and H2 identify sensor structure as one axis on representational content. H3 asks whether gaze *location* — where the foveated window is centred — is a second, dissociable axis. We test the two forms described in §3.3 and previewed in the Introduction: an active-gaze form (the minimal learned-gaze module should move its gaze to task-useful locations), and a gaze-location form (two foveated agents whose gaze is centred at *different static* locations should produce different content). The active-gaze form is out of scope of our minimal gaze decoder and we do not claim anything about it; the gaze-location form is what our controlled experiments test.

The learned-gaze module collapses Our minimal learned-gaze implementation (deterministic sigmoid MLP, slow-gaze updates once per 128-step rollout segment; Appendix B) does not produce active gaze: after 250M frames the gaze output is frozen at (0.49, 0.62) with per-dimension std $\leq 2.5 \times 10^{-4}$ (numerical noise). Two anti-collapse regularisers (Appendix B) restore gaze variance only at the cost of task performance. We explicitly do not claim this shows active gaze cannot be learned — our gaze decoder is minimal by design (to be a controlled variable, not a model of active vision), and richer architectures (stochastic gaze, information-seeking intrinsic reward, attention-based token selection) would need separate investigation. What the collapse gives us for free is an opportunity to test the gaze-location form: after training, we have an agent that is *effectively foveated with a static gaze 30 pixels below image centre*.

Gaze-location test (in progress) Both rich-encoder foveated variants (foveated-fixed and foveated-learned) sit in the pass-through regime under current-state probing (GPS/compass $R^2 \approx 0$; Table 1). So the H3 test is about whether, *within* the pass-through regime, gaze location still shifts representational content along other axes — for example, whether the LSTM’s non-spatial feature usage differs, or whether the behavioural memory-reliance pattern (shortcut experiment, Figure 5b) differs consistently with gaze location. To isolate gaze-location from training-dynamics confounds, we train




Figure 5: H3 content evidence. (a) Goal-vector probe R^2 : goal *distance* is decodable everywhere ($R^2 \geq 0.78$), but goal *direction* and the full 2-D ego-to-goal vector are at chance or negative for every condition — no condition stores the goal vector (the PointGoal task provides it as a per-step input). What differs is the integrated trajectory and scene-history content each recurrent state accumulates. (b) Shortcut discovery: persistent vs. reset memory on the second of two paired same-scene/different-goal episodes. Per condition we run 20 Gibson scenes $\times$ 10 paired episodes each, and report the mean of the 20 scene-level mean SPLs (not a flat 200-sample mean), to prevent high-episode-count scenes from dominating. Conditions are ordered by relative SPL drop (%). Blind loses most (52%); foveated-learned loses least (15%).

a foveated-shifted control: an agent identical to foveated-fixed except the hardcoded gaze location is (0.49, 0.62) instead of (0.5, 0.5) — matching foveated-learned’s final gaze position but without any learned-gaze module in the optimisation path. **[TODO: foveated-shifted training is in progress (~3-day walltime); the detailed probe and behavioural comparison will populate §4.4 once it completes.]**

Shortcut discovery (behavioural memory-reliance) We separately measure how tightly each policy couples action choice to integrated recurrent memory: run paired episodes in the same scene with fresh goals, compare second-episode SPL under *reset* (LSTM zeroed between episodes) vs. *persistent* (state carried over) memory. Larger relative SPL drop = more brittle memory coupling (Figure 5b): blind loses 52%, uniform 41%, foveated 21%, matched 18%, foveated-learned 15%. This is consistent with H1: blind, the most memory-reliant condition, is also the most brittle when that memory is stale; rich-encoder conditions get more of their decision-relevant signal from instantaneous visual input and are correspondingly less perturbed by stale recurrent state. (The matched-compute case is lower than expected from its strong GPS encoding alone; we read this as evidence that *using* a cognitive map for policy decisions is a separate property from *having* one — matched has the code but weights it less in policy output.)

4.5 Additional Analyses

Controls Two probes fail to discriminate between conditions and so define the boundary of our claims. Coarse occupancy (“have I visited this region?”) is trivially decodable everywhere ($R^2 \in [0.80, 0.83]$), so one-bit spatial working memory is ubiquitous and not part of the H1–H3 story. Metric 20×20 occupancy reconstruction fails everywhere ($R^2 < 0$), consistent with the SPACE benchmark’s finding that fine-grained layout requires non-linear probes or longer horizons [Ramakrishnan et al., 2025]: our claims are about representational *format* and *content*, not about high-resolution metric layout.

Place-cell-like units and multi-layer depth Two further analyses confirm that our findings are about representational *content*, not about where we probe from or how we probe (Figure 6). First, four of five conditions distribute spatial encoding broadly across LSTM units in the Banino et al. [Banino et al., 2018] sense ($> 90\%$ of units carry > 1 bit of spatial information); foveated-learned is the clear outlier (56% of units), consistent with its specialisation away from position. Second, foveated-learned’s heading $\gg$ position dissociation is present at *every* LSTM layer — this is a property of the representation, not of a linear probe on one vantage point.




Figure 6: *Left (Fig. 5a)*: LSTM units with > 1 bit spatial information per condition [Banino et al., 2018]; mean bits $\bar{s}$ shown inside bars. Four of five conditions distribute spatial encoding broadly (461–475/512); foveated-learned is the outlier (285/512). *Right (Fig. 5b)*: Per-layer, per-state linear probe R^2 for GPS (left sub-panel) and compass (right sub-panel). Foveated-learned’s compass- $\rightarrow$ -position dissociation is present at every layer, not just the top.




Figure 7: **[Stochastic-collection data; being re-collected under deterministic protocol.]** Probe generalisation from Gibson evaluation (unshaded) to held-out MP3D evaluation (hatched), 300 episodes/condition. Same trained checkpoints; no re-training. These results were collected with the stochastic action-selection protocol described in §3.3 and are shown here only as a directional indication that the H1 bottleneck ordering is preserved across scene sets. Updated numbers using the deterministic protocol will replace this figure; see §4.5.

Generalisation to MP3D **[TODO: MP3D numbers under the new deterministic-probe protocol are being re-collected; previous MP3D probes used the stochastic-sampling collection that we retired in §3.3. The updated MP3D comparison will populate this paragraph once available. Expected: the H1 bottleneck ordering (blind + matched vs. rich-encoder conditions) is preserved across datasets, since the mechanism is agent-internal and does not depend on scene statistics.]**

5 Discussion

5.1 Two orthogonal axes: sensor structure and gaze location

Our findings unify under a single mechanism: the encoder–memory race. Whether the LSTM encodes world-frame spatial variables depends on whether the visual encoder can resolve them from the current frame. When the encoder cannot — as in blind (no encoder) or matched-compute (1×1 encoder feature map collapsing spatial structure) — the LSTM takes over and maintains a persistent cognitive-map-like code across many time-steps of lag. When the encoder can — uniform, foveated, foveated-learned all use full-resolution ResNet18 encoders — the LSTM drifts to pass-through. H1 quantifies this directly: GPS decodes from the recurrent state at $R^2 = 0.95$ for blind and $R^2 = 0.78$ for matched-compute, and at chance for the three rich-encoder conditions. H2 shows that *what* the LSTM encodes — when it encodes anything — is linearly distinct across conditions (CKA < 0.004 every pair; transplant SPL drops 0.19–0.21), i.e. each condition develops its own representational

subspace even on the pass-through side. H3 asks whether, within the rich-encoder pass-through regime, gaze location biases what non-spatial features the LSTM carries; the foveated-shifted causal control that answers this is in training at submission time. The same encoder-capacity axis that collapses the “foveated vs. uniform” question in the original framing (both sit in the pass-through regime) surfaces as a clean scaling prediction: GPS decodability should fall monotonically with encoder resolution, which we begin to test in Appendix D.

5.2 Why this, not alternative accounts

We rule out four alternative explanations for the bottleneck-compensation pattern. First, *task-competence differences*: all five conditions solve PointGoal at 93–99% success, so differences in encoded content cannot be attributed to differences in navigation skill. Second, *probe-capacity artefacts*: Hewitt–Liang label-permutation selectivity exceeds 0.76 for every condition; DtG decodes robustly everywhere ($R^2 \geq 0.81$), ruling out a bottleneck in the probe’s ability to fit high- R^2 scalars. The chance-level GPS/compass R^2 for rich-encoder conditions is a property of the hidden state content, not of the probe. Third, *sample-size or trajectory-distribution artefacts*: the CV σ spans for all five conditions are from the same 500-episode, ~ 100 –400 step rollouts, using identical deterministic evaluation protocol. Fourth, *linear-probe artefacts*: the behavioural interventions (memory transplant in §4.3, shortcut discovery in §4.4) go outside the probe framework and give the same ordering, with cross-condition transplants dropping SPL by 0.19–0.21 between rich-encoder pairs.

5.3 Biological precedent as convergent evidence

Each of our three findings has a parallel in the animal-navigation literature, which we take as one piece of convergent (not conclusive) evidence that the effects are not specific to a single LSTM architecture. Compensatory memory under reduced sensory access is documented in congenitally blind humans, who develop enhanced hippocampal coupling and outperform sighted controls on non-visual spatial-working-memory tasks [Kupers and Ptito, 2014], and in rodents whose grid-cell periodicity collapses in darkness despite intact head-direction inputs [Chen et al., 2016]—both format-level, not just strength-level, effects. Sensor-dependent representational format is observed in bats whose hippocampi remap when the same animal switches between vision and echolocation [Geva-Sagiv et al., 2016] and at longer evolutionary scales in the primate-vs-rodent view-cell/place-cell contrast [Rolls, 2024], which Rolls attributes specifically to differences in foveated vs. near-panoramic sampling statistics. Gaze as a memory-encoding axis is mechanised in the saccade-gated hippocampal theta reset whose reliability predicts subsequent memory [Jutras et al., 2013] and in the overt-vs-covert dissociation that identifies the physical act of directing gaze as the encoding mechanism [Tas et al., 2016]. We do not claim our agents implement these biological mechanisms, nor that the parallel structure settles the generality question; we do claim that the parallels make a pure single-architecture artefact less likely, and that the controlled agent setting complements biological comparisons where sensor type, architecture, and training objective cannot be independently manipulated.

5.4 Implications for ML practice and architectural scope

Within the recurrent PointNav setting we study, three concrete implications follow. (i) *Encoder capacity is a representational-content axis*: LSTM spatial encoding is driven by whether the visual encoder can resolve world-frame position/heading on its own, not by task pressure alone. Rich encoders offload spatial bookkeeping to the visual stream; impoverished encoders force it into the recurrent state. This predicts the encoder-capacity scaling curve shown in Figure 9 and provides a practical lever: to induce a cognitive-map-like internal representation in an RL agent, one can impose an encoder bottleneck. (ii) *“More pixels” is not the right lens*: the relationship is not monotonic in pixel count — it is monotonic in encoder-preserved information. Uniform 256×256 (full encoder) lacks the LSTM spatial code; 48×48 with the same encoder stack (matched-compute, 1×1 feature map) has a strong one. (iii) *Attention policy (H3)*: within the rich-encoder regime where LSTM does *not* encode world-frame spatial info by default, shifting the static gaze location alters what the LSTM *does* encode (details in §4.4, pending foveated-shifted convergence).

Beyond the recurrent setting (untested) We use a Wijmans-style recurrent (LSTM) backbone for direct comparability with the emergent-map literature. The three implications above are established only in this setting, but we speculate about where they plausibly extend. Transformer-based

embodied agents build their internal state by accumulating information via learned attention, which is structurally similar to our gaze location in the sense that it selects which observation content is read into the representation; our H3 finding therefore suggests a testable prediction for transformer navigators—strong constraints on attended tokens should produce content-level dissociations of the sort we observe—but the prediction has not been verified and is left for future work. The sensor-structure axis (H1, H2) is a claim specifically about RL recurrent navigation agents trained under partial observability; whether equivalent representational-content effects arise in supervised visual learners, non-visual modalities, or multimodal systems is an open empirical question that the present experiments do not resolve.

5.5 Limitations

(i) *Intermediate checkpoint for foveated-fixed*: all foveated-fixed analyses use ckpt.36 at $\sim 174\text{M}$, the last checkpoint before the silent-corruption window (Appendix C); the other four conditions use the final converged checkpoint. (ii) *Minimal learned-gaze architecture*. Our learned-gaze module is a deterministic sigmoid MLP with slow-gaze rollout updates, chosen for computational compatibility with DD-PPO rollouts and simplicity as a controlled variable, not as a competitive active-gaze model. Its collapse under pure task supervision is specific to this architecture (and so far to a single seed; multi-seed replication is in progress); richer gaze mechanisms (stochastic gaze with entropy bonus, information-seeking intrinsic reward, attention-based token selection) would need separate investigation. We therefore do not claim that implicit task pressure cannot produce active gaze in general; we claim only that our minimal decoder collapsed, and we use the collapsed-gaze position as an empirical starting point for the gaze-location test (foveated-shifted control, §4.4) — a test whose logic does not depend on the gaze decoder working. (iii) *Cross-heading probe*: the proposal included a cross-heading generalisation probe (train on heading *A*, test on opposite heading). Our implementation returned “insufficient samples” sentinels at our probe sizes and does not appear in our results. (iv) *Foveation model*: Gaussian blur with quadratic falloff is a simplification; log-polar or multi-scale pyramid models may behave differently. (v) *Linear probes only*: non-linear probes might reveal structure Ridge regression misses, particularly for occupancy maps. (vi) *Architecture scope*: we use a recurrent (LSTM) backbone to remain comparable with the emergent-map literature we build on. The core claims (sensor structure and gaze location as representational-content axes) should hold for transformer-based navigation agents whose internal state is built by learned attention, but directly verifying this requires re-running the five-condition ablation on a transformer backbone and is left for future work; see §5 for the specific predictions.

6 Conclusion

We ran a controlled five-condition ablation that isolates the effect of visual-input capacity on the recurrent representations learned by an otherwise-identical DD-PPO navigation agent. Three findings emerge. (1) *Encoder-bottleneck compensation*. LSTM world-frame spatial encoding monotonically tracks visual-encoder information content: blind and matched-compute (1×1 feature-map encoder) develop persistent cognitive-map-like codes (GPS $R^2 = 0.95 \pm 0.02$, 0.78 ± 0.10 respectively, stable across 10+ steps of lag), while uniform, foveated, and foveated-learned all decode world-frame variables at chance. DtG (the ego-relative, task-relevant quantity) decodes robustly everywhere, confirming the agents’ LSTMs are not simply inert. (2) *Condition-specific representational format*. Even where rich-encoder LSTMs are spatially pass-through, their representations are not interchangeable: pairwise CKA stays below 0.004 and cross-condition mid-episode memory transplants drop SPL by 0.19–0.21. Whatever non-spatial task-relevant features each LSTM carries, it carries them in its own linear subspace. (3) *Gaze-location test (pending)*. Whether gaze location biases the non-spatial content that a rich-encoder LSTM retains is the question H3 asks; the cleanest test (foveated-shifted training with hardcoded gaze at the learned-gaze module’s collapsed position) is in progress at submission and will report in §4.4. Bio-analog: the finding replicates and generalises the blind-agent cognitive-map result of Wijmans et al. [Wijmans et al., 2023], and aligns with the enhanced-hippocampal-coupling literature in blind humans [Kupers and Ptito, 2014] and the darkness-grid-cell literature [Chen et al., 2016] — both cases of sensory absence forcing a shift to recurrent spatial bookkeeping. Open questions include (i) whether the monotonic encoder-bottleneck ordering extends to finer-grained resolution sweeps (Appendix D), (ii) whether H3 stands on a clean causal test, (iii) whether transformer-based navigation agents reproduce the bottleneck-compensation mecha-

nism, and (iv) whether analogous encoder-bottleneck effects shape representations in non-navigation embodied tasks.

References

- Andrei Atanov, Yicheng Fu, Gaurav Singh, Jianren Yu, Andrew Spielberg, and Amir Zamir. How far can a 1-pixel camera go? solving vision tasks using photoreceptors and computationally designed visual morphology. In *European Conference on Computer Vision (ECCV)*, 2024.
- Andrea Banino, Caswell Barry, Benigno Uria, Charles Blundell, Timothy Lillicrap, Piotr Mirowski, Alexander Pritzel, Martin J Chadwick, Thomas Degris, Joseph Modayil, et al. Vector-based navigation using grid-like representations in artificial agents. *Nature*, 557(7705):429–433, 2018.
- Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics*, 48(1):207–219, 2022.
- Guifen Chen, John A. King, Neil Burgess, and John O’Keefe. Absence of visual input results in the disruption of grid cell firing in the mouse. *Current Biology*, 26(17):2335–2342, 2016.
- Christopher J Cueva and Xue-Xin Wei. Emergence of grid-like representations by training recurrent neural networks to perform spatial localization. In *International Conference on Learning Representations (ICLR)*, 2018.
- Kshitij Dwivedi, Gemma Roig, Aniruddha Kembhavi, and Roozbeh Mottaghi. What do navigation agents learn about their environment? In *Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Maya Geva-Sagiv, Sandro Romani, Liora Las, and Nachum Ulanovsky. Hippocampal global remapping for different sensory modalities in flying bats. *Nature Neuroscience*, 19(7):952–958, 2016.
- James Gornet and Matt Thomson. Automated construction of cognitive maps with visual predictive coding. *Nature Machine Intelligence*, 6:820–833, 2024.
- Jacqueline Gottlieb, Pierre-Yves Oudeyer, Manuel Lopes, and Adrien Baranes. Information-seeking, curiosity, and attention: computational and neural mechanisms. *Trends in Cognitive Sciences*, 17(11):585–593, 2013.
- Mary M. Hayhoe and Jonathan S. Matthis. Control of gaze in natural environments: effects of rewards and costs, uncertainty and memory in target selection. *Philosophical Transactions of the Royal Society B*, 373(1718):20170368, 2018.
- Jay A Hennig, Sandra A Romero Pinto, Takahiro Yamaguchi, Scott W Linderman, Naoshige Uchida, and Samuel J Gershman. Emergence of belief-like representations through reinforcement learning. *PLOS Computational Biology*, 19(9):e1011067, 2023.
- John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2019.
- Michael J. Jutras, Pascal Fries, and Elizabeth A. Buffalo. Oscillatory activity in the monkey hippocampus during visual exploration and memory formation. *Proceedings of the National Academy of Sciences*, 110(32):13144–13149, 2013.
- Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *International Conference on Machine Learning (ICML)*, 2019.
- Ron Kupers and Maurice Ptito. Compensatory plasticity and cross-modal reorganization following early visual deprivation. *Neuroscience & Biobehavioral Reviews*, 41:36–52, 2014.
- Volodymyr Mnih, Nicolas Heess, Alex Graves, and Koray Kavukcuoglu. Recurrent models of visual attention. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2014.
- Jiri Najemnik and Wilson S. Geisler. Optimal eye movement strategies in visual search. *Nature*, 434(7031):387–391, 2005.

- Motahareh Pourrahimi and Pouya Bashivan. Emergent brain-like representations in a goal-directed neural network model of visual search. *bioRxiv preprint*, 2025.
- Santhosh Kumar Ramakrishnan, Erik Wijmans, Philipp Krähenbühl, and Vladlen Koltun. Does spatial cognition emerge in frontier models? In *International Conference on Learning Representations (ICLR)*, 2025.
- Edmund T. Rolls. Hippocampal spatial view cells for memory and navigation, and their underlying connectivity in humans. *Hippocampus*, 33(6):533–572, 2023.
- Edmund T. Rolls. Why do primates have view cells instead of place cells? *Trends in Cognitive Sciences*, 28(9):862–874, 2024.
- Ruth Rosenholtz. Capabilities and limitations of peripheral vision. *Annual Review of Vision Science*, 2:437–457, 2016.
- Manolis Savva, Abhishek Kadian, Oleksandr Maksymets, Yili Zhao, Erik Wijmans, Bhavana Jain, Julian Straub, Jia Liu, Vladlen Koltun, Jitendra Malik, Dhruv Batra, and Devi Parikh. Habitat: A platform for embodied AI research. In *International Conference on Computer Vision (ICCV)*, 2019.
- Jinghuan Shang and Michael S. Ryoo. Active vision reinforcement learning under limited visual observability. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Nathan Sprague and Dana H. Ballard. Eye movements for reward maximization. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2003.
- Hans Strasburger, Ingo Rentschler, and Martin Jüttner. Peripheral vision and pattern recognition: A review. *Journal of Vision*, 11(5):13, 2011.
- Amanda C. Tas, Steven J. Luck, and Andrew Hollingworth. The role of attention in the encoding of visual working memory. *Journal of Experimental Psychology: Human Perception and Performance*, 42(8):1121–1138, 2016.
- Erik Wijmans, Abhishek Kadian, Ari Morcos, Stefan Lee, Irfan Essa, Devi Parikh, Dhruv Batra, and Manolis Savva. DD-PPO: Learning near-perfect pointgoal navigators from 2.5 billion frames. In *International Conference on Learning Representations (ICLR)*, 2020.
- Erik Wijmans, Manolis Savva, Irfan Essa, Stefan Lee, Ari Morcos, and Dhruv Batra. Emergence of maps in the memories of blind navigation agents. In *International Conference on Learning Representations (ICLR)*, 2023.
- Sylvia Wirth, Pierre Baraduc, Angélique Planté, Serge Pinede, and Jean-René Duhamel. Gaze-informed, task-situated representation of space in primate hippocampus during virtual navigation. *PLOS Biology*, 15(2):e2001045, 2017.
- Fei Xia, Amir R Zamir, Zhiyang He, Alexander Sax, Jitendra Malik, and Silvio Savarese. Gibson Env: Real-world perception for embodied agents. In *Computer Vision and Pattern Recognition (CVPR)*, 2018.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer:
answerYes

Justification: The abstract and §1 state the three hypotheses H1/H2/H3 with explicit signatures (what pattern supports each), and preview the three findings with specific magnitudes (path-history $R^2=0.56$, CKA <0.004 , compass $R^2=0.94$ vs GPS $R^2=0.03$). These are expanded in §4 and §5 with no stronger claims than the data support.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer:

answerYes

Justification: §5.4 explicitly lists six limitations (i–vi) covering checkpoint staleness, single-seed H3, missing proposed analyses, foveation-model simplification, linear-probes-only, and LSTM-architecture scope.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer:

answerNA

Justification: The paper reports empirical results; it contains no theoretical claims or proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer:
answerYes

Justification: §3.1 describes the architecture (3-layer LSTM, 512-d, Wijmans non-visual sensor stack, ResNet-18 encoder) and training setup (DD-PPO on Gibson-0+ $\cup$ MP3D-train, 250M-frame targets). §3.2–§3.6 describe all probing and cross-condition protocols. Appendix A details the gaze-diversity ablation. Code will be released upon publication.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer:
answerYes

Justification: §3.7 commits to open-source release of the training and probing code. Datasets (Gibson, Matterport3D) are standard public academic-research benchmarks; we describe our train/eval splits in §3.1.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).

- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer:

answerYes

Justification: §3.1 specifies training pool and eval set; §3.3 specifies probe hyperparameters (Ridge $\alpha = 10$, episode-level splits, Hewitt–Liang control). Appendix A gives the ablation’s loss formulation and coefficient.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer:

answerNo

Justification: We report single-seed numbers per condition. This is documented as Limitation (ii) in §5.4. Multi-seed replication for tighter confidence intervals is in the immediate follow-up.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).

- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer:

answerYes

Justification: §3.7 reports per-condition training budget (5–7 days on 2×V100 for 250M frames) and probing budget (<1h per condition on one V100).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer:

answerYes

Justification: The research uses public academic-research benchmarks, publicly-released open-source RL frameworks, and involves no human subjects or deployment. We also report and remediate a numerical-stability bug upstream (Appendix B).

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer:

answerYes

Justification: §5.3 discusses implications — sensor design as a tunable axis for interpretable embodied agents. Negative societal impacts are minimal: the work is fundamental research on simulated indoor navigation with no direct deployment pathway.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer:

answerNA

Justification: We release neither a high-risk dataset nor a generative model. Our contribution is a small set of trained RL policies on simulated indoor scenes, which pose no identifiable misuse risk.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer:

answerYes

Justification: Gibson Database and Matterport3D are used under their released academic-research licenses; Habitat-Sim and habitat-lab are MIT-licensed. Our use complies with their terms; citations §3.1.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.

- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer:

answerYes

Justification: Code will be released with documentation upon publication. Figures are novel; they are released under the paper's license alongside the code. The gaze-diversity auxiliary loss and NaN-sanitisation patch are small standalone modules with docstrings.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer:

answerNA

Justification: The paper involves no crowdsourcing and no human-subjects data.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer:

answerNA

Justification: No human-subjects research; no IRB/equivalent review was required.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer:

answerNA

Justification: LLMs were used only for writing assistance on the manuscript. They are not part of the core research methodology; all empirical results come from standard RL training and linear probing.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.

A Supplementary visualisations




Figure 8: Supplementary visualisations referenced from the main text.

B Gaze-Diversity Auxiliary Loss (H3 Ablation)

Motivation In the first foveated-learned training run the gaze decoder collapsed (per-dimension $\text{std} \approx 5 \times 10^{-4}$). We attribute this to two mechanisms: (i) once the sigmoid output saturates, its derivative $\sigma'(x) \rightarrow 0$ attenuates gradient signal; (ii) consistent-gaze views yield more predictable visual features, so the reward-gradient path has no explicit incentive for variety.

Target-variance hinge We add $\mathcal{L}_{\text{div}} = \lambda \cdot \text{relu}(\sigma_{\text{target}}^2 - \text{Var}_{\text{batch}}[g])$ where $g \in [0, 1]^2$, $\text{Var}_{\text{batch}}$ averaged over the two gaze dimensions, $\sigma_{\text{target}}^2 = 0.01$ ($\text{std} \approx 0.1$), $\lambda = 0.01$. The hinge disappears once variance exceeds the target. We also evaluated an un-capped $-\lambda \cdot \text{Var}[g]$ variant as a negative control.

Pilot 1 (uncapped, job 2840256) After 10 M frames, gaze variance exceeded that of a uniform distribution over $[0, 1]^2$ ($\text{std} [0.43, 0.42]$ vs. theoretical max 0.289): gaze was essentially random. Task performance collapsed: SPL plateaued at 0.05, running-mean metrics transitioned to NaN at ~ 8 M frames. The uncapped regulariser destroyed the consistent-foveation signal the policy needs.

Pilot 2 (hinge, job 2842288) Gaze converged to an extreme corner of image space: mean (0.001, 0.998) with $\text{std} [0.020, 0.023]$ — well below the target std of 0.1 and therefore with the hinge active throughout. Task SPL plateaued at 0.05 as in pilot 1. The policy apparently found the minimum-task-cost gaze that partially satisfies the diversity pressure: a sigmoid saturated at one corner, producing a constant image-corner view.

Interpretation Both pilots argue that in our deterministic sigmoid-MLP gaze architecture with slow-gaze rollouts, gaze collapse is a pre-requisite for task learning rather than an accident. Task signal alone does not spontaneously produce informative fixation, and an exogenous anti-collapse regulariser moves gaze toward whatever region reduces regulariser loss fastest, which is generically *not* the task-useful region. A genuine H3 test likely requires (i) a stochastic gaze location with entropy bonus, so diversity is sampled at rollout time rather than enforced on a deterministic output; or (ii) a task-aligned intrinsic reward (e.g., information gain about goal-relative position); or (iii) an architectural change separating gaze generation from visual-feature extraction (e.g., attention-based gaze).

C Training Stability: Silent Weight Corruption in DD-PPO

During an initial foveated training run we observed a failure mode that did not manifest as a crash but as slow, silent corruption of the learned weights. After $\approx 170\text{M}$ frames of stable progress (reward ≈ 10 , SPL ≈ 0.83), one mini-batch on one DDP worker produced a NaN gradient. Because `torch.nn.utils.clip_grad_norm_` does not sanitise non-finite gradients—the global norm becomes NaN, the clip coefficient becomes NaN, and `grad *= NaN` remains NaN—the subsequent `optimizer.step()` wrote NaN into every parameter. DDP all-reduce propagated the corruption across workers. The training loop did not raise: it continued for 2.5 days with reward=NaN and SPL= 0 until `torch.multinomial` finally rejected the NaN probability tensor. All checkpoints saved during this window were unrecoverable (90 of 97 parameter tensors entirely NaN).

To isolate the origin of the NaN gradient we re-ran training from the last clean checkpoint with `torch.autograd.set_detect_anomaly(True)` and per-module forward-output hooks. The anomaly traceback pinned the failure to `MseLossBackward0`, i.e. the backward pass of the value-head loss $\frac{1}{2}(V(s) - R)^2$; rare return-vs-value discrepancies overflowed float32 in the squaring step, producing non-finite gradients downstream.

Our mitigation is a minimally invasive patch to `PPO.before_step` that zeroes any non-finite gradient element before clipping. A bad mini-batch becomes a no-op update rather than a catastrophic one. The patch is a no-op on clean training paths. Across the five conditions reported here, no NaN events were observed during full training after deployment, and all final checkpoints contain only finite weights. We reported the underlying issue upstream (habitat-lab issue #2226) with a minimal reproduction and a proposed native fix.

D Encoder-capacity scaling sweep

[TODO: The scaling sweep trains matched-compute variants at input resolutions {32, 48, 64, 96, 128, 192} alongside the existing blind (no encoder) and uniform 256×256 anchors. Probe results on the trained checkpoints will populate Figure 9 and this appendix section. Hypothesis: GPS probe R^2 decreases monotonically with input resolution, from $R^2 \approx 0.95$ at no-encoder to $R^2 \approx 0$ at 256×256 , crossing into the pass-through regime between 48 and 128 pixels.]




Figure 9: Encoder-capacity scaling curve. **[TODO: Pending data.]** x-axis: input resolution in pixels; y-axis: GPS and compass probe R^2 (5-fold CV, $\text{mean} \pm \sigma$).