

---

# How Foveated Vision Shapes Cognitive Maps in Navigation Agents

---

Anonymous authors

Paper under double-blind review

## Abstract

What determines the content of a recurrent navigation agent’s learned internal state—the task, the architecture, or the structure of the visual input? Prior work shows that spatial codes emerge from task pressure alone, almost exclusively under spatially uniform visual input, a simplifying convention met in no biological vision system. Meanwhile, efficient visual-intelligence systems increasingly break the uniform assumption (foveated rendering, event cameras, learned attention), raising a basic question: are such non-uniform designs purely computational tools for efficiency, or do they also change what the downstream system encodes? We answer this empirically in the narrow but controlled setting of recurrent PointNav agents. We train DD-PPO agents in Habitat under five visual conditions that differ only in sensor structure and gaze policy (blind, uniform, foveated-fixed, matched-compute, foveated-learned-gaze), share the same recurrent backbone, and are trained to the same task budget. All sighted conditions solve the task. Probing their internal representations with linear decoders, representational-similarity analysis, probe transfer, and two behavioural interventions reveals that sensor structure and gaze policy are two axes shaping the *format*—not merely the strength—of learned spatial memory. Non-uniform vision induces compensatory temporal memory that a matched-pixel uniform control cannot recreate; different sensor/gaze combinations produce distinct internal geometries not interchangeable via mid-episode memory transplant, even between agents sharing the same foveation transform; and a collapsed learned gaze steers the representation from position-dominant to heading-dominant content without altering task performance. The effects reproduce on held-out MP3D scenes and have direct analogs in primate, rodent, and bat spatial-memory literatures—convergent evidence within the tested scope. Whether the same axes shape representations in transformer-based navigators or in non-navigation systems is a natural question raised by these findings, not settled by them.

## 1 Introduction

Recurrent navigation agents trained with deep RL build internal states that support memory, planning, and generalisation, and the content of those internal states is at least as interesting as the task metrics they drive. A natural question is what *determines* that content: architecture, objective, training budget, or visual input? Wijmans et al. [Wijmans et al., 2023] showed that map-like spatial representations emerge from task pressure alone, even in blind agents that receive only goal displacement and proprioception; Dwivedi et al. [Dwivedi et al., 2022] probed sighted agents for related concepts; Banino et al. [Banino et al., 2018] and Cueva & Wei [Cueva and Wei, 2018] demonstrated grid-cell-like structure in path-integration networks. Together these results establish that *some* spatial structure is learned without being designed in. They leave open the complementary question that we take up here: within agents that all achieve high task success, how much of *what* recurrent memory encodes is determined by the structure of the visual input rather than by the task? In biology the answer is clear




Figure 1: The five visual conditions differ only in what reaches the agent’s encoder. (a) Blind: no visual input; only the non-visual sensor stack (goal-in-start-frame, GPS, compass, close-to-goal scalar) — the Wijmans [Wijmans et al., 2023] configuration. (b) Uniform:  $256 \times 256$  full-resolution RGB. (c) Foveated (fixed gaze): eccentricity-dependent Gaussian blur centred on the image middle ( $\sigma_{\max} = 8$ , quadratic falloff). The two additional conditions not shown are *matched-compute* (a down-sampled uniform view with the same total pixel budget as foveated) and *foveated (learned gaze)*, which adds a learned  $(x, y)$  gaze offset to (c). (d) All conditions solve PointGoal navigation in the same Habitat Matterport3D / Gibson scenes; the top-down view is shown for reference only and is not available to the agent.

— primate hippocampal cells code joint place–view–action state because primates sample the scene with a foveated sensor [Wirth et al., 2017, Rolls, 2024], and the same bat produces non-overlapping hippocampal populations in vision versus echolocation [Geva-Sagiv et al., 2016]. We investigate the analogous question in a controlled agent setting where the only free variables are the visual input structure and the gaze that directs it.

Our five-condition design isolates three degrees of freedom: sensor type (blind / uniform / foveated), information volume (uniform vs. matched-compute), and gaze policy (fixed vs. learned). Every condition uses the same DD-PPO learner, the same Gibson-0+MP3D training pool, the same LSTM backbone, and the same non-visual sensor stack (Figure 1). We adopt the LSTM architecture from Wijmans et al. [Wijmans et al., 2023] deliberately: the emergent-map literature we build on is recurrent, and matching that architecture is what lets our findings be interpretable as statements about *the structure of the task-shaped memory* rather than about the recurrent-versus-transformer design choice. We make our claims specifically about recurrent PointNav agents, and discuss in §5 whether the logic plausibly extends to transformer-backed or non-navigation systems (an empirical question we do not settle here). Within this setup, two accounts make divergent predictions for what a foveated agent’s memory should do. A *passive buffer* stores what it currently perceives, so decodable content should fall when current input is degraded. A *compensatory* memory invests in internal state when current input is unreliable, so decodable content should persist over longer horizons than for an agent whose current input is always reliable. Which pattern the agent converges to, and whether gaze policy can act as a separate representational-content lever even when task performance is held fixed, are the questions this paper asks empirically.

We test three hypotheses, each aligned with one of the two claimed axes. **H1 (sensor structure → compensatory memory)**. Spatially non-uniform visual input should induce a longer-lived spatial trace in internal memory than uniform input, controlling for total information volume. Signature: probing  $R^2$  for past positions stays high for non-uniform conditions at longer lags and drops for uniform; the matched-compute control (same pixel count, uniform distribution) does not recreate the effect. **H2 (sensor structure → format-level divergence)**. Different sensor conditions should

produce internal representations in distinct linear subspaces, not scaled versions of a common code. Signature: pairwise CKA near zero, probe weights fail to transfer across conditions (while within-condition probes succeed), and cross-condition memory transplants degrade task performance beyond same-condition self-transplants. **H3 (gaze policy  $\rightarrow$  orthogonal content lever)**. The policy that directs where a foveated sensor is centred should itself act as a representational-content axis, independent of the sensor transform. We test two forms: (a) a strong *epistemic* form — learned end-to-end, gaze should track per-step memory uncertainty, measurable as correlation between the learned gaze coordinate and the fixed-gaze variant’s per-step position-probe residual; (b) a broader *content-steering* form — the representation produced under a learned gaze should differ in content from the fixed-gaze variant even when task performance matches. Each hypothesis has a direct analog in the animal-navigation literature [Kupers and Ptito, 2014, Chen et al., 2016, Geva-Sagiv et al., 2016, Jutras et al., 2013, Najemnik and Geisler, 2005], which we return to in Discussion as evidence that the effects reflect general sensor–memory coupling principles rather than architectural artefacts.

**Findings** H1 and H2 are supported by clean large-magnitude effects (path-history  $R^2 = 0.56$  at lag 5 for foveated vs.  $-1.53$  for uniform; cross-condition CKA  $< 0.004$  on every pair). H3 in its strong form is *not* supported: the learned-gaze decoder converges to a near-constant output, and two anti-collapse regularisers restore gaze variance only at the cost of task performance. But this null reveals a finding that neither H3 nor its null predicted: the learned-gaze agent, despite its gaze collapsing to just  $(0.49, 0.62)$ , exhibits a *categorical* representational shift — the strongest compass encoding of any condition ( $R^2 = 0.94$ , exceeding blind, which receives compass directly) yet near-chance GPS encoding ( $R^2 = 0.03$ ), at every LSTM layer. Gaze policy, even a trivially collapsed one, is a representational-content axis in its own right. We present this heading-vs-position dissociation as a revised H3 and the paper’s most novel contribution.

## 2 Related Work

**Cognitive maps in RL navigation agents** Wijmans et al. [Wijmans et al., 2023] demonstrated that map-like representations emerge in the recurrent memories of DD-PPO agents even without visual input; Dwivedi et al. [Dwivedi et al., 2022] probed sighted agents for related spatial concepts; Banino et al. [Banino et al., 2018] and Cueva & Wei [Cueva and Wei, 2018] found grid-cell-like units in recurrent networks trained on path integration; Gornet & Thomson [Gornet and Thomson, 2024] extended this via visual predictive coding. Hennig et al. [Hennig et al., 2023] showed that RL under partial observability can induce belief-state-like representations without explicit probabilistic objectives. Ramakrishnan et al. [Ramakrishnan et al., 2025] examined spatial cognition in frontier models. All these studies use either absent or spatially uniform perception; none vary input quality spatially, and none test whether different sensors produce different representational formats.

**Foveated vision and gaze policies in deep learning** Foveated rendering [Strasburger et al., 2011, Rosenholtz, 2016] exploits eccentricity-dependent acuity falloff for efficiency. Learned glimpse policies appear in recurrent attention models [Mnih et al., 2014]; Atanov et al. [Atanov et al., 2024] showed that sensor layout changes downstream task performance. Pourrahimi & Bashivan [Pourrahimi and Bashivan, 2025] probed a foveated visual-search model for neuroscience-predicted features; Shang & Ryoo [Shang and Ryoo, 2023] jointly learned gaze and motor policies following the Sprague & Ballard [Sprague and Ballard, 2003] framework. These address visual search or task performance; none asks how sensor structure reshapes the learned spatial memory of a downstream navigation policy.

**Probing neural representations** Belinkov [Belinkov, 2022] surveys probing classifiers; Hewitt & Liang [Hewitt and Liang, 2019] introduced control tasks to distinguish genuine structure from probe memorisation; Kornblith et al. [Kornblith et al., 2019] introduced linear CKA for cross-model similarity. We adopt linear probing, the Hewitt–Liang selectivity metric, and unaligned linear CKA, and complement them with two behavioural interventions directly on the recurrent state (memory transplant and shortcut discovery) to validate the probe-based findings outside the linear-probe framework.

**Biological precedent** Our hypotheses and framing draw on three threads of animal-navigation research. (i) Primate hippocampal cells code joint place–view–action state [Wirth et al., 2017, Rolls, 2023]; Rolls [Rolls, 2024] argues that primate *spatial-view cells* and rodent *place cells* reflect the different sampling statistics of foveated versus near-panoramic vision. (ii) The same

bat in the same environment produces non-overlapping hippocampal populations when it switches between vision and echolocation [Geva-Sagiv et al., 2016]—format-level divergence across sensor modalities within a single animal. (iii) In primates, each saccade resets hippocampal 3–8 Hz rhythms and the reliability of the reset predicts subsequent memory [Jutras et al., 2013]; only overt (not covert) attention automatically writes to visual working memory [Tas et al., 2016]. Gaze targets themselves are selected in an information-gain-maximising fashion consistent with ideal-observer predictions [Najemnik and Geisler, 2005, Gottlieb et al., 2013, Hayhoe and Matthis, 2018]. We use these results only to frame our ML predictions and interpret our findings; our contribution is the controlled ablation in a learned agent that cross-species and cross-modality biological comparisons cannot cleanly supply.

### 3 Methods

#### 3.1 Agent Architecture and Training

All agents are trained on PointGoal navigation in Habitat [Savva et al., 2019] using DD-PPO [Wijmans et al., 2020] on a merged pool of Gibson-0+ [Xia et al., 2018] (411 scenes) and Matterport3D train (61 scenes), totalling 472 training scenes. Evaluation uses 18 held-out MP3D test scenes (1008 episodes).

All agents share an identical 3-layer LSTM (512-d) backbone and the Wijmans non-visual sensor stack: goal-in-start-frame, GPS, compass, and a close-to-goal scalar, each projected to 32-d and concatenated with a 32-d previous-action embedding. All sighted conditions are trained for a target of 250 M environment frames; the blind baseline is trained to 342 M frames to allow slower convergence from proprioception alone.

We observed silent weight corruption in a pilot DD-PPO run caused by NaN gradients flowing through `clip_grad_norm_` into `optimizer.step()`; we mitigate this with a gradient-sanitisation patch to `PP0.before_step`, described in Appendix B.

#### 3.2 Visual Conditions

Illustrative examples of the five conditions are shown in Figure 1. Specifically:

Five conditions vary only the visual input (Figure 1):

1. *Blind*: no visual input; non-visual sensor stack only, replicating Wijmans et al. [Wijmans et al., 2023].
2. *Uniform*: full-resolution  $256 \times 256$  RGB encoded by ResNet-18.
3. *Foveated (fixed)*: eccentricity-dependent Gaussian blur ( $\sigma_{\max} = 8$ , quadratic falloff) with gaze locked at image centre; same ResNet-18 encoder.
4. *Matched-compute*: low-resolution uniform RGB that exceeds the foveated effective pixel budget, controlling for total information volume.
5. *Foveated (learned)*: same blur model, but gaze centre predicted by a lightweight MLP trained end-to-end with the navigation policy.

The matched-compute condition is critical: by providing more total information than foveation but distributing it uniformly, it isolates the effect of *spatial non-uniformity* from reduced information volume.

#### 3.3 Probing Methodology

After training, we freeze all weights and collect per-step LSTM hidden states (top-layer  $\mathbf{h}$ , 512-d) paired with ground-truth spatial labels from Gibson held-out validation episodes (the split not used in training). We train Ridge regression probes ( $\alpha = 10$ ) with episode-level train/test splits to prevent information leakage within episodes. Probe quality is validated with the Hewitt–Liang selectivity metric [Hewitt and Liang, 2019]: we compare probe accuracy on the true task to accuracy on a control task (random label permutation), confirming that high  $R^2$  reflects genuine spatial structure rather than probe memorisation.

Probed variables: GPS position, compass heading, distance-to-goal, path history (position at lag  $k \in \{0, \dots, 5\}$  from current hidden state), visited-region occupancy, full occupancy grid ( $20 \times 20$ ), per-unit spatial information (place-cell analysis), and the ego-frame goal vector  $\mathbf{R}(-\theta)(\mathbf{g} - \mathbf{p})$  decomposed into scalar goal-distance and 2-D direction.

**Sample-size and truncation policy** Four of the five conditions (blind, uniform, foveated-fixed, matched-compute) produce short evaluation episodes (mean  $\leq 5$  steps in the Gibson val split) because the agents navigate efficiently to goal. Foveated-learned produces much longer episodes (mean 171, median 77 steps) because its collapsed gaze yields inefficient but ultimately successful trajectories. This causes an  $\sim 100\times$  difference in the effective spatial range covered by the probe-training set per episode between foveated-learned and the other four conditions. We handle this with a uniform policy, applied consistently throughout §4:

- *Cross-condition quantitative comparisons* (§4.2 baseline probes, §4.4 CKA and probe transfer, §4.5 goal-vector and heading-content claims, §4.6 MP3D) use *truncated-to-first-4-steps* for foveated-learned and full episodes for the other four; the truncation matches the natural horizon of the other conditions. Tables and figures are labelled “truncated” where this applies.
- *Within-condition analyses where foveated-learned is not compared against others* (§4.6 per-condition place-cell counts, §4.6 multi-layer heatmap) use each condition’s full-episode data.
- *Behavioural interventions* (§4.4 transplant, §4.5 shortcut) use full rollouts for every condition because they measure on-policy task performance, which cannot be meaningfully truncated.
- *Path-history lag- $k$  probe* (§4.3) requires  $k + 1$  consecutive steps per episode; with  $k = 5$  the 4-step truncation would exclude every sample. Foveated-learned is therefore *excluded from H1* rather than truncated, and we discuss the absent data point separately.

For foveated-fixed, all analyses use checkpoint ckpt.36 at  $\sim 174\text{M}$  frames, the last clean intermediate before the silent-corruption window (Appendix B). A small number of our earliest probe tables were collected from a slightly earlier intermediate ( $\sim 154\text{M}$ ) on the same pre-corruption training trajectory; the two are within the training curve’s noise band, and we are re-running the probes at ckpt.36 for full consistency.

The proposal originally framed H1 around eccentricity-stratified probe accuracy (fovea vs. periphery). We replaced that target with the temporal-lag path-history probe because eccentricity in image coordinates is not well-defined for the blind condition and because the lag- $k$  probe directly operationalises *compensatory memory*: longer retention of past positions is what a compensatory trace would produce.

### 3.4 Cross-Condition Analysis

To test H2, we employ two complementary methods. *Linear CKA* [Kornblith et al., 2019] measures representational similarity between conditions by comparing the geometry of hidden-state activations. Values near 1 indicate identical geometry; near 0 indicates unrelated representations. *Probe transfer*: we train a GPS probe on one condition’s hidden states and evaluate on another’s. If spatial information is encoded in a shared linear subspace, cross-condition transfer should succeed. Catastrophic failure indicates fundamentally different encoding formats. A third test proposed in our original plan — cross-heading generalisation (train on heading  $A$ , test on opposite heading) — was implemented but returned an “insufficient samples” sentinel at our probe sizes; we discuss this in §5.5.

### 3.5 H3 Ablation: Gaze-Diversity Auxiliary Loss

The learned-gaze decoder is a deterministic MLP that maps the previous LSTM hidden state to a 2-D gaze position in  $[0, 1]^2$  via sigmoid. Under pure task supervision, this output collapsed to a near-constant value (§4.5). To test whether H3 can be recovered by explicitly discouraging collapse, we add a per-batch auxiliary loss:  $\mathcal{L}_{\text{div}} = \lambda \cdot \text{relu}(\sigma_{\text{target}}^2 - \text{Var}_{\text{batch}}[g])$ . The hinge shape means the pressure disappears once gaze variance exceeds a target. Pilot failure modes are reported in Appendix A.

Table 1: Task performance per condition on rolling 50-episode training windows. The “Probe ckpt” column gives the checkpoint used for probing and behavioural-probe results in §4.3–§4.6.

| Condition       | Frames | Success             | SPL                 | Probe ckpt     |
|-----------------|--------|---------------------|---------------------|----------------|
| Blind           | 342M   | 0.928               | 0.588               | 342M           |
| Uniform         | 250M   | <b>0.986</b>        | <b>0.852</b>        | 250M           |
| Foveated (fix)  | 174M   | <b>[TODO: 0.98]</b> | <b>[TODO: 0.83]</b> | 174M (ckpt.36) |
| Matched-compute | 250M   | 0.979               | 0.670               | 250M           |
| Fov-learned     | 250M   | 0.978               | 0.820               | 250M           |




Figure 2: Training curves read from tensorboard event files. All four sighted conditions converge to similar SPL (0.82–0.85) and success (> 97%) by ~100M frames; blind converges to lower SPL (~0.50) over a longer horizon. A short glitch at frame ≈ 228M (fov-learned) is a resume-boundary artefact in the scalar log.

### 3.6 Behavioural Probes

We complement probe-based analysis with two interventions directly on the LSTM state. *Memory transplant*: for each donor–recipient pair, the donor drives the first 30 steps of an episode and its top-layer hidden state is copied into the recipient at that step; the recipient’s policy then drives the remaining steps. We compare *baseline* (recipient alone, full episode), *self-transplant* (recipient as its own donor, controls for transplant mechanics), and *cross-transplant* (different donor). Midpoint 30 was chosen because Gibson PointNav episodes average ~100–120 steps; a later midpoint left too few steps after transplant. We use 150 episodes per pair per condition, sampled uniformly from the held-out Gibson validation split. *Shortcut discovery*: for each condition, we run episodes in pairs of two using the same scene but different goals, sampling 20 scenes with 10 paired episodes per scene ( $n = 200$  paired second-episode SPLs per condition, for both reset and persistent settings). We compare *reset* (LSTM zeroed between the two episodes) vs. *persistent* (LSTM state carried over) on the second-episode SPL. A reduction under the persistent setting measures how tightly the stored representation is locked to the previous goal. For the foveated-fixed condition, both behavioural probes use the ckpt.36 checkpoint (~174M frames, the last clean intermediate before the silent-corruption window, §4); the linear-probing results in §4.3–4.6 use a slightly earlier ~154M intermediate from the same pre-corruption window.

## 4 Results

### 4.1 Training Performance

All four sighted conditions converge to comparable task performance (97–99% success) at matched training budgets; blind reaches 93% at a longer horizon, reproducing Wijmans et al. [Wijmans et al., 2023]. This matters for everything that follows: differences in what the *hidden state* encodes cannot be attributed to differences in what the agents *do*. Two training-stability issues are disclosed in full in Appendix B: the foveated-fixed agent suffered a NaN-gradient corruption at ≈182M frames (we use ckpt.36 at ~174M, the last clean intermediate), and matched-compute converged with a degenerate  $1 \times 1$  feature-map in its encoder head (a valid information-volume control, since our H2 question is about the *encoded* volume).

Table 2: Global linear probe accuracy (Ridge,  $\alpha = 10$ , episode-level split,  $n = 500$  held-out Gibson evaluation episodes per condition). Per the sample-size and truncation policy in §3.3, foveated-learned is reported on the first 4 steps per episode to match the natural horizon of the other four conditions, whose episodes are short by nature of successful navigation (mean  $\leq 5$  steps). Foveated-fixed uses ckpt.36 (174M); others use the final converged checkpoint.

|                 | GPS $R^2$    | GPS MAE | Compass $R^2$ | Compass MAE |
|-----------------|--------------|---------|---------------|-------------|
| Blind           | <b>0.899</b> | 0.020 m | 0.878         | 0.8°        |
| Foveated        | 0.841        | 0.034 m | 0.715         | 2.1°        |
| Matched-compute | 0.832        | 0.033 m | 0.607         | 1.3°        |
| Uniform         | 0.614        | 0.035 m | 0.423         | 1.6°        |
| Fov-learned     | 0.031        | 0.022 m | <b>0.941</b>  | —           |




Figure 3: Global GPS probe  $R^2$  (solid) and Hewitt–Liang selectivity (hatched). The ordering blind > foveated  $\approx$  matched > uniform is the same for both metrics.

## 4.2 Baseline Probes: GPS, Compass, Distance-to-Goal

Every condition’s hidden state linearly encodes spatial quantities well above chance, and distance-to-goal decodes near-perfectly ( $R^2 > 0.95$ ) everywhere: the probe machinery works. Hewitt–Liang selectivity exceeds 0.76 in all sighted conditions, so the signals are genuine structure rather than probe memorisation. Two patterns visible in Table 2 pre-empt the three hypotheses that follow. First, among sighted fixed-gaze conditions, *foveated decodes GPS better than uniform* (0.84 vs. 0.61) — the opposite of what “more pixels in  $\Rightarrow$  better decoding out” would predict, and the hook for H1. Second, *foveated-learned* reaches the highest compass  $R^2$  of any condition (0.94, exceeding even the blind agent that receives compass directly) while its GPS decoding collapses to chance (0.03) — a categorical content shift we develop in H3.

## 4.3 H1: Sensor Structure $\rightarrow$ Compensatory Memory

**Finding** Non-uniform vision pushes memory to retain past positions longer. Five steps into the past, foveated hidden states still contain a linearly decodable location code (66% of their instant-decoding quality retained); uniform hidden states do not (probe fails to generalise at the same lag). In plain terms: only the foveated-fixed agent’s memory, five steps after visiting a point, still “knows where it was” well enough for a linear readout to reconstruct the coordinates.

**Ruling out alternatives** This is neither a probe nor a training-budget artefact. (a) Lag-0 decoding is comparable across conditions ( $R^2 \geq 0.61$ ), so probe capacity is not the bottleneck. (b) Hewitt–Liang label-permutation selectivity exceeds 0.76 at every lag, ruling out probe memorisation. (c) The foveated > uniform gap is present at every  $k$ , not a single-lag fluctuation; the curves do not cross. (d) All four conditions use the same held-out set, Ridge hyperparameters, and episode splits. Task




Figure 4: Path-history probe  $R^2$  vs. lag  $k$ . Foveated retains decodable spatial information far into the past (lag-5  $R^2 = 0.556$ ); uniform decays catastrophically (lag-5  $R^2 = -1.53$ ). Negative values indicate probe failure to generalise. Foveated-learned (orange) is flat at  $R^2 \approx -2.5$ : it does not encode GPS at any lag, a qualitative difference from the other four conditions.

Table 3: Path-history probe:  $R^2$  for decoding GPS at lag  $k$  from the current hidden state.

|          | lag-0 | lag-1 | lag-2 | lag-3 | lag-4        | lag-5        |
|----------|-------|-------|-------|-------|--------------|--------------|
| Blind    | 0.899 | 0.905 | 0.867 | 0.748 | 0.536        | 0.043        |
| Foveated | 0.841 | 0.853 | 0.824 | 0.737 | <b>0.629</b> | <b>0.556</b> |
| Matched  | 0.832 | 0.855 | 0.807 | 0.581 | 0.445        | 0.274        |
| Uniform  | 0.614 | 0.595 | 0.534 | 0.386 | -0.105       | -1.532       |

competence also fails as an explanation: uniform achieves the highest task SPL of all five conditions (0.852) yet has the worst retention.

**Why: non-uniformity, not volume** Two contrasts localise the mechanism. Foveated  $>$  uniform even though both see the scene: a full-resolution stream removes the *need* for memory-based position, because each frame re-derives it. Foveated  $>$  matched-compute at every lag  $k \geq 3$  even though matched-compute delivers more total pixels: the difference is that matched distributes its pixels uniformly. Compensatory retention tracks the *non-uniformity* of the sensor, not the quantity of input. The signature is consistent with belief-state-like representations emerging from pure RL under partial observability [Hennig et al., 2023].

**Fov-learned exclusion** The lag- $k$  probe needs  $k + 1$  consecutive within-episode steps, which the 4-step truncation used for cross-condition probing (§3.3) removes from foveated-learned. Its full-episode lag- $k$  numbers are not comparable to the other four (they drift into the very large spatial range of its long wandering episodes, §4.4 caveat), so we neither claim nor deny H1 for foveated-learned. Its content signature — heading-dominant rather than position-dominant — is a different effect addressed in H3.

#### 4.4 H2: Sensor Structure $\rightarrow$ Format-Level Divergence

**Finding** Five agents that solve the same task at the same success rate do not share a spatial representation. Pairwise linear similarity (CKA) sits near zero for every condition pair (Figure 5); a GPS probe trained on one condition’s hidden states and evaluated on another’s does not generalise at all (Table 4; diagonal 0.90–0.97, off-diagonals all  $R^2 \leq -2.7$ ). In plain terms: what each agent calls "position" is written in a different linear subspace. Visual input type shapes *format*, not just strength.

**Ruling out alternatives** Self-probes succeed ( $R^2 > 0.90$ ) for every condition, so probe capacity is not the bottleneck. CKA is computed on a common-sample truncation, ruling out sample-size asymmetry. Task competence is tightly bunched (97–99% success), so the divergence is not a skill




Figure 5: Unaligned linear CKA [Kornblith et al., 2019] between top-layer LSTM hidden states ( $n = 1,907$ ). Every off-diagonal is below 0.004; the most-similar pair is the two foveated variants (fixed vs. learned gaze) at 0.0035.

Table 4: Probe-transfer  $R^2$  on GPS: train on row condition, test on column. Self-accuracy diagonal 0.90–0.97; every cross-condition transfer is catastrophically negative.

| Train \ Test    | Blind       | Uniform     | Fov-fix     | Fov-lrn     | Matched-compute |
|-----------------|-------------|-------------|-------------|-------------|-----------------|
| Blind           | <b>0.97</b> | −9.7        | −2.7        | −613        | −15.7           |
| Uniform         | −7.3        | <b>0.97</b> | −13.5       | −410        | −11.5           |
| Fov-fix         | −11.3       | −5.1        | <b>0.96</b> | −232        | −10.7           |
| Fov-lrn         | −482        | −430        | −456        | <b>0.90</b> | −855            |
| Matched-compute | −5.7        | −11.7       | −15.5       | −390        | <b>0.96</b>     |

difference. And transfer fails symmetrically (both directions), which rules out the possibility that one condition’s code simply embeds another’s.

**Foveated-learned caveat** Foveated-learned’s transfer values are one-to-two orders of magnitude worse than other pairs because its average episode is  $\sim 40\times$  longer, spreading its probe data over a much larger spatial range (§3.3). On a first-4-steps-per-episode subset matched to the others, its cross-transfer falls into the  $[-2, -15]$  band of every other pair — the qualitative conclusion (no shared code) is unchanged. t-SNE of pooled hidden states (Figure 6) confirms visually: the five conditions occupy non-overlapping regions of the 2-D manifold.

**Boundaries** Our similarity tests are linear (CKA) or near-linear (t-SNE default); non-linear similarity measures could reveal higher-order shared structure. Our five sensors are diverse but not exhaustive; hybrid sensors (coarse uniform periphery + sharp fovea, depth-only inputs) might produce intermediate geometries.

**Memory transplant** Probe-based similarity measures could still miss *functional* equivalence: two representations with low CKA might still be interchangeable if the downstream policy reads them the same way. We test this with memory transplant — paste a donor’s mid-episode hidden state into the recipient’s policy at step  $t$ , let the recipient drive the rest of the episode. The isolated condition-mismatch cost (cross-transplant SPL − self-transplant SPL, which subtracts both a small  $t=0$  protocol-noise floor and the rollout-divergence cost of re-initialising a recurrent state mid-episode) is stable across midpoints  $t \geq 15$  and carries a clear ranking (Figure 7b): foveated-learned→foveated incurs the largest drop (−0.17 SPL), foveated→uniform a smaller one (−0.10), and foveated→blind is essentially neutral (−0.02). The striking piece is the foveated-learned→foveated drop: this is the pair with the *highest* pairwise CKA (0.0035) of any pair we tested, yet it is the most behaviourally incompatible. Higher probe-based similarity does not imply higher functional interchangeability — the representations of agents that share the same sensor transform but differ in gaze policy are more compatible to look at than to use. (CKA and probe transfer use common-sample truncation for fair cross-condition comparison; transplant uses full on-policy rollouts because the task is not




Figure 6: t-SNE projection of top-layer LSTM hidden states, pooled across all five conditions (1,500 random steps per condition; perplexity 30). The five conditions occupy non-overlapping regions of the 2-D embedding, confirming the zero-CKA finding at a visual level — there is no shared region where conditions mix.




Figure 7: Memory transplant across a midpoint sweep  $\{0, 15, 30, 60, 90\}$  steps ( $n = 100\text{--}150$  episodes/cell). (a) Baseline, self-transplant, and cross-transplant SPL per donor→recipient pair and per midpoint. At  $t = 0$  the donor has accumulated no information yet, so self- and cross-transplant reduce to a protocol-noise floor (self  $\approx -0.03$ , cross  $\approx -0.06$  SPL vs. baseline). Self-transplant SPL drops monotonically with increasing midpoint: the recipient’s own past recurrent state genuinely diverges from the state it would have occupied had it driven the rollout from step 0. (b) Cross minus self (the isolated condition-mismatch component, with both protocol noise and rollout divergence subtracted) is roughly stable across midpoints  $t \geq 15$ : foveated-learned→foveated is the most behaviourally incompatible pair at every midpoint tested ( $-0.17$  SPL on average), *despite having the highest pairwise CKA* (0.0035) of any condition pair—higher probe-based similarity does not predict higher functional interchangeability.

truncatable. The two measurement scopes are deliberate, and their conclusions are consistent.) Together with CKA, probe transfer, and t-SNE, transplant provides a fourth convergent signal that the five representations are format-level distinct.

#### 4.5 H3: Gaze Policy as an Orthogonal Content Lever

H1 and H2 identify sensor structure as one axis on representational content. H3 asks whether the gaze policy that directs the sensor is a *second*, dissociable axis. We test two progressively weaker forms: a strong *epistemic* form (learned gaze should actively target high-uncertainty regions) and




Figure 8: Ridge probe  $R^2$  for three goal-related targets. Goal distance is linearly decodable in every condition ( $R^2 \geq 0.78$ ); goal *direction* and the full 2-D ego-to-goal vector are at chance or negative for every condition, including the ones most hurt by the shortcut manipulation below. No condition stores the goal vector; what differs across conditions is the integrated trajectory and scene-history content.

a weaker *content-steering* form (the gaze-policy-induced sampling statistics should change what is encoded even if the gaze itself does not move).

**Strong form: not supported** The epistemic form needs a non-trivially moving gaze, which our implementation does not produce. After 250M frames the gaze decoder collapses to a near-constant point (0.49, 0.62) with per-dimension std  $< 10^{-3}$ : functionally a fixed-gaze agent. Two anti-collapse regularisers (Appendix A) restore gaze variance only at the cost of task performance. A single-seed collapse is not definitive, and a multi-seed and alternative-decoder follow-up is in progress; what we can say so far is that implicit task pressure alone did not produce actively moving gaze in our setting.

**Content-steering form: supported** Fov-learned has the same foveation transform as foveated-fixed and matches its task-success range (SPL 0.820, 97.8%), yet its representation encodes qualitatively different content: the highest compass decoding of any condition ( $R^2 = 0.941$ , exceeding even blind which receives compass directly as input) and near-chance GPS decoding ( $R^2 = 0.031$ ). The dissociation holds at every LSTM layer (Figure 10), so it is not a top-layer probing artefact. In plain terms: the gaze policy that happens to collapse sends the LSTM a near-constant view of one image patch, which reliably carries orientation cues (what wall am I facing?) but not much position information (which room am I in?); the memory specialises accordingly. Gaze policy is a representational-content lever even when it does not move.

**Mechanism** It would be natural to read fov-learned as "heading-only because it does not store the goal vector." A goal-vector probe (Figure 8) rejects this reading in a revealing way: *no* condition stores the ego-to-goal direction linearly, because the PointGoal task provides it as a per-step input and the LSTM has no reason to memorise it. What each LSTM memorises is its own condition-specific trajectory and scene-history features. The heading-vs-position dissociation in fov-learned is therefore not a gap in its memory; it is a specialisation: the foveated-centre patch under a locked gaze preferentially feeds heading cues into the recurrent dynamics, and the full-visual stream under uniform gaze preferentially feeds position cues. Same architecture, same task, same sensor transform, different gaze  $\Rightarrow$  different content.

**Shortcut discovery** The content-steering claim makes a testable prediction: heading-dominant representations should be more robust to goal changes than position-dominant representations, because a heading estimate stays useful when the goal moves while a position-to-goal mapping does not. We test this by running paired episodes in the same scene with fresh goals, comparing second-episode SPL under *reset* vs. *persistent* memory (20 scenes  $\times$  10 paired episodes = 200 paired SPLs per condition). Persistent memory hurts every condition — the agent is primed by its previous trajectory — but by very different amounts (Figure 9): foveated-learned loses the smallest fraction of its baseline SPL (15%), uniform and blind the largest (41–52%), with foveated-fixed and matched-compute in between. Critically, this ordering is *not* what a "stronger memory  $\Rightarrow$  more locked-in" story would predict — uniform and foveated-fixed both decode GPS at  $R^2 \approx 0.95$  at the




Figure 9: Shortcut discovery ( $n = 200$  second-episode SPLs per bar: 20 Gibson scenes  $\times$  10 paired episodes). Persistent recurrent memory across episodes (new goal, same scene) vs. reset memory. Conditions are ordered by relative SPL drop (bold  $\Delta$ SPL). Heading-dominant fove-learned is least hurt (15%); position-dominant uniform and blind are most hurt (41–52%). The ordering confirms the content-steering form of H3: the learned-gaze representation is functionally different from the fixed-gaze representation, not just probe-different.




Figure 10: *Left* (Fig. 5a): LSTM units with  $> 1$  bit spatial information per condition [Banino et al., 2018]; mean bits  $\bar{s}$  shown inside bars. Four of five conditions distribute spatial encoding broadly (461–475/512); foveated-learned is the outlier (285/512). *Right* (Fig. 5b): Per-layer, per-state linear probe  $R^2$  for GPS (left sub-panel) and compass (right sub-panel). Foveated-learned’s compass- $\rightarrow$ -position dissociation is present at every layer, not just the top.

top layer, yet foveated-fixed loses half as much. It is what the heading-dominant content prediction gives us: the agent whose memory is oriented around what it is facing rather than where it is carries over best. The shortcut ordering therefore confirms H3’s content-steering form behaviourally.

#### 4.6 Additional Analyses

**Controls** Two probes fail to discriminate between conditions and so define the boundary of our claims. Coarse occupancy (“have I visited this region?”) is trivially decodable everywhere ( $R^2 \in [0.80, 0.83]$ ), so one-bit spatial working memory is ubiquitous and not part of the H1–H3 story. Metric  $20 \times 20$  occupancy reconstruction fails everywhere ( $R^2 < 0$ ), consistent with the SPACE benchmark’s finding that fine-grained layout requires non-linear probes or longer horizons [Ramakrishnan et al., 2025]: our claims are about representational *format* and *content*, not about high-resolution metric layout.

**Place-cell-like units and multi-layer depth** Two further analyses confirm that our findings are about representational *content*, not about where we probe from or how we probe (Figure 10). First, four of five conditions distribute spatial encoding broadly across LSTM units in the Banino et al. [Banino et al., 2018] sense ( $> 90\%$  of units carry  $> 1$  bit of spatial information); foveated-learned is the clear outlier (56% of units), consistent with its specialisation away from position. Second, foveated-learned’s heading  $\gg$  position dissociation is present at *every* LSTM layer — this is a property of the representation, not of a linear probe on one vantage point.




Figure 11: Probe generalisation from Gibson evaluation (unshaded) to held-out MP3D evaluation (hatched), 300 episodes/condition. Same trained checkpoints; no re-training. Across GPS, compass, and distance-to-goal, four of five conditions show  $|\Delta R^2| \leq 0.17$ , confirming that our H1–H3 findings are not Gibson-specific. Foveated-learned’s compass drop ( $-0.56$ ) is the one exception; the heading  $>$  position ordering is nonetheless preserved within MP3D.

**Generalisation to MP3D** To test whether our findings are Gibson-specific, we re-ran the probing pipeline on MP3D validation scenes [Savva et al., 2019]: larger and more complex apartments than Gibson, same trained checkpoints, no re-training (Figure 11). The probe signatures hold. Four of five conditions show  $|\Delta R^2| \leq 0.17$  on every target, keeping the qualitative ordering intact. Foveated-learned’s compass  $R^2$  drops from 0.94 (Gibson, matched-protocol) to 0.38 on MP3D — the largest dataset-transfer drop of any condition on any target, and plausibly because its collapsed gaze was fit to Gibson-specific image statistics. Even at these degraded absolute values, however, fov-learned’s heading-over-position ordering survives (compass 0.38  $\gg$  GPS  $-0.21$ ): the qualitative H3 content shift is dataset-robust where the quantitative magnitude is not. This is the pattern we would have hoped for: our strongest claims (about representational format and about qualitative content shifts) transfer across datasets; the one absolute-magnitude sensitivity we find localises cleanly onto the fragile gaze-collapse mechanism that H3 itself flags as dataset-dependent.

## 5 Discussion

### 5.1 Two orthogonal axes: sensor structure and gaze policy

Our three findings are unified under a two-axes picture. *Sensor structure* (tested by H1 and H2) controls the spatial distribution of information each frame delivers: when current perception is reliable and spatially uniform, the recurrent memory does not need to maintain spatial information internally because each frame re-derives it; when current perception is absent (blind) or spatially non-uniform (foveated-fixed), re-derivation fails and memory must carry the content instead. Retention is monotone in frame reliability (uniform  $<$  matched-compute  $<$  blind  $<$  foveated-fixed, with foveated-fixed’s memory the longest-lived of all) and, independently, the representational *format* diverges across conditions: no two share a common linear subspace, and mid-episode hidden states do not transfer even between the two foveated variants. *Gaze policy* (tested by H3) appears to act as a second axis, dissociable from the first: foveated-learned has the same foveation transform as the fixed-gaze variant, yet a near-constant learned gaze produces a near-constant view of one image region, preferentially carrying heading cues over position cues, and the representation specialises accordingly (compass  $R^2 = 0.941$ , GPS  $R^2 = 0.031$ ). In the conditions we tested, the two axes steer the *format* of learned spatial memory, not merely its strength; we note that the axes’ full independence would require testing gaze-policy variants on uniform-vision and matched-compute sensors, which is beyond our scope. Neither axis is visible if one studies only blind or uniform-sighted agents, which is why both have remained implicit in the emergent-map literature.

## 5.2 Why this, not alternative accounts

We rule out four alternative explanations. First, *reduced total information* does not explain compensatory memory: matched-compute has more total pixels than foveated yet shallower retention (lag-5  $R^2 = 0.27$  vs. 0.56). Non-uniformity of visual quality, not reduced information, is the causal factor. Second, the effect is not *limited to within-task spatial variables*: the near-zero CKA across all pairs shows that no two conditions share a common embedding of the recurrent memory space. Third, the learned-gaze heading->-position dissociation is not an *artefact of episode-length disparity*: on a matched-distribution subset (first 4 steps per episode) the dissociation persists, and it appears at every layer (Figure 10). Fourth, the divergence is not a *linear-probe artefact*: the behavioural interventions (memory transplant in §4.4, shortcut discovery in §4.5) go outside the probe framework and give the same ordering, with cross-condition transplants dropping SPL by 0.19–0.21 between the most probe-similar pair (the two foveated variants).

## 5.3 Biological precedent as convergent evidence

Each of our three findings has a parallel in the animal-navigation literature, which we take as one piece of convergent (not conclusive) evidence that the effects are not specific to a single LSTM architecture. Compensatory memory under reduced sensory access is documented in congenitally blind humans, who develop enhanced hippocampal coupling and outperform sighted controls on non-visual spatial-working-memory tasks [Kupers and Ptito, 2014], and in rodents whose grid-cell periodicity collapses in darkness despite intact head-direction inputs [Chen et al., 2016]—both format-level, not just strength-level, effects. Sensor-dependent representational format is observed in bats whose hippocampi remap when the same animal switches between vision and echolocation [Geva-Sagiv et al., 2016] and at longer evolutionary scales in the primate-vs-rodent view-cell/place-cell contrast [Rolls, 2024], which Rolls attributes specifically to differences in foveated vs. near-panoramic sampling statistics. Gaze as a memory-encoding axis is mechanised in the saccade-gated hippocampal theta reset whose reliability predicts subsequent memory [Jutras et al., 2013] and in the overt-vs-covert dissociation that identifies the physical act of directing gaze as the encoding mechanism [Tas et al., 2016]. We do not claim our agents implement these biological mechanisms, nor that the parallel structure settles the generality question; we do claim that the parallels make a pure single-architecture artefact less likely, and that the controlled agent setting complements biological comparisons where sensor type, architecture, and training objective cannot be independently manipulated.

## 5.4 Implications for ML practice and architectural scope

Within the recurrent PointNav setting we study, three concrete implications follow. (i) *Sensor design is a tunable representational axis*: the spatial distribution of visual quality predictably shapes the format of the learned spatial code, so foveation-like structures may be useful not only for compute efficiency but for inducing richer temporal memory in downstream recurrent policies. (ii) *Richer input does not imply richer representation*: adding pixels uniformly (matched-compute) does not recreate the temporal-memory signature produced by non-uniform input — “more input” is not a valid substitute for “structured input.” (iii) *Attention policy is not a task-performance-only choice*: even a trivially collapsed gaze that preserves task success (97.8%) can reshape what the downstream representation encodes (heading  $\gg$  position), in ways that matter for interpretability, transfer, and robustness.

**Beyond the recurrent setting (untested)** We use a Wijmans-style recurrent (LSTM) backbone for direct comparability with the emergent-map literature. The three implications above are established only in this setting, but we speculate about where they plausibly extend. Transformer-based embodied agents build their internal state by accumulating information via learned attention, which is structurally similar to our gaze policy in the sense that it selects which observation content is read into the representation; our H3 finding therefore suggests a testable prediction for transformer navigators—strong constraints on attended tokens should produce content-level dissociations of the sort we observe—but the prediction has not been verified and is left for future work. The sensor-structure axis (H1, H2) is a claim specifically about RL recurrent navigation agents trained under partial observability; whether equivalent representational-content effects arise in supervised visual learners, non-visual modalities, or multimodal systems is an open empirical question that the present experiments do not resolve.

## 5.5 Limitations

(i) *Intermediate checkpoint for foveated-fixed*: all foveated-fixed analyses use ckpt.36 at  $\sim 174\text{M}$ , the last checkpoint before the silent-corruption window (Appendix B); the other four conditions use the final converged checkpoint. (ii) *Single-seed H3*: our H3 observations use one seed and one gaze-decoder architecture, so we cannot distinguish a fundamental limit of implicit task pressure from a single-seed local minimum or sigmoid saturation. (iii) *Cross-heading probe*: the proposal included a cross-heading generalisation probe (train on heading  $A$ , test on opposite heading). Our implementation returned “insufficient samples” sentinels at our probe sizes and does not appear in our results. (iv) *Foveation model*: Gaussian blur with quadratic falloff is a simplification; log-polar or multi-scale pyramid models may behave differently. (v) *Linear probes only*: non-linear probes might reveal structure Ridge regression misses, particularly for occupancy maps. (vi) *Architecture scope*: we use a recurrent (LSTM) backbone to remain comparable with the emergent-map literature we build on. The core claims (sensor structure and gaze policy as representational-content axes) should hold for transformer-based navigation agents whose internal state is built by learned attention, but directly verifying this requires re-running the five-condition ablation on a transformer backbone and is left for future work; see §5 for the specific predictions.

## 6 Conclusion

We ran a controlled five-condition ablation that isolates the effect of visual-input structure on the representations learned by an otherwise-identical DD-PPO navigation agent. Three findings emerge. (1) *Compensatory memory*. Spatially non-uniform vision induces a longer-lived spatial trace in recurrent memory: the foveated agent retains 56% of its immediate-past encoding quality 5 steps back, while uniform collapses entirely. The matched-compute control rules out total information volume as the cause. (2) *Sensor-dependent representational format*. No two of our five conditions share a representational geometry ( $\text{CKA} < 0.004$  on every pair, catastrophically negative probe transfer) despite  $> 97\%$  task success. Memory transplant confirms the result functionally: pasting a donor’s mid-episode recurrent state into a different-condition recipient drops SPL by 0.19–0.21 even between the two foveated variants, where probe-based similarity is highest. (3) *Gaze-locked representational content*. When the learned-gaze decoder collapses to a near-constant output, memory encodes heading strongly ( $R^2 = 0.94$ ) at the expense of position ( $R^2 = 0.03$ ) — a categorical content shift, not a graded fidelity loss. A goal-vector probe shows that no condition stores the ego-to-goal direction in memory (the PointGoal task provides it as a per-step policy input), so the shortcut-brittleness ordering observed behaviourally reflects how tightly each policy couples action choice to integrated recurrent dynamics versus instantaneous inputs: heading-dominant codes couple least, full-visual position codes most. Each of these three findings has a parallel in the animal-navigation literature [Kupers and Ptito, 2014, Chen et al., 2016, Geva-Sagiv et al., 2016, Rolls, 2024, Jutras et al., 2013], which we read as convergent evidence (not proof) that the effects are not pure artefacts of our specific LSTM. Within the setting we tested—recurrent PointNav agents with five matched-protocol sensor/gaze configurations—the takeaway is that sensor structure and gaze policy are tunable axes along which the *format* of learned spatial memory can be steered; both are invisible under the uniform-vision assumption of prior work. Open questions include whether a genuinely epistemic gaze (requiring stochastic or attention-based architectures) recovers or reverses the heading-vs-position dissociation; whether the same five-condition ablation on a transformer backbone reproduces or refines our findings; whether the effects persist at longer navigation horizons and with other sensor topologies; and whether analogous representational-content effects arise in non-navigation tasks or non-visual modalities.

## References

- Andrei Atanov, Yicheng Fu, Gaurav Singh, Jianren Yu, Andrew Spielberg, and Amir Zamir. How far can a 1-pixel camera go? solving vision tasks using photoreceptors and computationally designed visual morphology. In *European Conference on Computer Vision (ECCV)*, 2024.
- Andrea Banino, Caswell Barry, Benigno Uria, Charles Blundell, Timothy Lillicrap, Piotr Mirowski, Alexander Pritzel, Martin J Chadwick, Thomas Degris, Joseph Modayil, et al. Vector-based navigation using grid-like representations in artificial agents. *Nature*, 557(7705):429–433, 2018.

- Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics*, 48(1):207–219, 2022.
- Guifen Chen, John A. King, Neil Burgess, and John O’Keefe. Absence of visual input results in the disruption of grid cell firing in the mouse. *Current Biology*, 26(17):2335–2342, 2016.
- Christopher J Cueva and Xue-Xin Wei. Emergence of grid-like representations by training recurrent neural networks to perform spatial localization. In *International Conference on Learning Representations (ICLR)*, 2018.
- Kshitij Dwivedi, Gemma Roig, Aniruddha Kembhavi, and Roozbeh Mottaghi. What do navigation agents learn about their environment? In *Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Maya Geva-Sagiv, Sandro Romani, Liora Las, and Nachum Ulanovsky. Hippocampal global remapping for different sensory modalities in flying bats. *Nature Neuroscience*, 19(7):952–958, 2016.
- James Gornet and Matt Thomson. Automated construction of cognitive maps with visual predictive coding. *Nature Machine Intelligence*, 6:820–833, 2024.
- Jacqueline Gottlieb, Pierre-Yves Oudeyer, Manuel Lopes, and Adrien Baranes. Information-seeking, curiosity, and attention: computational and neural mechanisms. *Trends in Cognitive Sciences*, 17(11):585–593, 2013.
- Mary M. Hayhoe and Jonathan S. Matthis. Control of gaze in natural environments: effects of rewards and costs, uncertainty and memory in target selection. *Philosophical Transactions of the Royal Society B*, 373(1718):20170368, 2018.
- Jay A Hennig, Sandra A Romero Pinto, Takahiro Yamaguchi, Scott W Linderman, Naoshige Uchida, and Samuel J Gershman. Emergence of belief-like representations through reinforcement learning. *PLOS Computational Biology*, 19(9):e1011067, 2023.
- John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2019.
- Michael J. Jutras, Pascal Fries, and Elizabeth A. Buffalo. Oscillatory activity in the monkey hippocampus during visual exploration and memory formation. *Proceedings of the National Academy of Sciences*, 110(32):13144–13149, 2013.
- Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *International Conference on Machine Learning (ICML)*, 2019.
- Ron Kupers and Maurice Ptito. Compensatory plasticity and cross-modal reorganization following early visual deprivation. *Neuroscience & Biobehavioral Reviews*, 41:36–52, 2014.
- Volodymyr Mnih, Nicolas Heess, Alex Graves, and Koray Kavukcuoglu. Recurrent models of visual attention. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2014.
- Jiri Najemnik and Wilson S. Geisler. Optimal eye movement strategies in visual search. *Nature*, 434(7031):387–391, 2005.
- Motahareh Pourrahimi and Pouya Bashivan. Emergent brain-like representations in a goal-directed neural network model of visual search. *bioRxiv preprint*, 2025.
- Santhosh Kumar Ramakrishnan, Erik Wijmans, Philipp Krähenbühl, and Vladlen Koltun. Does spatial cognition emerge in frontier models? In *International Conference on Learning Representations (ICLR)*, 2025.
- Edmund T. Rolls. Hippocampal spatial view cells for memory and navigation, and their underlying connectivity in humans. *Hippocampus*, 33(6):533–572, 2023.
- Edmund T. Rolls. Why do primates have view cells instead of place cells? *Trends in Cognitive Sciences*, 28(9):862–874, 2024.

- Ruth Rosenholtz. Capabilities and limitations of peripheral vision. *Annual Review of Vision Science*, 2:437–457, 2016.
- Manolis Savva, Abhishek Kadian, Oleksandr Maksymets, Yili Zhao, Erik Wijmans, Bhavana Jain, Julian Straub, Jia Liu, Vladlen Koltun, Jitendra Malik, Dhruv Batra, and Devi Parikh. Habitat: A platform for embodied AI research. In *International Conference on Computer Vision (ICCV)*, 2019.
- Jinghuan Shang and Michael S. Ryoo. Active vision reinforcement learning under limited visual observability. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Nathan Sprague and Dana H. Ballard. Eye movements for reward maximization. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2003.
- Hans Strasburger, Ingo Rentschler, and Martin Jüttner. Peripheral vision and pattern recognition: A review. *Journal of Vision*, 11(5):13, 2011.
- Amanda C. Tas, Steven J. Luck, and Andrew Hollingworth. The role of attention in the encoding of visual working memory. *Journal of Experimental Psychology: Human Perception and Performance*, 42(8):1121–1138, 2016.
- Erik Wijmans, Abhishek Kadian, Ari Morcos, Stefan Lee, Irfan Essa, Devi Parikh, Dhruv Batra, and Manolis Savva. DD-PPO: Learning near-perfect pointgoal navigators from 2.5 billion frames. In *International Conference on Learning Representations (ICLR)*, 2020.
- Erik Wijmans, Manolis Savva, Irfan Essa, Stefan Lee, Ari Morcos, and Dhruv Batra. Emergence of maps in the memories of blind navigation agents. In *International Conference on Learning Representations (ICLR)*, 2023.
- Sylvia Wirth, Pierre Baraduc, Angélique Planté, Serge Pinede, and Jean-René Duhamel. Gaze-informed, task-situated representation of space in primate hippocampus during virtual navigation. *PLOS Biology*, 15(2):e2001045, 2017.
- Fei Xia, Amir R Zamir, Zhiyang He, Alexander Sax, Jitendra Malik, and Silvio Savarese. Gibson Env: Real-world perception for embodied agents. In *Computer Vision and Pattern Recognition (CVPR)*, 2018.

## NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

**The checklist answers are an integral part of your paper submission.** They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer:  
answerYes

Justification: The abstract and §1 state the three hypotheses H1/H2/H3 with explicit signatures (what pattern supports each), and preview the three findings with specific magnitudes (path-history  $R^2=0.56$ , CKA $<0.004$ , compass  $R^2=0.94$  vs GPS  $R^2=0.03$ ). These are expanded in §4 and §5 with no stronger claims than the data support.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer:

answerYes

Justification: §5.4 explicitly lists six limitations (i–vi) covering checkpoint staleness, single-seed H3, missing proposed analyses, foveation-model simplification, linear-probes-only, and LSTM-architecture scope.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer:

answerNA

Justification: The paper reports empirical results; it contains no theoretical claims or proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer:  
answerYes

Justification: §3.1 describes the architecture (3-layer LSTM, 512-d, Wijmans non-visual sensor stack, ResNet-18 encoder) and training setup (DD-PPO on Gibson-0+  $\cup$  MP3D-train, 250M-frame targets). §3.2–§3.6 describe all probing and cross-condition protocols. Appendix A details the gaze-diversity ablation. Code will be released upon publication.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

## 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer:  
answerYes

Justification: §3.7 commits to open-source release of the training and probing code. Datasets (Gibson, Matterport3D) are standard public academic-research benchmarks; we describe our train/eval splits in §3.1.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).

- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer:

answerYes

Justification: §3.1 specifies training pool and eval set; §3.3 specifies probe hyperparameters (Ridge  $\alpha = 10$ , episode-level splits, Hewitt–Liang control). Appendix A gives the ablation’s loss formulation and coefficient.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer:

answerNo

Justification: We report single-seed numbers per condition. This is documented as Limitation (ii) in §5.4. Multi-seed replication for tighter confidence intervals is in the immediate follow-up.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).

- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer:

answerYes

Justification: §3.7 reports per-condition training budget (5–7 days on 2×V100 for 250M frames) and probing budget (<1h per condition on one V100).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer:

answerYes

Justification: The research uses public academic-research benchmarks, publicly-released open-source RL frameworks, and involves no human subjects or deployment. We also report and remediate a numerical-stability bug upstream (Appendix B).

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer:

answerYes

Justification: §5.3 discusses implications — sensor design as a tunable axis for interpretable embodied agents. Negative societal impacts are minimal: the work is fundamental research on simulated indoor navigation with no direct deployment pathway.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer:

answerNA

Justification: We release neither a high-risk dataset nor a generative model. Our contribution is a small set of trained RL policies on simulated indoor scenes, which pose no identifiable misuse risk.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer:

answerYes

Justification: Gibson Database and Matterport3D are used under their released academic-research licenses; Habitat-Sim and habitat-lab are MIT-licensed. Our use complies with their terms; citations §3.1.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.

- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

### 13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer:

answerYes

Justification: Code will be released with documentation upon publication. Figures are novel; they are released under the paper's license alongside the code. The gaze-diversity auxiliary loss and NaN-sanitisation patch are small standalone modules with docstrings.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

### 14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer:

answerNA

Justification: The paper involves no crowdsourcing and no human-subjects data.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

### 15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer:

answerNA

Justification: No human-subjects research; no IRB/equivalent review was required.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

#### 16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer:

answerNA

Justification: LLMs were used only for writing assistance on the manuscript. They are not part of the core research methodology; all empirical results come from standard RL training and linear probing.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.

## A Gaze-Diversity Auxiliary Loss (H3 Ablation)

**Motivation** In the first foveated-learned training run the gaze decoder collapsed (per-dimension  $\text{std} \approx 5 \times 10^{-4}$ ). We attribute this to two mechanisms: (i) once the sigmoid output saturates, its derivative  $\sigma'(x) \rightarrow 0$  attenuates gradient signal; (ii) consistent-gaze views yield more predictable visual features, so the reward-gradient path has no explicit incentive for variety.

**Target-variance hinge** We add  $\mathcal{L}_{\text{div}} = \lambda \cdot \text{relu}(\sigma_{\text{target}}^2 - \text{Var}_{\text{batch}}[g])$  where  $g \in [0, 1]^2$ ,  $\text{Var}_{\text{batch}}$  averaged over the two gaze dimensions,  $\sigma_{\text{target}}^2 = 0.01$  ( $\text{std} \approx 0.1$ ),  $\lambda = 0.01$ . The hinge disappears once variance exceeds the target. We also evaluated an un-capped  $-\lambda \cdot \text{Var}[g]$  variant as a negative control.

**Pilot 1 (uncapped, job 2840256)** After 10 M frames, gaze variance exceeded that of a uniform distribution over  $[0, 1]^2$  ( $\text{std} [0.43, 0.42]$  vs. theoretical max 0.289): gaze was essentially random. Task performance collapsed: SPL plateaued at 0.05, running-mean metrics transitioned to NaN at  $\sim 8$  M frames. The uncapped regulariser destroyed the consistent-foveation signal the policy needs.

**Pilot 2 (hinge, job 2842288)** Gaze converged to an extreme corner of image space: mean (0.001, 0.998) with  $\text{std} [0.020, 0.023]$  — well below the target  $\text{std}$  of 0.1 and therefore with the hinge active throughout. Task SPL plateaued at 0.05 as in pilot 1. The policy apparently found the minimum-task-cost gaze that partially satisfies the diversity pressure: a sigmoid saturated at one corner, producing a constant image-corner view.

**Interpretation** Both pilots argue that in our deterministic sigmoid-MLP gaze architecture with slow-gaze rollouts, gaze collapse is a pre-requisite for task learning rather than an accident. Task signal alone does not spontaneously produce informative fixation, and an exogenous anti-collapse regulariser moves gaze toward whatever region reduces regulariser loss fastest, which is generically *not* the task-useful region. A genuine H3 test likely requires (i) a stochastic gaze policy with entropy bonus, so diversity is sampled at rollout time rather than enforced on a deterministic output; or (ii) a task-aligned intrinsic reward (e.g., information gain about goal-relative position); or (iii) an architectural change separating gaze generation from visual-feature extraction (e.g., attention-based gaze).

## B Training Stability: Silent Weight Corruption in DD-PPO

During an initial foveated training run we observed a failure mode that did not manifest as a crash but as slow, silent corruption of the learned weights. After  $\approx 170$ M frames of stable progress (reward  $\approx 10$ , SPL  $\approx 0.83$ ), one mini-batch on one DDP worker produced a NaN gradient. Because `torch.nn.utils.clip_grad_norm_` does not sanitise non-finite gradients—the global norm becomes NaN, the clip coefficient becomes NaN, and `grad *= NaN` remains NaN—the subsequent `optimizer.step()` wrote NaN into every parameter. DDP all-reduce propagated the corruption across workers. The training loop did not raise: it continued for 2.5 days with reward=NaN and SPL=0 until `torch.multinomial` finally rejected the NaN probability tensor. All checkpoints saved during this window were unrecoverable (90 of 97 parameter tensors entirely NaN).

To isolate the origin of the NaN gradient we re-ran training from the last clean checkpoint with `torch.autograd.set_detect_anomaly(True)` and per-module forward-output hooks. The anomaly traceback pinned the failure to `MseLossBackward0`, i.e. the backward pass of the value-head loss  $\frac{1}{2}(V(s) - R)^2$ ; rare return-vs-value discrepancies overflowed float32 in the squaring step, producing non-finite gradients downstream.

Our mitigation is a minimally invasive patch to `PPO.before_step` that zeroes any non-finite gradient element before clipping. A bad mini-batch becomes a no-op update rather than a catastrophic one. The patch is a no-op on clean training paths. Across the five conditions reported here, no NaN events were observed during full training after deployment, and all final checkpoints contain only finite weights. We reported the underlying issue upstream (habitat-lab issue #2226) with a minimal reproduction and a proposed native fix.