
How Foveated Vision Shapes Cognitive Maps in Navigation Agents

Anonymous authors

Paper under double-blind review

Abstract

What shapes the content of a recurrent navigation agent’s learned internal state — the task, the architecture, or the structure of the visual input? Prior work has shown that blind PointNav agents spontaneously develop cognitive-map-like representations in their recurrent state under task pressure alone [Wijmans et al., 2023]. We show that the blind result is a special case of a broader principle: *the visual encoder’s information capacity determines whether the LSTM has to carry world-frame spatial structure*. Holding task, optimiser, and LSTM backbone identical and varying only the visual pathway (no vision; a resolution-collapsed encoder; uniform vision; foveated vision at a fixed gaze; foveated vision with a learned gaze), we probe the recurrent state for world-frame position and heading. All five conditions encode top-layer GPS in their LSTM in the typical-episode window. What separates conditions is what happens *later*: when the visual encoder cannot resolve world-frame structure (blind, a resolution-collapsed encoder), the LSTM *maintains* a persistent, integration-capable code across episode duration; when it can (full-resolution encoders), the LSTM *overwrites* the GPS code as visual features take over. We call this the *encoder-memory race*; it recasts the “cognitive maps emerge without vision” story as “cognitive maps emerge — and persist — when the visual pathway cannot resolve the map on its own.” A second finding cuts across conditions: five agents that solve the task equally well develop recurrent representations in essentially orthogonal linear subspaces, and cross-condition memory transplants incur task cost well above transplant mechanics — sensor structure shapes representational *format*, not only strength. A third question — whether static gaze location within a rich-encoder regime acts as a second content axis — is tested by a foveated-shifted causal control whose training is in progress. The H1 ordering reproduces on a held-out MP3D evaluation, consistent with an agent-internal rather than Gibson-specific mechanism. Three implications, each scoped to the recurrent PointNav setting we study. (i) Encoder capacity correlates inversely with linear, top-layer LSTM spatial encoding; *if* the correlation reflects causation (which the encoder-resolution scaling sweep, in progress, tests directly), it would also be a practical training-time lever for inducing cognitive-map-like internal representations. (ii) “More pixels, richer memory” is the wrong intuition within this setup; the relationship across our five conditions is inverse. (iii) The pattern parallels independent animal-navigation findings — enhanced hippocampal coupling in congenitally blind humans, the disruption of rodent grid-cell firing in the dark (with navigation nonetheless preserved), cross-modality hippocampal remapping in bats — without depending on their biophysics, suggestive of a sensor-memory coupling principle whose generality across architectures (transformers, non-RL learners, non-navigation tasks) we do not test here.




Figure 1: The five visual conditions differ only in what reaches the agent’s encoder. (a) Blind: no visual input; only the non-visual sensor stack (goal-in-start-frame, GPS, compass, close-to-goal scalar) — the Wijmans [Wijmans et al., 2023] configuration. (b) Uniform: 256×256 full-resolution RGB. (c) Foveated (fixed gaze): eccentricity-dependent Gaussian blur centred on the image middle ($\sigma_{\max} = 8$, quadratic falloff). The two additional conditions not shown are *matched-compute* (a down-sampled uniform view with the same total pixel budget as foveated) and *foveated (learned gaze)*, which adds a learned (x, y) gaze offset to (c). (d) All conditions solve PointGoal navigation in the same Habitat Matterport3D / Gibson scenes; the top-down view is shown for reference only and is not available to the agent.

1 Introduction

Recurrent navigation agents trained with deep RL build internal states that support memory, planning, and generalisation, and the content of those internal states is at least as interesting as the task metrics they drive. A natural question is what *determines* that content: architecture, objective, training budget, or visual input? Wijmans et al. [Wijmans et al., 2023] showed that map-like spatial representations emerge from task pressure alone, even in blind agents that receive only goal displacement and proprioception; Dwivedi et al. [Dwivedi et al., 2022] probed sighted agents for related concepts; Banino et al. [Banino et al., 2018] and Cueva & Wei [Cueva and Wei, 2018] demonstrated grid-cell-like structure in path-integration networks. Together these results establish that *some* spatial structure is learned without being designed in. They leave open the complementary question that we take up here: within agents that all achieve high task success, how much of *what* recurrent memory encodes is determined by the structure of the visual input rather than by the task? In biology the answer is clear — primate hippocampal cells code joint place–view–action state because primates sample the scene with a foveated sensor [Wirth et al., 2017, Rolls, 2024], and the same bat produces non-overlapping hippocampal populations in vision versus echolocation [Geva-Sagiv et al., 2016]. We investigate the analogous question in a controlled agent setting where the only free variables are the visual input structure and the gaze that directs it.

Our five-condition design isolates three degrees of freedom: sensor type (blind / uniform / foveated), encoder capacity (uniform’s full ResNet-18 vs. matched-compute’s 1×1 -collapsed feature map), and gaze location (fixed vs. learned). Every condition uses the same DD-PPO learner, the same Gibson-0+MP3D training pool, the same LSTM backbone, and the same non-visual sensor stack (Figure 1). We adopt the LSTM architecture from Wijmans et al. [Wijmans et al., 2023] deliberately: the emergent-map literature we build on is recurrent, and matching that architecture lets our findings be interpretable as statements about *the structure of the task-shaped memory* rather than about the recurrent-versus-transformer design choice. We discuss in §5 whether the logic plausibly extends to transformer-backed or non-navigation systems — an empirical question we do not settle here.

We test three hypotheses. **H1 (encoder–memory race)**. Agents whose visual encoder cannot resolve world-frame position or heading should develop these quantities in their recurrent state; agents with a rich visual encoder should not. Signature: current-state GPS / compass probe R^2 tracks visual-encoder information content monotonically across the five conditions, with persistent decoding at lags of ≥ 10 time-steps. **H2 (sensor structure $\rightarrow$ format divergence)**. Different sensor conditions should produce internal representations in distinct linear subspaces, not scaled versions of a common code. Signature: pairwise CKA near zero, probe weights fail to transfer across conditions while within-condition probes succeed, and cross-condition memory transplants degrade task performance beyond same-condition self-transplants. **H3 (gaze location $\rightarrow$ orthogonal content lever)**. Where a foveated sensor is centred should act as a representational-content axis independent of the sensor transform. We test this with a foveated-shifted control in which gaze is hardcoded to a different image location from the default central gaze; same foveation transform, same training budget, different gaze target. Our minimal learned-gaze module collapsed to a static output (Appendix B), so we use that collapsed location for the shifted-gaze control rather than as evidence about active gaze; the foveated-shifted training is in progress (§4.5). Each hypothesis has a direct analog in the animal-navigation literature [Kupers and Ptito, 2014, Chen et al., 2016, Geva-Sagiv et al., 2016, Jutras et al., 2013, Najemnik and Geisler, 2005], which we return to in Discussion as evidence that the effects reflect general sensor–memory coupling principles rather than architectural artefacts.

Findings H1 is supported by a clean monotonic effect at the level of *temporal stability* of the top-layer GPS code. All five conditions encode GPS in the LSTM in the typical-episode window (steps 50–200, $R^2 \in [0.6, 0.95]$); the bottleneck conditions then *maintain* that code across episode duration (blind $R^2 = 0.92$ even at step 800+; matched $R^2 = 0.68$), while rich-encoder conditions *overwrite* it as visual-encoder features take over (uniform / foveated / foveated-learned all collapse to chance or below by step 400–800). Aggregated across full-length episodes (Table 1, no step cap), this presents as $R^2 = 0.95/0.78$ for blind / matched-compute and $R^2 \in [-2.4, +0.06]$ for the rich-encoder conditions. Bottleneck conditions’ code is persistent at lags ≥ 20 time-steps. The same ordering reproduces on a held-out MP3D evaluation, confirming an agent-internal mechanism. H2 is supported by four independent measures of cross-condition divergence: pairwise CKA at the noise floor ($< 10^{-4}$), catastrophic probe-transfer failure ($R^2 \ll -800$ in every off-diagonal cell), perfect 1-NN cluster purity (1.000 across 7500 pooled hidden states, vs. chance 0.20), and behaviourally costly cross-condition memory transplants (SPL drops 0.19–0.21 between rich-encoder pairs, well above the self-transplant noise floor). H3 asks whether gaze location is a second content axis *within* the rich-encoder regime; the clean causal test (foveated-shifted control) is in training at submission and will populate §4.5. A non-linear (MLP) probe recovers position from the rich-encoder *top-layer* hidden states ($R^2 \in [0.51, 0.73]$), sharpening rather than weakening H1. We interpret this together with the per-layer probes (Figure 2c): every condition has linearly-decodable GPS at LSTM Layer 0 (where the GPS sensor embedding enters the input vector — so this is at least partly a sensor-pass-through effect, not a network “computation”); the divergence is what happens through Layers 1–2. Bottleneck conditions retain linearly-accessible GPS to the top layer used by DD-PPO’s linear action head; rich-encoder conditions do not. The MLP recovery at the top layer is consistent with position information persisting in a non-linear subspace, though we cannot from this alone rule out MLP exploitation of position-correlated features (wall distance, scene identifiers) rather than position itself — a question the encoder-resolution scaling sweep (Appendix E, in progress) will help resolve. H1 is most precisely stated as a claim about *linear, top-layer, policy-accessible* GPS encoding.

Contributions (i) A controlled five-condition ablation that decouples “no vision” from “visual-encoder bottleneck” via the matched-compute condition, making the encoder–memory race claim falsifiable. (ii) Convergent evidence (linear probes, MLP probes, CKA, probe transfer, 1-NN purity, memory transplant, shortcut discovery, MP3D generalisation, per-layer probes) for two of the three hypotheses, with a clean methodology disclosure including the deterministic-rollout protocol that fixes a probe-collection confound documented in §3.3. (iii) A mechanistic re-reading of Wijmans et al.’s blind-agent finding: cognitive maps emerge not because vision is absent but because the visual encoder cannot resolve the map, predicting an encoder-capacity scaling curve we begin to test in Appendix E.

2 Related Work

Cognitive maps in RL navigation agents Wijmans et al. [Wijmans et al., 2023] demonstrated that map-like representations emerge in the recurrent memories of DD-PPO agents even without visual input; Dwivedi et al. [Dwivedi et al., 2022] probed sighted agents for related spatial concepts; Banino et al. [Banino et al., 2018] and Cueva & Wei [Cueva and Wei, 2018] found grid-cell-like units in recurrent networks trained on path integration; Gornet & Thomson [Gornet and Thomson, 2024] extended this via visual predictive coding. Hennig et al. [Hennig et al., 2023] showed that RL under partial observability can induce belief-state-like representations without explicit probabilistic objectives. Ramakrishnan et al. [Ramakrishnan et al., 2025] examined spatial cognition in frontier models. All these studies use either absent or spatially uniform perception; none vary input quality spatially, and none test whether different sensors produce different representational formats.

Foveated vision and gaze policies in deep learning Foveated rendering [Strasburger et al., 2011, Rosenholtz, 2016] exploits eccentricity-dependent acuity falloff for efficiency. Learned glimpse policies appear in recurrent attention models [Mnih et al., 2014]; Atanov et al. [Atanov et al., 2024] showed that sensor layout changes downstream task performance. Pourrahimi & Bashivan [Pourrahimi and Bashivan, 2025] probed a foveated visual-search model for neuroscience-predicted features; Shang & Ryoo [Shang and Ryoo, 2023] jointly learned gaze and motor policies following the Sprague & Ballard [Sprague and Ballard, 2003] framework. These address visual search or task performance; none asks how sensor structure reshapes the learned spatial memory of a downstream navigation policy.

Probing neural representations Belinkov [Belinkov, 2022] surveys probing classifiers; Hewitt & Liang [Hewitt and Liang, 2019] introduced control tasks to distinguish genuine structure from probe memorisation; Kornblith et al. [Kornblith et al., 2019] introduced linear CKA for cross-model similarity. We adopt linear probing, the Hewitt–Liang selectivity metric, and unaligned linear CKA, and complement them with two behavioural interventions directly on the recurrent state (memory transplant and shortcut discovery) to validate the probe-based findings outside the linear-probe framework.

Biological precedent Our hypotheses and framing draw on three threads of animal-navigation research. (i) Primate hippocampal cells code joint place–view–action state [Wirth et al., 2017, Rolls, 2023]; Rolls [Rolls, 2024] argues that primate *spatial-view cells* and rodent *place cells* reflect the different sampling statistics of foveated versus near-panoramic vision. (ii) The same bat in the same environment produces non-overlapping hippocampal populations when it switches between vision and echolocation [Geva-Sagiv et al., 2016]—format-level divergence across sensor modalities within a single animal. (iii) In primates, each saccade resets hippocampal 3–8 Hz rhythms and the reliability of the reset predicts subsequent memory [Jutras et al., 2013]; only overt (not covert) attention automatically writes to visual working memory [Tas et al., 2016]. Gaze targets themselves are selected in an information-gain-maximising fashion consistent with ideal-observer predictions [Najemnik and Geisler, 2005, Gottlieb et al., 2013, Hayhoe and Matthis, 2018]. We use these results only to frame our ML predictions and interpret our findings; our contribution is the controlled ablation in a learned agent that cross-species and cross-modality biological comparisons cannot cleanly supply.

What is new here Three things distinguish this work from the studies above. (i) A *controlled* five-condition ablation with everything but the visual pathway held identical, including the matched-compute condition that decouples “no vision” from “visual encoder bottleneck” — the lever that makes the encoder–memory race claim falsifiable. (ii) A protocol-corrected probing methodology: deterministic action selection at probe time (the policy used by every other component of our eval pipeline) rather than stochastic sampling, which we found inflates probe R^2 for high-entropy policies via a 4-step quasi-static-trajectory artefact (§3.3, Sampling protocol). (iii) Convergent behavioural validation of the probe-based H1 and H2 findings via memory transplant and shortcut discovery, both of which sit *outside* the linear-probe framework.

3 Methods

3.1 Agent Architecture and Training

All agents are trained on PointGoal navigation in Habitat [Savva et al., 2019] using DD-PPO [Wijmans et al., 2020] on a merged pool of Gibson-0+ [Xia et al., 2018] (411 scenes) and Matterport3D train (61 scenes), totalling 472 training scenes. Evaluation uses 18 held-out MP3D test scenes (1008 episodes).

All agents share an identical 3-layer LSTM (512-d) backbone and the Wijmans non-visual sensor stack: goal-in-start-frame, GPS, compass, and a close-to-goal scalar, each projected to 32-d and concatenated with a 32-d previous-action embedding. All sighted conditions are trained for a target of 250 M environment frames; the blind baseline is trained to 342 M frames to allow slower convergence from proprioception alone.

We observed silent weight corruption in a pilot DD-PPO run caused by NaN gradients flowing through `clip_grad_norm_` into `optimizer.step()`; we mitigate this with a gradient-sanitisation patch to `PP0.before_step`, described in Appendix C.

3.2 Visual Conditions

Five conditions vary only the visual input (Figure 1):

1. *Blind*: no visual input; non-visual sensor stack only, replicating Wijmans et al. [Wijmans et al., 2023].
2. *Uniform*: full-resolution 256×256 RGB encoded by ResNet-18.
3. *Foveated (fixed)*: eccentricity-dependent Gaussian blur ($\sigma_{\max} = 8$, quadratic falloff) with gaze locked at image centre; same ResNet-18 encoder.
4. *Matched-compute*: 48×48 uniform RGB. At this resolution, ResNet-18’s 32-fold spatial downsampling collapses the feature map to 1×1 , removing the spatial dimensions of the encoder’s output (channel information is preserved as a flat 512-d vector). **[TODO: Verify whether channel info alone in matched-compute encodes any usable position signal — a small probe on the encoder feature map (not LSTM state) would settle this.]**
5. *Foveated (learned)*: same blur model, but gaze centre predicted by a lightweight MLP trained end-to-end with the navigation policy.

The matched-compute condition is central to H1: its 1×1 encoder feature map provides a visual observation while leaving the encoder unable to resolve world-frame position from the current frame, separating the “no vision” account from the “encoder bottleneck” account. We attribute its strong LSTM GPS encoding to the feature-map collapse specifically, but the 48×48 input also has lower spatial resolution than uniform’s 256×256 , so we cannot from this single condition fully separate “ 1×1 feature-map collapse” from “low input resolution.” The encoder-resolution scaling sweep (Appendix E, in progress) varies input resolution at fixed encoder stack to test this attribution directly.

Implementation note The foveated conditions (fixed, shifted, learned) disable the running-mean-and-variance input normaliser used by uniform and matched-compute, to avoid an in-place autograd conflict along the gaze-decoder gradient path. All reported results were obtained with this setting; we have not directly verified that re-enabling the normaliser would leave the H1 ordering unchanged. **[TODO: Verify normaliser invariance: rerun foveated-fixed with the normaliser re-enabled (the autograd conflict is gaze-decoder-specific, so feasible for fov-fix), and confirm the H1 GPS R^2 does not move.]**

3.3 Probing Methodology

After training, we freeze all weights and collect per-step LSTM hidden states (top-layer h , 512-d) paired with ground-truth spatial labels from $n=500$ Gibson held-out validation episodes (the split not used in training). Rollouts use *deterministic* (argmax) action selection, matching the eval protocol of our shortcut-discovery, memory-transplant, and standard habitat-baselines evaluators. Under this protocol the rollout lengths span roughly ~ 100 – 400 mean steps depending on condition (longer for blind, shorter for high-SPL conditions), so we use 5-fold episode-level CV to keep the train/test splits balanced per scene rather than length-matching across conditions. **[TODO: A controlled**

length-matching ablation (truncate every condition to a common length, re-run probes) is left for follow-up; we expect it to leave the H1 ordering unchanged because the bottleneck conditions decode at high R^2 even at the smallest sub-sample, but have not verified it.] We train Ridge regression probes ($\alpha = 10$) with episode-level 5-fold cross-validation, reporting mean $\pm \sigma$ across folds, and verify probe specificity with Hewitt–Liang label-permutation selectivity [Hewitt and Liang, 2019].

Probed variables: GPS position (episodic), compass heading (episodic), distance-to-goal, lag- k path history (target value at $t - k$ decoded from $\mathbf{h}_t$, $k \in \{0, \dots, 20\}$), visited-region occupancy, full occupancy grid (20×20), per-unit spatial information (place-cell analysis), and the ego-frame goal vector $\mathbf{R}(-\theta)(\mathbf{g} - \mathbf{p})$ decomposed into scalar goal-distance and 2-D direction.

Sampling protocol An earlier draft of this work used stochastic (policy-distribution) action selection for probe rollouts. Under stochastic sampling, conditions with higher action-entropy (specifically the rich-encoder conditions foveated, uniform, matched, blind under their trained stochastic policies) sampled the STOP action early with probability ~ 0.25 per step, producing mean episode lengths of ~ 4 steps and target variance two orders of magnitude below the episodic range. The resulting probe R^2 values were trivially high across all conditions (any predictor fits a near-constant target). This confounded cross-condition comparisons, and we switched to the deterministic protocol used here and elsewhere in the probe/eval pipeline. All results in this paper use the deterministic protocol.

For foveated-fixed, all reported analyses (probes, behavioural, eval) use checkpoint ckpt.36 at $\sim 174\text{M}$ frames, the last clean intermediate before the silent-corruption window (Appendix C).

3.4 Cross-Condition Analysis

To test H2, we employ two complementary methods. *Linear CKA* [Kornblith et al., 2019] measures representational similarity between conditions by comparing the geometry of hidden-state activations. Values near 1 indicate identical geometry; near 0 indicates unrelated representations. *Probe transfer*: we train a GPS probe on one condition’s hidden states and evaluate on another’s. If spatial information is encoded in a shared linear subspace, cross-condition transfer should succeed. Catastrophic failure indicates fundamentally different encoding formats. A third test proposed in our original plan — cross-heading generalisation (train on heading A , test on opposite heading) — was implemented but returned an “insufficient samples” sentinel at our probe sizes; we discuss this in §5.5.

3.5 Behavioural Probes

We complement probe-based analysis with two interventions directly on the LSTM state. *Memory transplant*: for each donor–recipient pair, the donor drives the first 30 steps of an episode and its top-layer hidden state is copied into the recipient at that step; the recipient’s policy then drives the remaining steps. We compare *baseline* (recipient alone, full episode), *self-transplant* (recipient as its own donor, controls for transplant mechanics), and *cross-transplant* (different donor). Midpoint 30 was chosen because Gibson PointNav episodes average ~ 100 – 120 steps; a later midpoint left too few steps after transplant. We use 150 episodes per pair per condition, sampled uniformly from the held-out Gibson validation split. *Shortcut discovery*: for each condition, we run episodes in pairs of two using the same scene but different goals, sampling 20 scenes with 10 paired episodes per scene (200 paired runs per condition per setting). Per-condition SPL is reported as the mean of the 20 scene-level mean SPLs (prevents high-episode-count scenes from dominating). We compare *reset* (LSTM zeroed between the two episodes) vs. *persistent* (LSTM state carried over) on the second-episode SPL. A reduction under the persistent setting measures how tightly the stored representation is locked to the previous goal. For the foveated-fixed condition, all probes and behavioural runs use the ckpt.36 checkpoint ($\sim 174\text{M}$ frames, the last clean intermediate before the silent-corruption window, §4).

4 Results

4.1 Five conditions, same task

Table 1 collects the headline per-condition numbers this section returns to. All sighted conditions solve PointGoal at 96–99% success; blind reaches 93% at a longer training budget, reproducing Wijmans et al. [Wijmans et al., 2023]. Success is tightly bunched, so “the agents differ in what they

Table 1: Per-condition headline. Probe R^2 columns are the 5-fold episode-level CV mean $\pm$ std on full-length deterministic-rollout trajectories (no per-episode step cap). *GPS* and *Compass* are world-frame. The bottleneck-vs-rich-encoder gap is in temporal *stability* of the code, not its mid-episode existence (Figure 2b): all five conditions decode top-layer GPS at $R^2 \in [0.6, 0.95]$ in the typical-episode window (steps 50–200); the rich-encoder long-tail collapse is what produces the sub-chance numbers below. Bold marks the H1-relevant ordering (blind $>$ matched $\gg$ rich-encoder).

Condition	Frames	SPL	Success	GPS R^2 (no-cap)	Compass R^2 (no-cap)
Blind	342M	0.59	0.93	$+0.95 \pm 0.02$	$+0.81 \pm 0.08$
Matched-compute	250M	0.67	0.98	$+0.78 \pm 0.10$	$+0.64 \pm 0.10$
Uniform	250M	0.85	0.99	-0.31 ± 0.86	$+0.36 \pm 0.23$
Foveated (fix)	174M	0.78	0.96	$+0.06 \pm 0.88$	$+0.07 \pm 0.69$
Fov-learned	250M	0.82	0.98	-2.43 ± 3.98	-1.34 ± 3.14

do” is an unlikely explanation for the encoded-content gap we report (it does not strictly rule it out). The probe machinery itself is calibrated: distance-to-goal decodes robustly across all five conditions ($R^2 \in [0.81, 0.90]$, Figure 2a, right), confirming that Ridge has the capacity to fit high- R^2 scalars when the hidden state contains them; and Hewitt–Liang label-permutation selectivity is high in the conditions where the GPS code itself is high (blind GPS / compass selectivity $+0.95 / +0.81$; matched $+0.72 / +0.56$), consistent with a genuine signal rather than label memorisation. Two training-stability caveats are disclosed in full in Appendix C: foveated-fixed uses ckpt.36 at ~ 174 M frames (last clean checkpoint before a silent NaN-gradient corruption at ≈ 182 M); and matched-compute’s 1×1 feature-map collapse, originally noticed as a training quirk, is what makes matched-compute a useful test of the encoder-bottleneck framing (we cannot from one condition alone fully separate the 1×1 feature-map effect from the lower input resolution; see §3.2 and §4.4).

4.2 Encoder–memory race: a temporal-stability claim (H1)

Finding Visual-encoder capacity inversely tracks LSTM spatial encoding across our five conditions (Figure 2a). Blind agents encode current GPS in the recurrent state at $R^2 = 0.95 \pm 0.02$; matched-compute agents (48×48 RGB input, 1×1 encoder feature map) at $R^2 = 0.78 \pm 0.10$; all three full-resolution-encoder conditions (uniform, foveated, foveated-learned) decode GPS at chance ($R^2 \in [-2.4, +0.06]$ with CV $\sigma \in [0.7, 4.0]$). Compass shows the same ordering at smaller effect size ($0.81 \rightarrow 0.64 \rightarrow 0.36 \rightarrow 0.07 \rightarrow -1.3$); DtG decodes robustly everywhere ($R^2 \in [0.81, 0.90]$). **[TODO: Whether this correlation reflects causation (encoder capacity drives LSTM encoding) is what the foveation strength sweep, log-polar control, and encoder feature-map probe in §4.4, and the encoder-resolution scaling sweep in Appendix E, test directly. Multi-seed runs of the existing five conditions are in progress to bound seed-to-seed variability.]**

Temporal stability across episode duration (H1’s precise claim) The no-cap probe in Figure 2a does not separate “does the LSTM ever encode GPS” from “does the LSTM maintain that encoding as the episode lengthens.” We test the temporal version directly with a step-binned out-of-fold probe (Figure 2b; OOF prediction so each bin’s R^2 is from a probe never trained on that bin’s episode). *All five conditions* encode top-layer GPS in the typical-episode window (steps 50–200, $R^2 = 0.6$ – 0.95); rich-encoder conditions are not at chance at this scale. They diverge from the bottleneck conditions in what happens *later*. Blind and matched-compute hold $R^2 > 0.7$ across all step bins out to 800+. Uniform / foveated / foveated-learned peak at steps 100–200, then decay; by step 400–800 they are at or near chance, and in the long-tail 800+ bin they are catastrophically negative (-0.37 to -5.67). The “rich-encoder $\approx$ chance” headline in Table 1 is therefore predominantly long-tail behaviour. *H1 is most precisely a claim about the temporal stability of the linearly-decodable top-layer GPS code, not its existence at any moment.* A separate lag- k probe corroborates: decoding GPS_{t-k} from $\mathbf{h}_t$ holds $R^2 = 0.92$ at $k = 20$ for blind and $R^2 > 0.5$ at $k = 20$ for matched, while rich-encoder conditions decode at chance for every $k \leq 10$.

Disambiguating Layer-0 sensor pass-through (H1’s precise level of the network) A natural objection is that all five conditions have linearly-decodable GPS at LSTM Layer 0 ($R^2 > 0.91$, Figure 2c) — so how can H1 distinguish them? Layer 0 is where the GPS sensor embedding enters the LSTM input vector, so this lower bound is at least partly a sensor-pass-through effect rather than the network “computing” position. The interesting signal is the per-layer trajectory: bottleneck




Figure 2: H1 evidence in one figure. (a) Current-state probe R^2 (5-fold CV, Ridge $\alpha = 10$, full-length deterministic rollouts) for GPS, compass, and DtG. Blind and matched-compute (the bottleneck conditions) decode world-frame GPS / compass robustly; uniform / foveated / foveated-learned do not. DtG decodes everywhere because every agent must track it. (b) Top-layer GPS R^2 binned by step-in-episode (out-of-fold prediction): all five conditions encode top-layer GPS in the typical-episode window (steps ~ 25 –200), and the rich-encoder conditions then *overwrite* the code (negative R^2 in the long-tail 800+ bin), while bottleneck conditions hold $R^2 > 0.7$ to step 800+. The headline “rich-encoder $\approx$ chance” in panel (a) is therefore primarily long-tail behaviour. (c) Per-LSTM-layer GPS R^2 (no-cap, single-fit on top-layer hidden states): all five conditions have GPS at Layer 0 (where the GPS sensor embedding enters the LSTM input vector — this lower-bound is at least partly a sensor-pass-through effect, not the network “computing” position); the divergence is at Layer 2, the policy readout. (d) Same checkpoints, evaluated on held-out MP3D (300 episodes/condition, deterministic protocol). The bottleneck-vs-pass-through ordering is preserved; foveated does recover to a mild positive ($R^2 = +0.35$) on MP3D, consistent with Gibson-specific feature use by the foveated encoder.

conditions *retain* the GPS code through Layers 1–2 to the top-layer policy readout, while rich-encoder conditions discard it (uniform Layer 2 GPS $R^2 = -1.56$, foveated -1.64 , foveated-learned -10). Since the standard DD-PPO action head is a linear projection from the top-layer state $\mathbf{h}_2$, “linearly decodable at Layer 2” is precisely what the policy itself can use. (*Side observation, not yet explained: blind’s Layer 1 GPS R^2 is lower (0.61) than its Layer 0 (0.99) or Layer 2 (0.95); the LSTM does not pass GPS through monotonically. Flagged as an artefact-or-mechanism question for follow-up.*)

Non-linear-probe disambiguation An MLP probe on the same top-layer state recovers position from rich-encoder networks (foveated $R^2 = 0.51$, foveated-learned $R^2 = 0.73$), *consistent with* position information persisting in a non-linear subspace — though we cannot from this alone exclude MLP exploitation of position-correlated features (wall distance, scene identifiers, episode-phase markers) rather than position itself. What this establishes is that the bottleneck-vs-pass-through divergence is not just about whether *any* probe can recover position; it is specifically about whether a *linear, top-layer* probe can. The encoder-resolution scaling sweep (Appendix E, in progress) and the foveation-strength sweep (§4.4) test the bottleneck framing more directly.

Robustness to dataset shift (MP3D) The H1 ordering is preserved on a held-out MP3D evaluation (300 episodes per condition, deterministic protocol; same trained checkpoints, no re-training; Figure 2d). Blind GPS R^2 moves from 0.95 (Gibson) to 0.93 (MP3D); matched moves from 0.78 to 0.71. Foveated recovers from 0.06 to a mild positive $+0.35$ on MP3D, and foveated-learned compass

Table 2: Probe-transfer GPS R^2 on deterministic-rollout hidden states: train on row condition, test on column. Self-accuracy diagonal is 0.89–0.99; every cross-condition transfer is catastrophically negative ($R^2 \ll -800$, often $\ll -10^4$), reflecting that a probe trained on one condition’s hidden-state geometry does not just degrade on another — it predicts off-manifold values.

Train \ Test	Blind	Uniform	Fov-fix	Fov-lrn	Matched-compute
Blind	+0.99	−844	−11,301	−10,731	−55,706
Uniform	−1,383	+0.90	−7,046	−3,305	−110,317
Fov-fix	−2,091	−2,792	+0.92	−1,263	−39,100
Fov-lrn	−6,722	−2,423	−6,205	+0.89	−38,842
Matched-compute	−10,621	−4,832	−4,984	−58,169	+0.96

shows a large cross-dataset swing ($-1.34 \rightarrow +0.41$). We mark these single-seed cross-dataset shifts as observations requiring multi-seed replication before attributing them to a specific mechanism (Gibson-specific feature use by the foveated encoder is one candidate); the qualitative bottleneck-vs-pass-through partition is unaffected. Compass tells the same qualitative story across datasets (numbers in Appendix A).

Proposed mechanism We *propose* the following account, consistent with the data but not directly tested. In PointNav, GPS and compass are observations the policy reads each step; the recurrent state need only re-encode them if it cannot rely on the visual stream to do so. A full-resolution RGB encoder can re-derive position and heading from the current frame — so as the episode progresses and the rich-encoder LSTM accumulates other task-relevant features (visual landmarks, scene history, goal-direction integration), it has no incentive to keep refreshing the GPS code, and the top-layer state drifts off it. An impoverished encoder cannot re-derive position from vision alone; the LSTM has to keep the GPS code linearly readable for the policy throughout the episode, and the top-layer code stays stable. Blind and matched-compute bracket this picture: blind has *no* visual encoder (pure compensation); matched has a *collapsed* 1×1 encoder (partial compensation); foveated, uniform, and foveated-learned all have full encoders (pass-through). Task success is tightly bunched, so a skill-gap artefact is unlikely to explain the encoding gap. The temporal dynamic in Figure 2b (rich-encoder GPS code emerging early then decaying) is the direct fingerprint of this account. **[TODO: A direct causal test of the mechanism — e.g., training a rich-encoder agent with GPS perturbed mid-rollout, or training a bottleneck agent with the GPS sensor removed — is left for future work.]**

Relation to Wijmans et al. Wijmans et al. [Wijmans et al., 2023] showed that blind PointNav agents develop spatial representations in LSTM memory. Our blind result ($R^2 = 0.95$) replicates this; matched-compute extends it beyond the blind special case to a condition that has visual input but cannot resolve world-frame structure from it. We interpret this as evidence that the relevant trigger is the visual encoder’s information capacity, not the binary presence/absence of vision. §4.4 and Appendix E test the trigger directly.

Bio parallel The pattern parallels the animal-navigation literature without depending on its mechanisms: congenitally blind humans develop enhanced hippocampal coupling and outperform sighted controls on non-visual spatial tasks [Kupers and Ptito, 2014]; rodent grid cells lose their hexagonal periodicity in darkness even with head-direction input intact [Chen et al., 2016], yet the animal still navigates — so spatial bookkeeping that the (vision-dependent) grid code can no longer carry has to come from somewhere else. We use these only as analogy: in both cases visual-pathway absence shifts spatial bookkeeping into a non-visual substrate, the same qualitative direction we measure in LSTM hidden states. We do not claim the agent and biological mechanisms are the same.

4.3 Format-level divergence across conditions (H2)

Lead with the behavioural test The most informative H2 measurement is also the most direct test of “do these representations matter?”: *memory transplant* (Figure 3, right). For each donor–recipient pair, the donor drives the first t steps of an episode, its top-layer hidden state is copied into the recipient at step t , and the recipient’s policy drives the rest. The isolated condition-mismatch cost — cross-transplant SPL minus self-transplant SPL, which subtracts both a small $t=0$ protocol-noise floor and the rollout-divergence cost of mid-episode state replacement — is roughly stable across midpoints




Figure 3: H2 convergent evidence. *Left*: Unaligned linear CKA [Kornblith et al., 2019] between top-layer LSTM hidden states ($n=95,640$ steps per condition under deterministic-rollout probes); every off-diagonal is below 10^{-4} . *Right*: Memory transplant across midpoints $\{0, 15, 30, 60, 90\}$ ($n = 100\text{--}150$ episodes/cell). (a) Baseline / self-transplant / cross-transplant SPL; at $t=0$ self- and cross- reduce to a protocol-noise floor. (b) Cross minus self (the isolated condition-mismatch component) is roughly stable for $t \geq 15$: foveated-learned $\rightarrow$ foveated incurs -0.17 SPL on average even though all pairwise CKA values sit at the noise floor. Probe-based similarity is uninformative at this scale; transplant reveals the functional gap.

$t \geq 15$ and carries a clear ranking: foveated-learned $\rightarrow$ foveated incurs -0.17 SPL, foveated $\rightarrow$ uniform -0.10 , and foveated $\rightarrow$ blind essentially neutral (-0.02). Pairs that look identically near-zero in linear-similarity measures differ markedly under transplant. Probe-based similarity is uninformative at this scale; the functional gap is what matters. **[TODO: The remaining cross-condition transplant pairs (e.g., blind $\rightarrow$ matched, matched $\rightarrow$ uniform) are not yet measured; running them would let us state the format-divergence finding as a complete 5×5 behavioural matrix rather than a representative subset.]**

Convergent linear-similarity evidence Three further linear-similarity measures converge on the same “different format” conclusion. *Pairwise CKA* [Kornblith et al., 2019] sits at the noise floor for every condition pair (Figure 3, left; off-diagonals all $< 10^{-4}$); within-condition split-half CKA is ≈ 1 , so the near-zero off-diagonals are not estimator noise. *Cross-condition probe transfer* is catastrophically negative in every off-diagonal cell (Table 2, $R^2 \ll -800$, often $\ll -10^4$): a probe trained on one condition’s hidden states does not just degrade on another, it predicts off-manifold values. *1-NN purity*: pooling 1500 random top-layer hidden states from each of the five conditions and asking whether each sample’s nearest neighbour (Euclidean, full 512-d state) comes from the same condition, the answer is yes for 7500/7500 in our sample (1.000 vs. chance 0.20). A t-SNE projection of the same pool (Appendix A) confirms visually that the five conditions occupy non-overlapping regions. **[TODO: Whether 1-NN purity holds at much larger sample size or under domain shift (MP3D-pooled, fov-shifted-pooled when available) is a further-investigate item.]**

Ruling out alternatives Self-probes succeed for every condition, so probe capacity is not the bottleneck. CKA is computed on a common-sample truncation, ruling out sample-size asymmetry. Task competence is tightly bunched (96–99% success across sighted conditions), so the divergence is not a skill difference. Transfer fails symmetrically (both directions), which rules out the possibility that one condition’s code simply embeds another’s. (CKA and probe transfer use common-sample truncation for fair cross-condition comparison; transplant uses full on-policy rollouts because the task is not truncatable. The two measurement scopes are deliberate, and their conclusions are consistent.)

Boundaries Our similarity tests are linear (CKA) or near-linear (t-SNE default); non-linear similarity measures could reveal higher-order shared structure. Our five sensors are diverse but not exhaustive; hybrid sensors (coarse uniform periphery + sharp fovea, depth-only inputs) might produce intermediate geometries.

4.4 Foveation: where it sits, and where the cleaner tests sit

The five-condition results above place our two foveated variants (foveated-fixed, foveated-learned) in the rich-encoder pass-through regime alongside uniform: under $\sigma_{\max} = 8$ Gaussian-blur foveation with the same ResNet-18 encoder stack, the LSTM does not develop a stable top-layer GPS code (Figure 2). This is itself a finding — it bounds where a soft Gaussian-blur foveation falls on the

encoder–memory race spectrum — but it leaves three sharper questions about foveation that the present five conditions cannot resolve. We submitted controlled experiments for each; the following four sub-paragraphs describe the experimental design and what the result will mean. Numbers will populate this subsection in revision; current status is tracked in Appendix D.

Foveation strength sweep (F1–F4, in progress) We train four additional foveated agents at $\sigma_{\max} \in \{2, 4, 12, 20\}$ (alongside the existing $\sigma_{\max} = 8$), holding everything else fixed — gaze location, encoder stack, training budget, training pool. The expected signature: at low $\sigma_{\max}$ the foveated agent should look like uniform (encoder still resolves world-frame); at high $\sigma_{\max}$ it should approach matched-compute (encoder cannot resolve world-frame, LSTM compensates). The R^2 vs. $\sigma_{\max}$ curve directly tests the encoder–memory race as a continuous lever, not a binary one.

Spatial sampling vs. blur (F3 log-polar control, in progress) Gaussian blur removes high-frequency content but preserves spatial layout: a foveated 256×256 image still has 256×256 pixels feeding the encoder, which can still resolve coarse position from peripheral structure. A log-polar foveation [Strasburger et al., 2011, Rosenholtz, 2016] resamples the image with a non-uniform spatial grid: high spatial sampling at the fovea, low sampling at the periphery. This is the architectural commitment that “foveated vision” makes in biology. We train an agent with the same ResNet-18 stack but a log-polar input, and ask whether *spatial-sampling foveation* (rather than blur-only foveation) triggers the encoder–memory race. If yes, it sharpens H1 from a “bottleneck of any kind” claim to a “spatial-resolution bottleneck specifically” claim.

Where is position information held in the foveated agents? (encoder feature-map probe, in progress) Our H1 evidence comes from probing the LSTM hidden state. The complementary question — can the encoder’s output already resolve position, or is the LSTM the only place world-frame structure lives? — is settled by probing the encoder feature map directly (post ResNet-18, pre-LSTM). We forward-hook the encoder’s compression layer and run the same Ridge probe on the flattened feature vector. Three outcomes are diagnostic: (i) encoder feature map decodes GPS for uniform / foveated but not matched — supports the claim that matched’s 1×1 collapse is the trigger; (ii) encoder feature map decodes GPS for all rich-encoder conditions including foveated — foveated agents have position in the encoder, the LSTM just does not promote it to the top layer (which would mean H1 is downstream of where the foveated bottleneck is); (iii) encoder feature map decodes GPS for matched too — the bottleneck framing of matched needs revision.

Foveated-shifted (gaze location, in progress) Gaze location is treated separately as H3 (§4.5), but the foveated-shifted control directly informs the foveation slot too: training a foveated agent with hardcoded static gaze at $(0.49, 0.62)$ matches foveated-learned’s collapsed gaze location while removing the learned-gaze module from the optimisation path. If foveated-shifted does *not* reproduce foveated-learned’s specific behaviours (Gibson–MP3D compass swing, transplant ranking), those behaviours can be attributed to the optimisation-path interaction and not to the static gaze location alone.

Why this slot exists Our title and framing focus on foveated vision; the present five conditions place foveated-fixed and foveated-learned within a framework but do not isolate foveation-specific effects (under simple Gaussian blur both behave like uniform). The four experiments above target specific foveation-specific predictions of the encoder–memory race framework. Each is a falsifiable test, designed and submitted at the time of writing. None depends on the others to be informative.

4.5 Gaze location as a second axis (H3)

H1 and H2 identify sensor structure as one axis on representational content. H3 asks whether gaze *location* — where the foveated window is centred — is a second, dissociable axis. We test the two forms described in §3.3 and previewed in the Introduction: an active-gaze form (the minimal learned-gaze module should move its gaze to task-useful locations), and a gaze-location form (two foveated agents whose gaze is centred at *different static* locations should produce different content). The active-gaze form is out of scope of our minimal gaze decoder and we do not claim anything about it; the gaze-location form is what our controlled experiments test.

The learned-gaze module collapses Our minimal learned-gaze implementation (deterministic sigmoid MLP, slow-gaze updates once per 128-step rollout segment; Appendix B) does not produce active gaze: after 250M frames the gaze output is frozen at $(0.49, 0.62)$ with per-dimension std $\leq 2.5 \times 10^{-4}$ (numerical noise). Two anti-collapse regularisers (Appendix B) restore gaze variance




Figure 4: Shortcut discovery: persistent vs. reset memory on the second of two paired same-scene/different-goal episodes. 20 Gibson scenes $\times$ 10 paired episodes each, bar height is the mean of the 20 scene-level mean SPLs. Conditions ordered by relative SPL drop (%): blind loses most (52%), foveated-learned loses least (15%). Larger drop = more brittle memory coupling.

only at the cost of task performance. We explicitly do not claim this shows active gaze cannot be learned — our gaze decoder is minimal by design (to be a controlled variable, not a model of active vision), and richer architectures (stochastic gaze, information-seeking intrinsic reward, attention-based token selection) would need separate investigation. What the collapse gives us for free is an opportunity to test the gaze-location form: after training, we have an agent that is *effectively foveated with a static gaze 30 pixels below image centre*.

Gaze-location test (foveated-shifted, in progress) Both rich-encoder foveated variants sit in the pass-through regime under current-state probing (Table 1), so the H3 test is whether *within* the pass-through regime, gaze location shifts content along other axes — non-spatial feature usage, behavioural memory-reliance (shortcut, Figure 4), or transplant compatibility. To isolate gaze location from training-dynamics confounds, we train a foveated-shifted control: an agent identical to foveated-fixed except gaze is hardcoded at (0.49, 0.62) instead of (0.5, 0.5) — matching foveated-learned’s collapsed gaze position but with no learned-gaze module in the optimisation path. **[TODO: Training in progress at submission; results will populate this paragraph.]**

Shortcut discovery (behavioural memory-reliance) We separately measure how tightly each policy couples action choice to integrated recurrent memory: run paired episodes in the same scene with fresh goals, compare second-episode SPL under *reset* (LSTM zeroed between episodes) vs. *persistent* (state carried over) memory. Larger relative SPL drop = more brittle memory coupling (Figure 4): blind loses 52%, uniform 41%, foveated 21%, matched 18%, foveated-learned 15%. The blind ranking is consistent with H1: blind has the strongest top-layer GPS code of the five conditions and is also the most brittle when that memory is stale; the three rich-encoder conditions, which decode GPS at chance from the top-layer state, are correspondingly less perturbed. *Note that “decodable from h” and “policy uses for decision-making” are not the same construct* — a probe shows what is *readable* from the state, not what the policy actually *relies on*; the shortcut experiment is the closer behavioural test of the latter. **[TODO: A direct measurement of “policy GPS reliance” (e.g., ablating GPS sensor input mid-rollout and measuring the SPL drop) would complement the shortcut signal; left for follow-up.]** *Anomaly: matched-compute drops only 18%, lower than its strong GPS encoding might predict.* A literal reading is that *using* a cognitive map for policy decisions is a separate property from *having* one — matched has the code but weights it less in policy output — but this rests on a single-seed shortcut run; we mark it as a candidate dissociation pending multi-seed replication of both the GPS probe and the shortcut measurement before drawing the “having vs. using” conclusion.

4.6 Boundaries and additional probes

Occupancy controls Two probes fail to discriminate between conditions and so define the boundary of our claims. Coarse occupancy (“have I visited this region?”) decodes above chance in every condition, so one-bit spatial working memory is ubiquitous and not part of the H1–H3 story. Metric

20×20 occupancy reconstruction fails everywhere, consistent with the SPACE benchmark’s finding that fine-grained layout requires non-linear probes or longer horizons [Ramakrishnan et al., 2025]: our claims are about representational *format* and *content*, not about high-resolution metric layout.

Goal vector No condition stores the ego-frame goal vector linearly: goal *distance* decodes strongly across conditions ($R^2 \in [0.81, 0.90]$, the DtG column), while goal *direction* and the full 2-D ego-to-goal vector are at chance or negative for every condition. This is unsurprising — PointGoal provides the goal vector as a per-step input, so the policy has no incentive to re-encode it — and is not part of the H1/H2/H3 story. The supplementary plot is in Appendix A.

Per-unit spatial information The number of LSTM units carrying > 1 bit of spatial information [Banino et al., 2018] is higher in bottleneck conditions than in rich-encoder conditions (Appendix A for the bar plot), consistent with a population code rather than a few dedicated “place neurons.” *Caveat:* this depends on the chosen 1-bit threshold and the per-scene grid resolution; we do not claim a sharp dissociation between “population code” and “few dedicated cells.”

5 Discussion

5.1 Two orthogonal axes: sensor structure and gaze location

Our findings unify under a single mechanism: the *encoder–memory race*. Whether the LSTM *maintains* a linearly-decodable world-frame spatial code at the top layer depends on whether the visual encoder can resolve that information from the current frame. When the encoder cannot — blind (no encoder) or matched-compute (1×1 feature map) — the LSTM keeps the GPS code stable across episode duration (H1, §4.2; Figure 2b). When the encoder can — uniform, foveated, foveated-learned all use full-resolution ResNet-18 encoders — the LSTM still encodes top-layer GPS in mid-episode (steps 50–200), but *overwrites* the code as visual-encoder features accumulate, leaving long-tail steps at chance or below. The encoder–memory race is therefore best read as a temporal-stability claim, not an existence claim. The Layer-0 finding (Figure 2c) is consistent: all conditions have GPS at Layer 0 (where the sensor embedding enters); the divergence is at Layers 1–2 *and* as the episode progresses. H2 (§4.3) cuts orthogonally: *what* the LSTM encodes is in different linear subspaces across conditions even where current-state position decodes at chance, and that linear difference matters for downstream behaviour (cross-condition memory transplants drop SPL more than transplant mechanics alone). H3 (§4.5) asks whether, *within* the rich-encoder pass-through regime, gaze location is a second representational-content axis; the clean foveated-shifted causal control that answers this is in training at submission time. The same encoder-capacity axis that collapses the original “foveated vs. uniform” question (both sit in the pass-through regime) surfaces a clean scaling prediction: GPS decodability should fall monotonically with encoder resolution, which we begin to test in Appendix E.

5.2 Why this, not alternative accounts

We rule out four alternative explanations.

Task-competence differences. Every condition solves PointGoal at 93–99% success, so the encoded-content gap cannot be attributed to differences in navigation skill (Table 1).

Probe-capacity / probe-fitting artefacts. DtG decodes robustly across all five conditions ($R^2 \geq 0.81$), demonstrating that Ridge has the capacity to fit high- R^2 scalars when the hidden state contains them. Where GPS does decode (blind, matched), Hewitt–Liang label-permutation selectivity is correspondingly high; where GPS is at chance, selectivity is also at chance because real and control probes face the same near-zero target-information ceiling. The chance-level GPS/compass R^2 for rich-encoder conditions is therefore a property of the hidden state, not of the probe.

Sample-size / trajectory-distribution artefacts. CV σ spans come from the same 500-episode pool with ~ 100 –400-step rollouts under identical deterministic evaluation; the protocol confound that haunted an earlier draft (4-step quasi-static trajectories under stochastic sampling) is documented and retired in §3.3.

Linear-probe artefacts. Two complementary checks address this. (a) *Behavioural:* the memory-transplant intervention (§4.3) goes outside the probe framework and reproduces the bottleneck-vs-pass-through ordering, with cross-condition transplants dropping SPL by 0.19–0.21 between

rich-encoder pairs. (b) *Non-linear probe*: a 2-layer MLP on the same top-layer hidden states recovers $R^2 \in [0.51, 0.73]$ from the rich-encoder conditions where Ridge sees chance. We interpret this as consistent with position information persisting in a non-linear subspace at the top layer, but emphasise the limits: an MLP probe can also exploit features merely *correlated* with position (wall distance, scene identifiers, episode-phase markers), so this number alone cannot disentangle “non-linear position code” from “MLP regressing through position-correlated features.” What it does establish is that the bottleneck-vs-pass-through divergence is not just about whether *any* probe can recover position — it is specifically about whether a *linear* top-layer probe can. Since the standard DD-PPO action head is a linear projection from h_2 , the linear-probe metric is precisely what the policy itself can use — so the H1 distinction maps directly onto policy-accessible encoding regardless of how the MLP-probe interpretation resolves. The encoder-resolution scaling sweep (Appendix E, in progress) tests the bottleneck framing more directly by varying input resolution while holding everything else fixed.

5.3 Biological precedent as convergent evidence

Each of our three findings has a parallel in the animal-navigation literature, which we take as one piece of convergent (not conclusive) evidence that the effects are not specific to a single LSTM architecture. Compensatory memory under reduced sensory access is documented in congenitally blind humans, who develop enhanced hippocampal coupling and outperform sighted controls on non-visual spatial-working-memory tasks [Kupers and Ptito, 2014], and in rodents whose grid-cell periodicity collapses in darkness despite intact head-direction inputs [Chen et al., 2016]—both format-level, not just strength-level, effects. Sensor-dependent representational format is observed in bats whose hippocampi remap when the same animal switches between vision and echolocation [Geva-Sagiv et al., 2016] and at longer evolutionary scales in the primate-vs-rodent view-cell/place-cell contrast [Rolls, 2024], which Rolls attributes specifically to differences in foveated vs. near-panoramic sampling statistics. Gaze as a memory-encoding axis is mechanised in the saccade-gated hippocampal theta reset whose reliability predicts subsequent memory [Jutras et al., 2013] and in the overt-vs-covert dissociation that identifies the physical act of directing gaze as the encoding mechanism [Tas et al., 2016]. We do not claim our agents implement these biological mechanisms, nor that the parallel structure settles the generality question; we do claim that the parallels make a pure single-architecture artefact less likely, and that the controlled agent setting complements biological comparisons where sensor type, architecture, and training objective cannot be independently manipulated.

5.4 Implications for ML practice and architectural scope

Within the recurrent PointNav setting we study, three implications follow. (i) *Encoder capacity correlates inversely with linear, top-layer, policy-accessible LSTM spatial encoding* across our five conditions. If the correlation reflects causation in this setting — which the encoder-resolution scaling sweep (Appendix E, in progress) tests directly — it would also be a practical training-time lever for inducing cognitive-map-like internal representations: impose an encoder bottleneck. We do not yet claim this generalises beyond PointNav with our specific architecture and training regime. (ii) *“More pixels” is not the right lens within this setting*: the linear top-layer GPS code is non-monotonic in raw pixel count — uniform 256×256 (full encoder) lacks it; 48×48 with the same encoder stack (matched-compute, 1×1 feature map) has a strong one. We attribute this to encoder-preserved spatial information, with the scaling sweep needed to confirm. (iii) *Attention policy (H3, pending)*: within the rich-encoder regime, we predict that shifting the static gaze location alters what the LSTM encodes; the foveated-shifted causal control is in training at submission time and will test this directly (§4.5).

Beyond the recurrent setting (untested) We use a Wijmans-style recurrent (LSTM) backbone for direct comparability with the emergent-map literature. The three implications above are established only in this setting, but we speculate about where they plausibly extend. Transformer-based embodied agents build their internal state by accumulating information via learned attention, which is structurally similar to our gaze location in the sense that it selects which observation content is read into the representation; our H3 finding therefore suggests a testable prediction for transformer navigators—strong constraints on attended tokens should produce content-level dissociations of the sort we observe—but the prediction has not been verified and is left for future work. The sensor-structure axis (H1, H2) is a claim specifically about RL recurrent navigation agents trained under partial observability; whether equivalent representational-content effects arise in supervised visual

learners, non-visual modalities, or multimodal systems is an open empirical question that the present experiments do not resolve.

5.5 Limitations

(i) *Intermediate checkpoint for foveated-fixed*: all foveated-fixed analyses use ckpt.36 at $\sim 174\text{M}$, the last checkpoint before the silent-corruption window (Appendix C); the other four conditions use the final converged checkpoint. (ii) *Minimal learned-gaze architecture*. Our learned-gaze module is a deterministic sigmoid MLP with slow-gaze rollout updates, chosen for computational compatibility with DD-PPO rollouts and simplicity as a controlled variable, not as a competitive active-gaze model. Its collapse under pure task supervision is specific to this architecture (and so far to a single seed; multi-seed replication is in progress); richer gaze mechanisms (stochastic gaze with entropy bonus, information-seeking intrinsic reward, attention-based token selection) would need separate investigation. We therefore do not claim that implicit task pressure cannot produce active gaze in general; we claim only that our minimal decoder collapsed, and we use the collapsed-gaze position as an empirical starting point for the gaze-location test (foveated-shifted control, §4.5) — a test whose logic does not depend on the gaze decoder working. (iii) *Cross-heading probe*: the proposal included a cross-heading generalisation probe (train on heading A , test on opposite heading). Our implementation returned “insufficient samples” sentinels at our probe sizes and does not appear in our results. (iv) *Foveation model*: Gaussian blur with quadratic falloff is a simplification; log-polar or multi-scale pyramid models may behave differently. (v) *Linear probes only*: non-linear probes might reveal structure Ridge regression misses, particularly for occupancy maps. (vi) *Architecture scope*: we use a recurrent (LSTM) backbone to remain comparable with the emergent-map literature we build on. The core claims (sensor structure and gaze location as representational-content axes) should hold for transformer-based navigation agents whose internal state is built by learned attention, but directly verifying this requires re-running the five-condition ablation on a transformer backbone and is left for future work; see §5 for the specific predictions.

6 Conclusion

We ran a controlled five-condition ablation isolating the effect of visual-input structure on the recurrent representations learned by an otherwise-identical DD-PPO navigation agent. Three findings emerge.

(1) *Encoder–memory race*. LSTM top-layer world-frame GPS encoding is *temporally stable* when the visual encoder cannot resolve world-frame structure on its own, and *transient* when it can. All five conditions encode top-layer GPS in the typical-episode window (steps 50–200, $R^2 \in [0.6, 0.95]$). Blind and matched-compute then maintain that code across episode duration (blind $R^2 = 0.92$ at step 800+; matched $R^2 = 0.68$); uniform / foveated / foveated-learned overwrite it (long-tail R^2 at chance or below). The headline “rich-encoder $\approx$ chance” GPS R^2 in Table 1 is therefore primarily a long-tail statistic, not a mid-episode property. Per-layer probing additionally shows all five conditions have linearly-decodable GPS at Layer 0 (where the sensor embedding enters), and an MLP probe at Layer 2 still recovers position from rich-encoder conditions ($R^2 \in [0.51, 0.73]$); we cannot from MLP recovery alone exclude exploitation of position-correlated features. H1 is most precisely a claim about what a *linear, top-layer, late-episode* probe can read — which is also what the standard DD-PPO action head reads.

(2) *Condition-specific representational format*. Even where rich-encoder LSTMs are spatially pass-through, their hidden-state subspaces are perfectly disjoint: pairwise CKA at the noise floor ($< 10^{-4}$), cross-condition probe transfer catastrophically negative ($R^2 \ll -800$), 1-NN purity 1.000 across pooled samples, and cross-condition memory transplants that drop SPL by 0.19–0.21 between rich-encoder pairs (isolation metric up to -0.17). Whatever non-spatial task-relevant features each LSTM carries, it carries them in its own linear subspace — so much so that a probe trained on one condition predicts off-manifold values on another.

(3) *Gaze-location test (pending)*. Whether gaze location is a second representational-content axis *within* the rich-encoder regime is what H3 asks; the cleanest test (foveated-shifted training with hardcoded gaze at the learned-gaze module’s collapsed location) is in progress at submission and will report in §4.5.

The pattern parallels three threads of animal-navigation research — enhanced hippocampal coupling in congenitally blind humans, the disruption of rodent grid-cell firing in darkness with navigation nonetheless preserved, and cross-modality hippocampal remapping in bats. We use these as analogy only: the convergent direction across systems is suggestive of a sensor–memory coupling principle, but our experiments do not test architectures other than LSTM, so we do not claim the principle is architecture-independent. Open questions: (i) whether the encoder-bottleneck ordering extends to finer-grained resolution sweeps (Appendix E); (ii) whether H3 stands on the clean causal test; (iii) whether transformer-based navigation agents reproduce the encoder–memory race; (iv) whether equivalent effects shape representations in non-navigation embodied tasks; (v) multi-seed replication of all five conditions to bound seed-to-seed variability of the H1 numbers.

References

- Andrei Atanov, Yicheng Fu, Gaurav Singh, Jianren Yu, Andrew Spielberg, and Amir Zamir. How far can a 1-pixel camera go? solving vision tasks using photoreceptors and computationally designed visual morphology. In *European Conference on Computer Vision (ECCV)*, 2024.
- Andrea Banino, Caswell Barry, Benigno Uria, Charles Blundell, Timothy Lillicrap, Piotr Mirowski, Alexander Pritzel, Martin J Chadwick, Thomas Degris, Joseph Modayil, et al. Vector-based navigation using grid-like representations in artificial agents. *Nature*, 557(7705):429–433, 2018.
- Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics*, 48(1):207–219, 2022.
- Guifen Chen, John A. King, Neil Burgess, and John O’Keefe. Absence of visual input results in the disruption of grid cell firing in the mouse. *Current Biology*, 26(17):2335–2342, 2016.
- Christopher J Cueva and Xue-Xin Wei. Emergence of grid-like representations by training recurrent neural networks to perform spatial localization. In *International Conference on Learning Representations (ICLR)*, 2018.
- Kshitij Dwivedi, Gemma Roig, Aniruddha Kembhavi, and Roozbeh Mottaghi. What do navigation agents learn about their environment? In *Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Maya Geva-Sagiv, Sandro Romani, Liora Las, and Nachum Ulanovsky. Hippocampal global remapping for different sensory modalities in flying bats. *Nature Neuroscience*, 19(7):952–958, 2016.
- James Gornet and Matt Thomson. Automated construction of cognitive maps with visual predictive coding. *Nature Machine Intelligence*, 6:820–833, 2024.
- Jacqueline Gottlieb, Pierre-Yves Oudeyer, Manuel Lopes, and Adrien Baranes. Information-seeking, curiosity, and attention: computational and neural mechanisms. *Trends in Cognitive Sciences*, 17(11):585–593, 2013.
- Mary M. Hayhoe and Jonathan S. Matthis. Control of gaze in natural environments: effects of rewards and costs, uncertainty and memory in target selection. *Philosophical Transactions of the Royal Society B*, 373(1718):20170368, 2018.
- Jay A Hennig, Sandra A Romero Pinto, Takahiro Yamaguchi, Scott W Linderman, Naoshige Uchida, and Samuel J Gershman. Emergence of belief-like representations through reinforcement learning. *PLOS Computational Biology*, 19(9):e1011067, 2023.
- John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2019.
- Michael J. Jutras, Pascal Fries, and Elizabeth A. Buffalo. Oscillatory activity in the monkey hippocampus during visual exploration and memory formation. *Proceedings of the National Academy of Sciences*, 110(32):13144–13149, 2013.
- Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *International Conference on Machine Learning (ICML)*, 2019.

- Ron Kupers and Maurice Ptito. Compensatory plasticity and cross-modal reorganization following early visual deprivation. *Neuroscience & Biobehavioral Reviews*, 41:36–52, 2014.
- Volodymyr Mnih, Nicolas Heess, Alex Graves, and Koray Kavukcuoglu. Recurrent models of visual attention. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2014.
- Jiri Najemnik and Wilson S. Geisler. Optimal eye movement strategies in visual search. *Nature*, 434(7031):387–391, 2005.
- Motahareh Pourrahimi and Pouya Bashivan. Emergent brain-like representations in a goal-directed neural network model of visual search. *bioRxiv preprint*, 2025.
- Santhosh Kumar Ramakrishnan, Erik Wijmans, Philipp Krähenbühl, and Vladlen Koltun. Does spatial cognition emerge in frontier models? In *International Conference on Learning Representations (ICLR)*, 2025.
- Edmund T. Rolls. Hippocampal spatial view cells for memory and navigation, and their underlying connectivity in humans. *Hippocampus*, 33(6):533–572, 2023.
- Edmund T. Rolls. Why do primates have view cells instead of place cells? *Trends in Cognitive Sciences*, 28(9):862–874, 2024.
- Ruth Rosenholtz. Capabilities and limitations of peripheral vision. *Annual Review of Vision Science*, 2:437–457, 2016.
- Manolis Savva, Abhishek Kadian, Oleksandr Maksymets, Yili Zhao, Erik Wijmans, Bhavana Jain, Julian Straub, Jia Liu, Vladlen Koltun, Jitendra Malik, Dhruv Batra, and Devi Parikh. Habitat: A platform for embodied AI research. In *International Conference on Computer Vision (ICCV)*, 2019.
- Jinghuan Shang and Michael S. Ryoo. Active vision reinforcement learning under limited visual observability. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Nathan Sprague and Dana H. Ballard. Eye movements for reward maximization. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2003.
- Hans Strasburger, Ingo Rentschler, and Martin Jüttner. Peripheral vision and pattern recognition: A review. *Journal of Vision*, 11(5):13, 2011.
- Amanda C. Tas, Steven J. Luck, and Andrew Hollingworth. The role of attention in the encoding of visual working memory. *Journal of Experimental Psychology: Human Perception and Performance*, 42(8):1121–1138, 2016.
- Erik Wijmans, Abhishek Kadian, Ari Morcos, Stefan Lee, Irfan Essa, Devi Parikh, Dhruv Batra, and Manolis Savva. DD-PPO: Learning near-perfect pointgoal navigators from 2.5 billion frames. In *International Conference on Learning Representations (ICLR)*, 2020.
- Erik Wijmans, Manolis Savva, Irfan Essa, Stefan Lee, Ari Morcos, and Dhruv Batra. Emergence of maps in the memories of blind navigation agents. In *International Conference on Learning Representations (ICLR)*, 2023.
- Sylvia Wirth, Pierre Baraduc, Angélique Planté, Serge Pinede, and Jean-René Duhamel. Gaze-informed, task-situated representation of space in primate hippocampus during virtual navigation. *PLOS Biology*, 15(2):e2001045, 2017.
- Fei Xia, Amir R Zamir, Zhiyang He, Alexander Sax, Jitendra Malik, and Silvio Savarese. Gibson Env: Real-world perception for embodied agents. In *Computer Vision and Pattern Recognition (CVPR)*, 2018.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer:
answerYes

Justification: The abstract and §1 state the three hypotheses H1/H2/H3 with explicit signatures (what pattern supports each), and preview the three findings with specific magnitudes (path-history $R^2=0.56$, CKA <0.004 , compass $R^2=0.94$ vs GPS $R^2=0.03$). These are expanded in §4 and §5 with no stronger claims than the data support.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer:

answerYes

Justification: §5.4 explicitly lists six limitations (i–vi) covering checkpoint staleness, single-seed H3, missing proposed analyses, foveation-model simplification, linear-probes-only, and LSTM-architecture scope.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer:

answerNA

Justification: The paper reports empirical results; it contains no theoretical claims or proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer:
answerYes

Justification: §3.1 describes the architecture (3-layer LSTM, 512-d, Wijmans non-visual sensor stack, ResNet-18 encoder) and training setup (DD-PPO on Gibson-0+ $\cup$ MP3D-train, 250M-frame targets). §3.2–§3.6 describe all probing and cross-condition protocols. Appendix A details the gaze-diversity ablation. Code will be released upon publication.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer:
answerYes

Justification: §3.7 commits to open-source release of the training and probing code. Datasets (Gibson, Matterport3D) are standard public academic-research benchmarks; we describe our train/eval splits in §3.1.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).

- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer:

answerYes

Justification: §3.1 specifies training pool and eval set; §3.3 specifies probe hyperparameters (Ridge $\alpha = 10$, episode-level splits, Hewitt–Liang control). Appendix A gives the ablation’s loss formulation and coefficient.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer:

answerNo

Justification: We report single-seed numbers per condition. This is documented as Limitation (ii) in §5.4. Multi-seed replication for tighter confidence intervals is in the immediate follow-up.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).

- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer:

answerYes

Justification: §3.7 reports per-condition training budget (5–7 days on 2×V100 for 250M frames) and probing budget (<1h per condition on one V100).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer:

answerYes

Justification: The research uses public academic-research benchmarks, publicly-released open-source RL frameworks, and involves no human subjects or deployment. We also report and remediate a numerical-stability bug upstream (Appendix B).

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer:

answerYes

Justification: §5.3 discusses implications — sensor design as a tunable axis for interpretable embodied agents. Negative societal impacts are minimal: the work is fundamental research on simulated indoor navigation with no direct deployment pathway.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer:

answerNA

Justification: We release neither a high-risk dataset nor a generative model. Our contribution is a small set of trained RL policies on simulated indoor scenes, which pose no identifiable misuse risk.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer:

answerYes

Justification: Gibson Database and Matterport3D are used under their released academic-research licenses; Habitat-Sim and habitat-lab are MIT-licensed. Our use complies with their terms; citations §3.1.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.

- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer:

answerYes

Justification: Code will be released with documentation upon publication. Figures are novel; they are released under the paper's license alongside the code. The gaze-diversity auxiliary loss and NaN-sanitisation patch are small standalone modules with docstrings.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer:

answerNA

Justification: The paper involves no crowdsourcing and no human-subjects data.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer:

answerNA

Justification: No human-subjects research; no IRB/equivalent review was required.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer:

answerNA

Justification: LLMs were used only for writing assistance on the manuscript. They are not part of the core research methodology; all empirical results come from standard RL training and linear probing.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.

A Supplementary visualisations




Figure 5: Supplementary visualisations referenced from the main text.




Figure 6: Population-coding indicator and goal-vector probe (referenced from §4.6).

B Gaze-Diversity Auxiliary Loss (H3 Ablation)

Motivation In the first foveated-learned training run the gaze decoder collapsed (per-dimension std $\approx 5 \times 10^{-4}$). We attribute this to two mechanisms: (i) once the sigmoid output saturates, its derivative $\sigma'(x) \rightarrow 0$ attenuates gradient signal; (ii) consistent-gaze views yield more predictable visual features, so the reward-gradient path has no explicit incentive for variety.

Target-variance hinge We add $\mathcal{L}_{\text{div}} = \lambda \cdot \text{relu}(\sigma_{\text{target}}^2 - \text{Var}_{\text{batch}}[g])$ where $g \in [0, 1]^2$, $\text{Var}_{\text{batch}}$ averaged over the two gaze dimensions, $\sigma_{\text{target}}^2 = 0.01$ (std ≈ 0.1), $\lambda = 0.01$. The hinge disappears once variance exceeds the target. We also evaluated an uncapped $-\lambda \cdot \text{Var}[g]$ variant as a negative control.

Pilot 1 (uncapped, job 2840256) After 10M frames, gaze variance exceeded that of a uniform distribution over $[0, 1]^2$ (std $[0.43, 0.42]$ vs. theoretical max 0.289): gaze was essentially random. Task performance collapsed: SPL plateaued at 0.05, running-mean metrics transitioned to NaN at ~ 8 M frames. The uncapped regulariser destroyed the consistent-foveation signal the policy needs.




Figure 7: Per-layer and MP3D companion figures (referenced from §4.2).

Pilot 2 (hinge, job 2842288) Gaze converged to an extreme corner of image space: mean $(0.001, 0.998)$ with std $[0.020, 0.023]$ — well below the target std of 0.1 and therefore with the hinge active throughout. Task SPL plateaued at 0.05 as in pilot 1. The policy apparently found the minimum-task-cost gaze that partially satisfies the diversity pressure: a sigmoid saturated at one corner, producing a constant image-corner view.

Interpretation Both pilots argue that in our deterministic sigmoid-MLP gaze architecture with slow-gaze rollouts, gaze collapse is a pre-requisite for task learning rather than an accident. Task signal alone does not spontaneously produce informative fixation, and an exogenous anti-collapse regulariser moves gaze toward whatever region reduces regulariser loss fastest, which is generically *not* the task-useful region. A genuine H3 test likely requires (i) a stochastic gaze location with entropy bonus, so diversity is sampled at rollout time rather than enforced on a deterministic output; or (ii) a task-aligned intrinsic reward (e.g., information gain about goal-relative position); or (iii) an architectural change separating gaze generation from visual-feature extraction (e.g., attention-based gaze).

C Training Stability: Silent Weight Corruption in DD-PPO

During an initial foveated training run we observed a failure mode that did not manifest as a crash but as slow, silent corruption of the learned weights. After ≈ 170 M frames of stable progress (reward ≈ 10 , SPL ≈ 0.83), one mini-batch on one DDP worker produced a NaN gradient. Because `torch.nn.utils.clip_grad_norm_` does not sanitise non-finite gradients—the global norm becomes NaN, the clip coefficient becomes NaN, and `grad *= NaN` remains NaN—the subsequent `optimizer.step()` wrote NaN into every parameter. DDP all-reduce propagated the corruption across workers. The training loop did not raise: it continued for 2.5 days with reward=NaN and SPL=0 until `torch.multinomial` finally rejected the NaN probability tensor. All checkpoints saved during this window were unrecoverable (90 of 97 parameter tensors entirely NaN).

To isolate the origin of the NaN gradient we re-ran training from the last clean checkpoint with `torch.autograd.set_detect_anomaly(True)` and per-module forward-output hooks. The anomaly traceback pinned the failure to `MseLossBackward0`, i.e. the backward pass of the value-head loss $\frac{1}{2}(V(s) - R)^2$; rare return-vs-value discrepancies overflowed float32 in the squaring step, producing non-finite gradients downstream.

Our mitigation is a minimally invasive patch to `PP0.before_step` that zeroes any non-finite gradient element before clipping. A bad mini-batch becomes a no-op update rather than a catastrophic one. The patch is a no-op on clean training paths. Across the five conditions reported here, no NaN events were observed during full training after deployment, and all final checkpoints contain only finite weights. We reported the underlying issue upstream (habitat-lab issue #2226) with a minimal reproduction and a proposed native fix.

D Foveation experiments: status and design

This appendix documents the four controlled foveation experiments described in §4.4. All four were submitted to the cluster at the time of writing; the table below tracks their state. Numbers will populate the corresponding paragraphs in §4.4 in revision.

Table 3: Foveation experiment status (submission time).

Experiment	Tests	Status
F1–F4 strength sweep ($\sigma_{\max} \in \{2, 4, 12, 20\}$)	encoder–memory race as continuous lever	in training
F3 log-polar foveation	spatial-sampling vs. blur bottleneck	in training
Encoder feature-map probe (matched / uniform / foveated)	where position information is held	in probing
Foveated-shifted (gaze (0.49, 0.62))	gaze-location axis (links to H3)	in training

Why disclose this rather than wait. The four experiments are designed to falsify or refine specific predictions of the encoder–memory race framework, and each has a falsifiable signature. We disclose the design openly so that (i) the present submission is honest about what foveation-specific evidence already exists vs. what is in flight, and (ii) any reader can identify which result would change the paper’s interpretation. The current evidence already places foveation-fixed and foveation-learned within the framework; the experiments above sharpen the “foveation” part of the story specifically.

E Encoder-capacity scaling sweep

[TODO: The scaling sweep trains matched-compute variants at input resolutions $\{32, 48, 64, 96, 128, 192\}$ alongside the existing blind (no encoder) and uniform 256×256 anchors. Probe results on the trained checkpoints will populate Figure 8 and this appendix section. Hypothesis: GPS probe R^2 decreases monotonically with input resolution, from $R^2 \approx 0.95$ at no-encoder to $R^2 \approx 0$ at 256×256 , crossing into the pass-through regime between 48 and 128 pixels.]




Figure 8: Encoder-capacity scaling curve. **[TODO: Pending data.]** x-axis: input resolution in pixels; y-axis: GPS and compass probe R^2 (5-fold CV, $\text{mean} \pm \sigma$).