
How Foveated Vision Shapes Cognitive Maps in Navigation Agents

Anonymous authors

Paper under double-blind review

Abstract

What shapes the content of a recurrent navigation agent’s learned internal state — the task, the architecture, or the structure of the visual input? Prior work has shown that blind PointNav agents spontaneously develop cognitive-map-like representations in their recurrent state under task pressure alone [Wijmans et al., 2023]. We show that the blind result is a special case of a broader principle: *the visual encoder’s information capacity determines whether the LSTM has to carry world-frame spatial structure*. Holding task, optimiser, and LSTM backbone identical and varying only the visual pathway (no vision; a resolution-collapsed encoder; uniform vision; foveated vision at a fixed gaze; foveated vision with a learned gaze), we probe the recurrent state for world-frame position and heading. All five conditions encode top-layer GPS in their LSTM in the typical-episode window. What separates conditions is what happens *later*: when the encoder hands the LSTM minimal spatial-feature variety per step (blind: no encoder; matched-compute: 1×1 feature map), the LSTM *maintains* a persistent, integration-capable code across episode duration; when the encoder hands it richer spatial features (full-resolution encoders, 8×8 feature map), the LSTM *overwrites* the GPS code as those visual features take over. We call this the *encoder–memory race*; it recasts the “cognitive maps emerge without vision” story as “cognitive maps emerge — and persist — when the visual pathway cannot resolve the map on its own.” A second finding cuts across conditions: five agents that solve the task equally well develop recurrent representations in essentially orthogonal linear subspaces, and cross-condition memory transplants incur task cost well above transplant mechanics — sensor structure shapes representational *format*, not only strength. A third question — whether static gaze location within a rich-encoder regime acts as a second content axis — is tested by a foveated-shifted causal control whose training is in progress. The H1 ordering reproduces on a held-out MP3D evaluation, consistent with an agent-internal rather than Gibson-specific mechanism. Three implications, each scoped to the recurrent PointNav setting we study. (i) Encoder capacity correlates inversely with linear, top-layer LSTM spatial encoding; *if* the correlation reflects causation (which the encoder-resolution scaling sweep, in progress, tests directly), it would also be a practical training-time lever for inducing cognitive-map-like internal representations. (ii) “More pixels, richer memory” is the wrong intuition within this setup; the relationship across our five conditions is inverse. (iii) The pattern parallels independent animal-navigation findings — enhanced hippocampal coupling in congenitally blind humans, the disruption of rodent grid-cell firing in the dark (with navigation nonetheless preserved), cross-modality hippocampal remapping in bats — without depending on their biophysics, suggestive of a sensor-memory coupling principle whose generality across architectures (transformers, non-RL learners, non-navigation tasks) we do not test here.


1 Introduction

Recurrent navigation agents trained with deep RL build internal states that support memory, planning, and generalisation, and the content of those internal states is at least as interesting as the task metrics they drive. A natural question is what *determines* that content: architecture, objective, training budget, or visual input? Wijmans et al. [Wijmans et al., 2023] showed that map-like spatial representations emerge from task pressure alone, even in blind agents that receive only goal displacement and proprioception; Dwivedi et al. [Dwivedi et al., 2022] probed sighted agents for related concepts; Banino et al. [Banino et al., 2018] and Cueva & Wei [Cueva and Wei, 2018] demonstrated grid-cell-like structure in path-integration networks. Together these results establish that *some* spatial structure is learned without being designed in. They leave open the complementary question that we take up here: within agents that all achieve high task success, how much of *what* recurrent memory encodes is determined by the structure of the visual input rather than by the task? In biology the answer is clear — primate hippocampal cells code joint place–view–action state because primates sample the scene with a foveated sensor [Wirth et al., 2017, Rolls, 2024], and the same bat produces non-overlapping hippocampal populations in vision versus echolocation [Geva-Sagiv et al., 2016]. We investigate the analogous question in a controlled agent setting where the only free variables are the visual input structure and the gaze that directs it.


Our five-condition design isolates three degrees of freedom: sensor type (blind / uniform / foveated), encoder capacity (uniform’s full ResNet-18 vs. matched-compute’s 1×1 -collapsed feature map), and gaze location (fixed vs. learned). Every condition uses the same DD-PPO learner, the same Gibson-0+MP3D training pool, the same LSTM backbone, and the same non-visual sensor stack (Figure 1). We adopt the LSTM architecture from Wijmans et al. [Wijmans et al., 2023] deliberately: the emergent-map literature we build on is recurrent, and matching that architecture lets our findings be interpretable as statements about *the structure of the task-shaped memory* rather than about the recurrent-versus-transformer design choice. We discuss in §5 whether the logic plausibly extends to transformer-backed or non-navigation systems — an empirical question we do not settle here.

We test three hypotheses. **H1 (encoder–memory race)**. Agents whose visual encoder cannot resolve world-frame position or heading should develop these quantities in their recurrent state; agents with a rich visual encoder should not. Signature: current-state GPS / compass probe R^2 tracks visual-encoder information content monotonically across the five conditions, with persistent decoding at lags of ≥ 10 time-steps. **H2 (sensor structure $\rightarrow$ format divergence)**. Different sensor conditions should produce internal representations in distinct linear subspaces, not scaled versions of a common code. Signature: pairwise CKA near zero, probe weights fail to transfer across conditions while within-condition probes succeed, and cross-condition memory transplants degrade task performance beyond same-condition self-transplants. **H3 (gaze location $\rightarrow$ orthogonal content lever)**. Where a foveated sensor is centred should act as a representational-content axis independent of the sensor transform. We test this with a foveated-shifted control in which gaze is hardcoded to a different image location from the default central gaze; same foveation transform, same training budget, different gaze target. Our minimal learned-gaze module collapsed to a static output (Appendix B), so we use that collapsed location for the shifted-gaze control rather than as evidence about active gaze; the foveated-shifted training is in progress (§4.5). Each hypothesis has a direct analog in the animal-navigation literature [Kupers and Ptito, 2014, Chen et al., 2016, Geva-Sagiv et al., 2016, Jutras et al., 2013, Najemnik and Geisler, 2005], which we return to in Discussion as evidence that the effects reflect general sensor–memory coupling principles rather than architectural artefacts.

Findings H1 is supported as a *temporal-stability* claim, not an existence claim: all five conditions encode top-layer GPS in the typical-episode window, and only the bottleneck conditions (blind, matched-compute) maintain it across episode duration; rich-encoder conditions overwrite it as visual features accumulate. The pattern is preserved on a held-out MP3D evaluation. The encoder feature-map probe (§4.4) sharpens the mechanism: *none* of the three sighted encoders linearly decode GPS from their visual stream, so the H1 divergence is not about encoder-level position retention but about how much spatial-feature variety the encoder hands to the LSTM — and whether that variety is enough to substitute for integrating the GPS-sensor signal across time. H2 is supported by four convergent cross-condition divergence measures: pairwise CKA at the noise floor, probe-transfer failure to off-manifold values, 1-NN purity 1.000, and behaviourally costly cross-condition memory transplants whose 5×5 matrix is asymmetric (blind donations hurt rich-encoder agents far more than vice versa). H3 asks whether gaze location is a second content axis within the rich-encoder



Same Gibson val episode (scene E9uDofAP3SH, ep 414)



(d) The 5 conditions on the same MP3D-train episode (scene E9uDofAP3SH, ep 40422 in the env iterator; same paired index 414 across our probing data), trajectories overlaid on the Habitat-rendered top-down occupancy map. Each condition's SPL on this episode is within ± 0.10 of its per-condition median, so all five are typical successful runs. Bottleneck conditions (blind 231 steps, matched 146) take longer paths through the corridor; rich-encoder conditions (uniform / foveated / fov-learned, 77–80 steps) take near-direct paths along the same corridor.

Figure 1: The five visual conditions differ only in what reaches the agent's encoder, but produce visibly different navigation behaviours. *Inputs:* (a) Blind: no visual input; only the non-visual sensor stack (goal-in-start-frame, GPS, compass, close-to-goal scalar) — the Wijmans [Wijmans et al., 2023] configuration. (b) Uniform: 256×256 full-resolution RGB. (c) Foveated (fixed gaze): eccentricity-dependent Gaussian blur centred on the image middle ($\sigma_{\max} = 8$, quadratic falloff). Two additional conditions not pictured: *matched-compute* (a down-sampled uniform view at 48×48 , whose total pixel budget matches foveated and whose ResNet-18 encoder produces a 1×1 feature map) and *foveated (learned gaze)*, which adds a learned (x, y) gaze offset to (c). *Behaviour:* (d) all five conditions solve the same PointGoal navigation task; the trajectory shape and length differ systematically, previewing the encoder–memory race finding (§4.2).

regime; the clean causal test (foveated-shifted) is in training. A non-linear (MLP) probe at the top layer recovers position from rich-encoder networks ($R^2 \in [0.51, 0.73]$), so H1 is most precisely a claim about *linear, top-layer* GPS encoding (i.e. what the linear DD-PPO action head can use); we cannot from MLP recovery alone exclude exploitation of position-correlated features (wall distance, scene identifiers) rather than position itself.

Contributions (i) A controlled five-condition ablation that decouples “no vision” from “visual-encoder bottleneck” via the matched-compute condition, making the encoder–memory race claim falsifiable. (ii) Convergent evidence (linear probes, MLP probes, CKA, probe transfer, 1-NN purity, memory transplant, shortcut discovery, MP3D generalisation, per-layer probes) for two of the three hypotheses, with a clean methodology disclosure including the deterministic-rollout protocol that fixes a probe-collection confound documented in §3.3. (iii) A mechanistic re-reading of Wijmans et al.’s blind-agent finding: cognitive maps emerge not because vision is absent but because the visual encoder cannot resolve the map, predicting an encoder-capacity scaling curve we begin to test in Appendix E.

2 Related Work

Cognitive maps in RL navigation agents Wijmans et al. [Wijmans et al., 2023] demonstrated that map-like representations emerge in the recurrent memories of DD-PPO agents even without visual input; Dwivedi et al. [Dwivedi et al., 2022] probed sighted agents for related spatial concepts; Banino et al. [Banino et al., 2018] and Cueva & Wei [Cueva and Wei, 2018] found grid-cell-like units in recurrent networks trained on path integration; Gornet & Thomson [Gornet and Thomson, 2024] extended this via visual predictive coding. Hennig et al. [Hennig et al., 2023] showed that RL under partial observability can induce belief-state-like representations without explicit probabilistic objectives. Ramakrishnan et al. [Ramakrishnan et al., 2025] examined spatial cognition in frontier models. All these studies use either absent or spatially uniform perception; none vary input quality spatially, and none test whether different sensors produce different representational formats.

Foveated vision and gaze policies in deep learning Foveated rendering [Strasburger et al., 2011, Rosenholtz, 2016] exploits eccentricity-dependent acuity falloff for efficiency. Learned glimpse policies appear in recurrent attention models [Mnih et al., 2014]; Atanov et al. [Atanov et al., 2024] showed that sensor layout changes downstream task performance. Pourrahimi & Bashivan [Pourrahimi and Bashivan, 2025] probed a foveated visual-search model for neuroscience-predicted features; Shang & Ryoo [Shang and Ryoo, 2023] jointly learned gaze and motor policies following the Sprague & Ballard [Sprague and Ballard, 2003] framework. These address visual search or task performance; none asks how sensor structure reshapes the learned spatial memory of a downstream navigation policy.

Probing neural representations Belinkov [Belinkov, 2022] surveys probing classifiers; Hewitt & Liang [Hewitt and Liang, 2019] introduced control tasks to distinguish genuine structure from probe memorisation; Kornblith et al. [Kornblith et al., 2019] introduced linear CKA for cross-model similarity. We adopt linear probing, the Hewitt–Liang selectivity metric, and unaligned linear CKA, and complement them with two behavioural interventions directly on the recurrent state (memory transplant and shortcut discovery) to validate the probe-based findings outside the linear-probe framework.

Biological precedent Our hypotheses and framing draw on three threads of animal-navigation research. (i) Primate hippocampal cells code joint place–view–action state [Wirth et al., 2017, Rolls, 2023]; Rolls [Rolls, 2024] argues that primate *spatial-view cells* and rodent *place cells* reflect the different sampling statistics of foveated versus near-panoramic vision. (ii) The same bat in the same environment produces non-overlapping hippocampal populations when it switches between vision and echolocation [Geva-Sagiv et al., 2016]—format-level divergence across sensor modalities within a single animal. (iii) In primates, each saccade resets hippocampal 3–8 Hz rhythms and the reliability of the reset predicts subsequent memory [Jutras et al., 2013]; only overt (not covert) attention automatically writes to visual working memory [Tas et al., 2016]. Gaze targets themselves are selected in an information-gain-maximising fashion consistent with ideal-observer predictions [Najemnik and Geisler, 2005, Gottlieb et al., 2013, Hayhoe and Matthis, 2018]. We use these results only to frame our ML predictions and interpret our findings; our contribution is the

controlled ablation in a learned agent that cross-species and cross-modality biological comparisons cannot cleanly supply.

What is new here Three things distinguish this work from the studies above. (i) A *controlled* five-condition ablation with everything but the visual pathway held identical, including the matched-compute condition that decouples “no vision” from “visual encoder bottleneck” — the lever that makes the encoder-memory race claim falsifiable. (ii) A protocol-corrected probing methodology: deterministic action selection at probe time (the policy used by every other component of our eval pipeline) rather than stochastic sampling, which we found inflates probe R^2 for high-entropy policies via a 4-step quasi-static-trajectory artefact (§3.3, Sampling protocol). (iii) Convergent behavioural validation of the probe-based H1 and H2 findings via memory transplant and shortcut discovery, both of which sit *outside* the linear-probe framework.

3 Methods

3.1 Agent Architecture and Training

All agents are trained on PointGoal navigation in Habitat [Savva et al., 2019] using DD-PPO [Wijmans et al., 2020] on a merged pool of Gibson-0+ [Xia et al., 2018] (411 scenes) and Matterport3D train (61 scenes), totalling 472 training scenes. Evaluation uses 18 held-out MP3D test scenes (1008 episodes).

All agents share an identical 3-layer LSTM (512-d) backbone and the Wijmans non-visual sensor stack: goal-in-start-frame, GPS, compass, and a close-to-goal scalar, each projected to 32-d and concatenated with a 32-d previous-action embedding. All sighted conditions are trained for a target of 250 M environment frames; the blind baseline is trained to 342 M frames to allow slower convergence from proprioception alone.

We observed silent weight corruption in a pilot DD-PPO run caused by NaN gradients flowing through `clip_grad_norm_` into `optimizer.step()`; we mitigate this with a gradient-sanitisation patch to `PP0.before_step`, described in Appendix C.

3.2 Visual Conditions

Five conditions vary only the visual input (Figure 1):

1. *Blind*: no visual input; non-visual sensor stack only, replicating Wijmans et al. [Wijmans et al., 2023].
2. *Uniform*: full-resolution 256×256 RGB encoded by ResNet-18.
3. *Foveated (fixed)*: eccentricity-dependent Gaussian blur ($\sigma_{\max} = 8$, quadratic falloff) with gaze locked at image centre; same ResNet-18 encoder.
4. *Matched-compute*: 48×48 uniform RGB. At this resolution, ResNet-18’s 32-fold spatial downsampling collapses the feature map to 1×1 , removing the spatial dimensions of the encoder’s output (channel information is preserved as a flat 512-d vector). **[TODO: Verify whether channel info alone in matched-compute encodes any usable position signal — a small probe on the encoder feature map (not LSTM state) would settle this.]**
5. *Foveated (learned)*: same blur model, but gaze centre predicted by a lightweight MLP trained end-to-end with the navigation policy.

The matched-compute condition is central to H1: its 1×1 encoder feature map provides a visual observation while leaving the encoder unable to resolve world-frame position from the current frame, separating the “no vision” account from the “encoder bottleneck” account. We attribute its strong LSTM GPS encoding to the feature-map collapse specifically, but the 48×48 input also has lower spatial resolution than uniform’s 256×256 , so we cannot from this single condition fully separate “ 1×1 feature-map collapse” from “low input resolution.” The encoder-resolution scaling sweep (Appendix E, in progress) varies input resolution at fixed encoder stack to test this attribution directly.

Implementation note The foveated conditions (fixed, shifted, learned) disable the running-mean-and-variance input normaliser used by uniform and matched-compute, to avoid an in-place autograd conflict along the gaze-decoder gradient path. All reported results were obtained with this setting; we have not directly verified that re-enabling the normaliser would leave the H1 ordering unchanged.

[TODO: Verify normaliser invariance: rerun foveated-fixed with the normaliser re-enabled (the autograd conflict is gaze-decoder-specific, so feasible for fov-fix), and confirm the H1 GPS R^2 does not move.]

3.3 Probing Methodology

After training, we freeze all weights and collect per-step LSTM hidden states (top-layer $\mathbf{h}$, 512-d) paired with ground-truth spatial labels from $n=500$ Gibson held-out validation episodes (the split not used in training). Rollouts use *deterministic* (argmax) action selection, matching the eval protocol of our shortcut-discovery, memory-transplant, and standard habitat-baselines evaluators. Under this protocol the rollout lengths span roughly ~ 100 – 400 mean steps depending on condition (longer for blind, shorter for high-SPL conditions), so we use 5-fold episode-level CV to keep the train/test splits balanced per scene rather than length-matching across conditions. **[TODO: A controlled length-matching ablation (truncate every condition to a common length, re-run probes) is left for follow-up; we expect it to leave the H1 ordering unchanged because the bottleneck conditions decode at high R^2 even at the smallest sub-sample, but have not verified it.]** We train Ridge regression probes ($\alpha = 10$) with episode-level 5-fold cross-validation, reporting mean $\pm \sigma$ across folds, and verify probe specificity with Hewitt–Liang label-permutation selectivity [Hewitt and Liang, 2019].

Probed variables: GPS position (episodic), compass heading (episodic), distance-to-goal, lag- k path history (target value at $t - k$ decoded from $\mathbf{h}_t$, $k \in \{0, \dots, 20\}$), visited-region occupancy, full occupancy grid (20×20), per-unit spatial information (place-cell analysis), and the ego-frame goal vector $\mathbf{R}(-\theta)(\mathbf{g} - \mathbf{p})$ decomposed into scalar goal-distance and 2-D direction.

Sampling protocol An earlier draft of this work used stochastic (policy-distribution) action selection for probe rollouts. Under stochastic sampling, conditions with higher action-entropy (specifically the rich-encoder conditions foveated, uniform, matched, blind under their trained stochastic policies) sampled the STOP action early with probability ~ 0.25 per step, producing mean episode lengths of ~ 4 steps and target variance two orders of magnitude below the episodic range. The resulting probe R^2 values were trivially high across all conditions (any predictor fits a near-constant target). This confounded cross-condition comparisons, and we switched to the deterministic protocol used here and elsewhere in the probe/eval pipeline. All results in this paper use the deterministic protocol.

For foveated-fixed, all reported analyses (probes, behavioural, eval) use checkpoint ckpt.36 at $\sim 174\text{M}$ frames, the last clean intermediate before the silent-corruption window (Appendix C).

3.4 Cross-Condition Analysis

To test H2, we employ two complementary methods. *Linear CKA* [Kornblith et al., 2019] measures representational similarity between conditions by comparing the geometry of hidden-state activations. Values near 1 indicate identical geometry; near 0 indicates unrelated representations. *Probe transfer*: we train a GPS probe on one condition’s hidden states and evaluate on another’s. If spatial information is encoded in a shared linear subspace, cross-condition transfer should succeed. Catastrophic failure indicates fundamentally different encoding formats. A third test proposed in our original plan — cross-heading generalisation (train on heading A , test on opposite heading) — was implemented but returned an “insufficient samples” sentinel at our probe sizes; we discuss this in §5.5.

3.5 Behavioural Probes

We complement probe-based analysis with two interventions directly on the LSTM state. *Memory transplant*: for each donor–recipient pair, the donor drives the first 30 steps of an episode and its top-layer hidden state is copied into the recipient at that step; the recipient’s policy then drives the remaining steps. We compare *baseline* (recipient alone, full episode), *self-transplant* (recipient as its own donor, controls for transplant mechanics), and *cross-transplant* (different donor). Midpoint 30 was chosen because Gibson PointNav episodes average ~ 100 – 120 steps; a later midpoint left too few steps after transplant. We use 150 episodes per pair per condition, sampled uniformly from the held-out Gibson validation split. *Shortcut discovery*: for each condition, we run episodes in pairs of two using the same scene but different goals, sampling 20 scenes with 10 paired episodes per scene (200 paired runs per condition per setting). Per-condition SPL is reported as the mean of the 20 scene-level mean SPLs (prevents high-episode-count scenes from dominating). We compare *reset* (LSTM

zeroed between the two episodes) vs. *persistent* (LSTM state carried over) on the second-episode SPL. A reduction under the persistent setting measures how tightly the stored representation is locked to the previous goal. For the foveated-fixed condition, all probes and behavioural runs use the ckpt.36 checkpoint ($\sim 174\text{M}$ frames, the last clean intermediate before the silent-corruption window, §4).

4 Results

4.1 Five conditions, same task

Table 1 collects the headline per-condition numbers this section returns to. All sighted conditions solve PointGoal at 96–99% success; blind reaches 93% at a longer training budget, reproducing Wijmans et al. [Wijmans et al., 2023]. Success is tightly bunched, so “the agents differ in what they do” is an unlikely explanation for the encoded-content gap we report (it does not strictly rule it out). The probe machinery itself is calibrated: distance-to-goal decodes robustly across all five conditions ($R^2 \in [0.81, 0.90]$; Appendix A for the companion plot), confirming that Ridge has the capacity to fit high- R^2 scalars when the hidden state contains them; and Hewitt–Liang label-permutation selectivity is high in the conditions where the GPS code itself is high (blind GPS / compass selectivity $+0.95 / +0.81$; matched $+0.72 / +0.56$), consistent with a genuine signal rather than label memorisation. Two training-stability caveats are disclosed in full in Appendix C: foveated-fixed uses ckpt.36 at $\sim 174\text{M}$ frames (last clean checkpoint before a silent NaN-gradient corruption at $\approx 182\text{M}$); and matched-compute’s 1×1 feature-map collapse, originally noticed as a training quirk, is what makes matched-compute a useful test of the encoder-bottleneck framing (we cannot from one condition alone fully separate the 1×1 feature-map effect from the lower input resolution; see §3.2 and §4.4).

Table 1: Per-condition headline. Probe R^2 columns are the 5-fold episode-level CV mean $\pm$ std on full-length deterministic-rollout trajectories (no per-episode step cap). *GPS* and *Compass* are world-frame. The bottleneck-vs-rich-encoder gap is in temporal *stability* of the code, not its mid-episode existence (Figure 2b): all five conditions decode top-layer GPS at $R^2 \in [0.6, 0.95]$ in the typical-episode window (steps 50–200); the rich-encoder long-tail collapse is what produces the sub-chance numbers below. Bold marks the H1-relevant ordering (blind $>$ matched $\gg$ rich-encoder).

Condition	Frames	SPL	Success	GPS R^2 (no-cap)	Compass R^2 (no-cap)
Blind	342M	0.59	0.93	$+0.95 \pm 0.02$	$+0.81 \pm 0.08$
Matched-compute	250M	0.67	0.98	$+0.78 \pm 0.10$	$+0.64 \pm 0.10$
Uniform	250M	0.85	0.99	-0.31 ± 0.86	$+0.36 \pm 0.23$
Foveated (fix)	174M	0.78	0.96	$+0.06 \pm 0.88$	$+0.07 \pm 0.69$
Fov-learned	250M	0.82	0.98	-2.43 ± 3.98	-1.34 ± 3.14

4.2 Encoder–memory race: a temporal-stability claim (H1)

Finding Visual-encoder capacity inversely tracks LSTM spatial encoding across our five conditions: bottleneck conditions (blind, matched) encode GPS / compass at the top layer; rich-encoder conditions (uniform, foveated, foveated-learned) decode at chance (Figure 2a, Table 1). DtG decodes robustly in every condition because every agent must track it. **[TODO: Whether the correlation reflects causation is what the §4.4 experiments and the Appendix E resolution sweep test; multi-seed gap-fills are in progress.]**

Temporal stability is the precise claim The no-cap probe conflates “does the LSTM ever encode GPS” with “does it maintain the encoding across episode duration.” Step-binned out-of-fold probing separates them (Figure 2b): all five conditions encode top-layer GPS in the typical-episode window; rich-encoder conditions then *overwrite* the code (chance by step 400–800, deeply negative in the long-tail bin), while bottleneck conditions hold it stable to step 800+. The Table 1 “rich-encoder $\approx$ chance” headline is therefore long-tail behaviour, not mid-episode. *H1 is most precisely a temporal-stability claim about the linearly-decodable top-layer GPS code.* A lag- k probe corroborates: blind GPS decodes at $R^2 = 0.92$ at $k = 20$, matched > 0.5 ; rich-encoder conditions are at chance for every $k \leq 10$.




Figure 2: H1 evidence in three panels. (a) GPS R^2 (5-fold CV, Ridge $\alpha = 10$, full-length deterministic rollouts) on Gibson and held-out MP3D (300 episodes/condition; same checkpoints, no re-training). Blind and matched-compute (bottleneck conditions, blue background) decode world-frame GPS robustly on both datasets; uniform / foveated / foveated-learned (rich-encoder, grey background) do not, except for foveated recovering to a mild positive ($R^2 = +0.35$) on MP3D. The Compass and DtG numbers are in Table 1 (compass shows the same ordering at smaller magnitude; DtG decodes everywhere because every agent must track it; Appendix A for the companion bar plots). (b) Top-layer GPS R^2 binned by step-in-episode (out-of-fold prediction): all five conditions encode top-layer GPS in the typical-episode window (steps ~ 25 –200); the rich-encoder conditions then *overwrite* the code (negative R^2 in the long-tail 800+ bin), while bottleneck conditions hold $R^2 > 0.7$ to step 800+. The headline “rich-encoder $\approx$ chance” in panel (a) is therefore primarily long-tail behaviour. (c) Per-LSTM-layer GPS R^2 (single-fit on top-layer hidden states): all five conditions have GPS at Layer 0 (where the GPS sensor embedding enters the LSTM input vector — this lower-bound is at least partly a sensor-pass-through effect, not the network “computing” position); the divergence is at Layer 2, the policy readout.

Layer 0 sensor pass-through is shared; the divergence is at Layer 2 All five conditions decode GPS at Layer 0 (Figure 2c) because the GPS sensor embedding enters the LSTM input vector there directly. The H1 distinction is what happens through Layers 1–2: bottleneck conditions retain the code to the top layer, rich-encoder conditions discard it. The standard DD-PPO action head is a linear projection from $\mathbf{h}_2$, so “linearly decodable at Layer 2” maps to “policy-accessible”. (*Side observation: blind’s Layer 1 R^2 is lower than Layer 0 or Layer 2; not yet explained.*)

Non-linear probe at the top layer recovers position A 2-layer MLP on rich-encoder top-layer states recovers GPS at $R^2 \in [0.51, 0.73]$. The bottleneck-vs-pass-through divergence is therefore not about whether any probe can recover position — it is specifically about whether a *linear, top-layer* probe can. We cannot from MLP recovery alone exclude exploitation of features merely correlated with position (wall distance, scene identifiers, episode-phase); §4.4 and Appendix E test the bottleneck framing more directly.

MP3D dataset shift The bottleneck-vs-pass-through partition is preserved on held-out MP3D (Figure 2a, MP3D bars). Two single-seed cross-dataset shifts are worth flagging: foveated GPS recovers from chance to $+0.35$ on MP3D, and foveated-learned compass shows a $+1.75$ Gibson→MP3D swing. Both are candidates for Gibson-specific feature use by the foveated encoder; multi-seed replication needed before attributing to a specific mechanism. Compass numbers in Appendix A.

Proposed mechanism We *propose* the following account, consistent with the data but not directly tested by a causal intervention. In PointNav, GPS and compass enter the LSTM input at Layer 0 directly, via a 32-d sensor embedding concatenated with the visual encoder’s flattened output (Figure 2c shows GPS at Layer 0 across all five conditions, $R^2 > 0.91$). The encoder feature-map probe (§4.4, Figure 4) shows that none of the three sighted encoders linearly re-derive GPS from their visual stream: matched $R^2 = -3.14$, uniform -0.32 , foveated -0.65 . The encoder-memory race is therefore not about whether the encoder *re-derives* position; it is about how much *visual feature variety* the encoder hands to the LSTM, and whether that variety is enough to substitute for integrating the Layer-0 GPS-sensor signal across time. Bottleneck conditions (blind, no encoder; matched-compute, 1×1 spatial output $\rightarrow$ minimal feature variety per step) leave the LSTM with little visual structure to substitute, so it integrates the Layer-0 GPS signal across time and *maintains* a top-layer linearly-readable GPS code. Rich-encoder conditions (uniform / foveated / foveated-learned, 8×8 spatial output $\rightarrow$ many position-correlated visual features) give the LSTM enough visual variety

to base actions on those features instead, and the top-layer GPS code drifts off as the episode progresses. Task success is tightly bunched, so a skill-gap artefact is unlikely to explain the encoding gap. The temporal dynamic in Figure 2b (rich-encoder GPS code emerging early at Layer 0 then decaying through Layers 1–2 to chance at the top layer) is the direct fingerprint of this account. **[TODO: A direct causal test of the mechanism — ablating the visual stream mid-rollout in rich-encoder agents to see whether the top-layer GPS code re-emerges, or perturbing the GPS sensor input mid-rollout in bottleneck agents to see whether SPL collapses — is left for future work.]**

Relation to Wijmans et al. Wijmans et al. [Wijmans et al., 2023] showed that blind PointNav agents develop spatial representations in LSTM memory. Our blind result ($R^2 = 0.95$) replicates this; matched-compute extends it to a condition that has visual input but a 1×1 spatial encoder output. The shared trigger across blind and matched is therefore the LSTM having minimal visual feature variety per step — not absence of vision per se, and not the encoder’s ability to resolve position (no sighted encoder does, §4.4). §4.4 and Appendix E test the spatial-feature-variety reading more directly.

Bio parallel The pattern parallels the animal-navigation literature without depending on its mechanisms: congenitally blind humans develop enhanced hippocampal coupling and outperform sighted controls on non-visual spatial tasks [Kupers and Ptito, 2014]; rodent grid cells lose their hexagonal periodicity in darkness even with head-direction input intact [Chen et al., 2016], yet the animal still navigates — so spatial bookkeeping that the (vision-dependent) grid code can no longer carry has to come from somewhere else. We use these only as analogy: in both cases visual-pathway absence shifts spatial bookkeeping into a non-visual substrate, the same qualitative direction we measure in LSTM hidden states. We do not claim the agent and biological mechanisms are the same.

4.3 Format-level divergence across conditions (H2)




Figure 3: H2 convergent evidence, two views. *Left*: Unaligned linear CKA [Kornblith et al., 2019] between top-layer LSTM hidden states ($n=95,640$ steps per condition under deterministic-rollout probes); every off-diagonal is below 10^{-4} . *Right*: 5×5 cross-condition memory-transplant matrix at midpoint $t=30$. Cells off-diagonal report cross-transplant SPL minus the recipient’s self-transplant SPL (the isolated condition-mismatch component, with rollout-divergence mechanics subtracted); diagonal cells report self-transplant SPL minus baseline (the rollout-divergence noise floor, ≈ -0.04 to -0.06). The matrix is asymmetric: blind $\rightarrow$ uniform incurs -0.38 SPL while uniform $\rightarrow$ blind is essentially neutral ($+0.02$); foveated-learned $\rightarrow$ uniform -0.29 , blind $\rightarrow$ foveated -0.19 . *Pairs that are indistinguishable to linear similarity (CKA noise floor) differ markedly under transplant.* Cells with ‘—’ are not yet measured (matched as recipient, jobs in flight). A midpoint-sweep view of three representative pairs across midpoints $\{0, 15, 30, 60, 90\}$ confirming the noise-floor and stability arguments is in Appendix A (Figure 8).

Behavioural test (memory transplant) The most direct test of “do these representations matter?” is to paste a donor’s mid-episode hidden state into a recipient’s policy at step 30 and let the recipient drive the rest (Figure 3, right). The isolated condition-mismatch cost — cross-transplant SPL minus

Table 2: Probe-transfer GPS R^2 on deterministic-rollout hidden states: train on row condition, test on column. Self-accuracy diagonal is 0.89–0.99; every cross-condition transfer is catastrophically negative ($R^2 \ll -800$, often $\ll -10^4$), reflecting that a probe trained on one condition’s hidden-state geometry does not just degrade on another — it predicts off-manifold values.

Train \ Test	Blind	Uniform	Fov-fix	Fov-lrn	Matched-compute
Blind	+0.99	−844	−11,301	−10,731	−55,706
Uniform	−1,383	+0.90	−7,046	−3,305	−110,317
Fov-fix	−2,091	−2,792	+0.92	−1,263	−39,100
Fov-lrn	−6,722	−2,423	−6,205	+0.89	−38,842
Matched-compute	−10,621	−4,832	−4,984	−58,169	+0.96

recipient’s self-transplant SPL — shows three patterns. *Asymmetry*. Blind→uniform incurs −0.38 SPL while uniform→blind is neutral within self-noise; bottleneck-donor states hurt rich-encoder recipients far more than the reverse. *Recipient ranking*. Uniform suffers most from *any* foreign donor (column mean ~ -0.20); blind least (~ -0.02 to -0.05). Blind tolerates foreign donors; rich-encoder recipients do not. *Order-of-magnitude probe-similarity blind spot*. Pairs that are indistinguishable under linear similarity ($\text{CKA} < 10^{-4}$ across the board) differ by 0.01–0.38 SPL under transplant. **[TODO: 4 matched-recipient cells of the 5×5 are queued; asymmetry will be tested for completeness on landing.]**

Convergent linear-similarity evidence Three further measures point to the same “different format” conclusion. *Pairwise CKA* [Kornblith et al., 2019] at the noise floor across all condition pairs (Figure 3, left; within-condition split-half $\text{CKA} \approx 1$, so the near-zero off-diagonals are not estimator noise). *Cross-condition probe transfer* (Table 2) is catastrophically negative in every off-diagonal cell, so a probe trained on one condition’s hidden states predicts off-manifold values on another. *1-NN purity*: among $1500 \times 5 = 7500$ pooled top-layer states, every sample’s Euclidean nearest neighbour shares its condition (purity 1.000 vs. chance 0.20); a t-SNE projection (Appendix A) confirms this visually. **[TODO: 1-NN at larger sample size and under domain shift (MP3D / fov-shifted) is a further-investigate item.]**

Ruling out alternatives Probe capacity is not the bottleneck (self-probes succeed everywhere). Sample-size asymmetry is ruled out by common-sample truncation. Task competence is bunched (96–99%), so the divergence is not a skill difference. Probe transfer fails symmetrically, ruling out one condition’s code simply embedding another’s.

Boundaries Our similarity tests are linear or near-linear; non-linear measures could reveal higher-order shared structure. Our five sensors are diverse but not exhaustive; hybrid sensors (sharp fovea + coarse periphery, depth-only) might produce intermediate geometries.

4.4 Foveation: convergence on H1, divergence on H2 + behaviour

Foveated-fixed and uniform fall in the same rich-encoder pass-through regime on the headline H1 metric (top-layer GPS R^2 at chance), but they are not interchangeable on H2 / behavioural / dataset-shift axes (Table 3). The take-away: *within the rich-encoder regime defined by H1, foveation and uniform produce internally distinct representations whose differences only the H2-and-beyond machinery makes visible*.

The four in-flight experiments below test *which* aspect of foveation drives these differences — sensor transform strength, spatial-sampling vs. blur, gaze location, or encoder feature-map representation. Numbers will populate the sub-paragraphs in revision; current status in Appendix D.

Foveation strength sweep (F1–F4, in progress) We train four additional foveated agents at $\sigma_{\max} \in \{2, 4, 12, 20\}$ (alongside the existing $\sigma_{\max} = 8$), holding everything else fixed — gaze location, encoder stack, training budget, training pool. The expected signature: at low $\sigma_{\max}$ the foveated agent should look like uniform (encoder still resolves world-frame); at high $\sigma_{\max}$ it should approach matched-compute (encoder cannot resolve world-frame, LSTM compensates). The R^2 vs. $\sigma_{\max}$ curve directly tests the encoder–memory race as a continuous lever, not a binary one.

Spatial sampling vs. blur (F3 log-polar control, in training) Gaussian blur removes high-frequency content but preserves *spatial layout*: a $\sigma_{\max} = 8$ foveated image still feeds the encoder

Table 3: Foveation vs. uniform at $\sigma_{\max} = 8$ Gaussian blur with the same ResNet-18 encoder. Same on H1, different on H2 / behavioural / generalisation. Sources: Table 1, Figure 3, Figure 5, Figure 4. Data are deterministic-rollout single-seed.

Measure	Foveated (fix)	Uniform	Verdict
LSTM GPS R^2 (Gibson)	$+0.06 \pm 0.88$	-0.31 ± 0.86	both chance, <i>indistinguishable</i>
Encoder $\rightarrow$ GPS R^2	-0.65 ± 0.28	-0.32 ± 0.08	both negative, overlapping
LSTM compass R^2 (Gibson)	$+0.07 \pm 0.69$	$+0.36 \pm 0.23$	uniform $\sim 5\times$ better
LSTM GPS R^2 (MP3D, held out)	$+0.35 \pm 0.27$	-0.36 ± 1.38	differ by 0.71
Shortcut SPL drop %	21%	41%	uniform $2\times$ more brittle
Memory-transplant cost (fov $\rightarrow$ uni)	-0.10 SPL beyond self-noise		not interchangeable
Pairwise CKA(fov, uni)	$< 10^{-4}$		near-orthogonal subspaces
1-NN purity (pooled)	1.000 vs. chance 0.20		linearly separable

256 $\times$ 256 pixels, so the ResNet-18 stack still produces an 8 $\times$ 8 spatial feature map. A log-polar foveation [Strasburger et al., 2011, Rosenholtz, 2016] resamples the image with a non-uniform retinal grid (e.g. 64 $\times$ 64): the encoder’s spatial output drops to $\sim 2\times 2$, much closer to matched-compute’s 1 $\times$ 1 than to uniform’s 8 $\times$ 8. This is the architectural commitment “foveated vision” makes in biology, and it isolates the variable our sharpened H1 mechanism (§4.4, encoder feature-map paragraph) identifies as the trigger: *spatial-feature variety in the encoder output*.

Predictions, falsifiable. If the encoder–memory race is driven by encoder spatial output dimensionality (not by “bottleneck of any kind”), then log-polar should produce an LSTM GPS R^2 *between* matched-compute (current +0.78) and uniform (current ≈ 0), likely ≥ 0.3 , with corresponding behavioural shifts (longer paths than uniform, larger shortcut SPL drop closer to blind / matched). If instead log-polar matches uniform’s near-chance H1 readout, the mechanism is not pure encoder-output-dimensionality and we need to reframe (e.g. allow the encoder to learn task-useful features from any spatial layout, rendering “bandwidth” the wrong axis). Either outcome is informative; F3 is in training ($\sim 14\text{M}$ / 250M frames at submission, ETA 2–3 days). Numbers will populate this paragraph in revision; the prediction was written before the result was available.




Figure 4: Encoder feature-map probe (post ResNet-18, pre-LSTM, 5-fold CV Ridge). Encoder dim $d = 2048$ (compressed flatten) for all three sighted conditions. None of the three encoders linearly decode GPS (matched $R^2 = -3.14$, uniform -0.32 , foveated -0.65); the encoder–memory race is therefore not about encoder-level GPS information availability. Compass shows a small positive on uniform / foveated ($+0.11$, $+0.13$); matched compass is also chance.

Encoder feature-map probe: where world-frame info actually sits Probing the visual encoder’s output directly (post ResNet-18, pre-LSTM, via forward-hook on the compression layer) gives a striking result: *none of the three sighted encoders linearly decode GPS* (Figure 4). Compass also fails

or is well below ceiling. *This sharpens the H1 mechanism*: the divergence happens inside the LSTM (or at the GPS-sensor pass-through entering Layer 0), not at the visual encoder. Bottleneck conditions integrate the GPS-sensor input across time because their visual stream gives the LSTM minimal feature variety to substitute; rich-encoder conditions have richer (though still position-blind) visual features that the LSTM uses instead, overwriting the Layer-0 GPS code as the episode progresses. Foveated does not differ from uniform on this measure (overlapping uncertainty), so its pass-through behaviour is consistent with “rich-encoder” rather than a foveation-specific effect at the encoder level.

Foveated-shifted (gaze location, in progress) Gaze location is treated separately as H3 (§4.5), but the foveated-shifted control directly informs the foveation slot too: training a foveated agent with hardcoded static gaze at (0.49, 0.62) matches foveated-learned’s collapsed gaze location while removing the learned-gaze module from the optimisation path. If foveated-shifted does *not* reproduce foveated-learned’s specific behaviours (Gibson–MP3D compass swing, transplant ranking), those behaviours can be attributed to the optimisation-path interaction and not to the static gaze location alone.

Why this slot exists Our title and framing focus on foveated vision; the present five conditions place foveated-fixed and foveated-learned within a framework but do not isolate foveation-specific effects (under simple Gaussian blur both behave like uniform). The four experiments above target specific foveation-specific predictions of the encoder–memory race framework. Each is a falsifiable test, designed and submitted at the time of writing. None depends on the others to be informative.

What we already conclude From existing data (Table 3), foveation under $\sigma_{\max} = 8$ Gaussian blur is *not a copy of uniform*: it converges with uniform on H1 but diverges on H2 (subspace orthogonality, 1-NN separability), behavioural memory reliance (shortcut drop), transplant compatibility, and dataset-shift recovery. The encoder feature-map probe further shows the encoder discards linear position regardless of condition — so *whatever foveation does differently from uniform, it does it inside the LSTM in subspaces invisible to a linear-GPS probe*. The in-flight experiments below probe *what* that “something different” is.

4.5 Gaze location as a second axis (H3)

H1 and H2 identify sensor structure as one axis on representational content. H3 asks whether gaze *location* is a second, dissociable axis. We test the gaze-location form (two foveated agents with gaze centred at different static positions should produce different content); the active-gaze form is out of scope of our minimal gaze decoder.

The learned-gaze module collapses Our minimal learned-gaze implementation (deterministic sigmoid MLP, slow updates per 128-step segment; Appendix B) freezes at (0.49, 0.62) with per-dimension std $\leq 2.5 \times 10^{-4}$ after 250M frames; anti-collapse regularisers restore variance only at task-cost. We do not claim active gaze cannot be learned — our decoder is minimal by design — but we use the collapsed gaze position as the empirical anchor for the gaze-location test: the trained agent is *effectively foveated with a static gaze 30 pixels below image centre*.

Gaze-location test (foveated-shifted, in progress) Both rich-encoder foveated variants sit in the pass-through regime under current-state probing (Table 1), so the H3 test is whether *within* the pass-through regime, gaze location shifts content along other axes — non-spatial feature usage, behavioural memory-reliance (shortcut, Figure 5), or transplant compatibility. To isolate gaze location from training-dynamics confounds, we train a foveated-shifted control: an agent identical to foveated-fixed except gaze is hardcoded at (0.49, 0.62) instead of (0.5, 0.5) — matching foveated-learned’s collapsed gaze position but with no learned-gaze module in the optimisation path. **[TODO: Training in progress at submission; results will populate this paragraph.]**

Shortcut: a 2×2 dissociation between probe-readable and policy-used memory Plotting probe-readable GPS R^2 against behavioural memory reliance (shortcut SPL drop) reveals a 2×2 pattern (Figure 5). Blind sits in the “has GPS, uses memory” diagonal; foveated and foveated-learned in the “no GPS, ignores memory” diagonal — both consistent with H1. The off-diagonal cells are the interesting ones: *matched-compute has GPS but ignores memory* (linear top-layer cognitive map, low policy weight on it), and *uniform has no linear GPS but heavy memory reliance* (integrates non-spatial scene-history / visual-landmark features). “Probe-readable” and “policy-relied-on” are therefore distinct constructs that can dissociate in either direction.




Figure 5: Probe-readable GPS (H1) vs. behavioural memory reliance (shortcut). The shortcut experiment runs paired episodes in the same scene with fresh goals and compares second-episode SPL under *reset* (LSTM zeroed) vs. *persistent* memory; larger drop = more memory-reliant policy. *x*-axis: GPS R^2 (5-fold CV mean) from Figure 2a. *y*-axis: relative shortcut SPL drop (%). 20 Gibson scenes $\times$ 10 paired episodes per condition. Fov-learned’s $R^2 = -2.43$ is clipped at -1.5 (arrow). The 5 conditions populate a 2×2 dissociation: blind “has GPS code, uses memory”; foveated-fixed and foveated-learned “no linear GPS, ignores memory”; **matched** and **uniform** are the two opposite anomalies (matched: “has GPS but ignores memory”; uniform: “no linear GPS yet still memory-reliant”).




Figure 6: Persistent-memory failure made visual: trajectories from one paired-episode run per condition where reset memory succeeds and persistent memory fails. Same scene, same start position, same goal between the two memory conditions; the only difference is whether the LSTM is zeroed (reset) or carries over from the previous episode (persistent). Blind (left) shows the cleanest interpretation of the shortcut SPL drop: persistent agent oscillates around the *old* goal location (light star, $N-1$) instead of going to the new one (gold star, N), consistent with the “LSTM still encodes previous task” reading. Uniform (centre) and foveated (right) persistent runs wander more diffusely without obviously locking onto the old goal — consistent with their hidden state encoding scene-history / visual features rather than a position-readable memory of the previous task. The contrast in trajectory length is also informative (blind reset 65 steps vs. persistent 443; uniform 38 vs. 2000 max; foveated 55 vs. 2000).

Figure 6 shows the failure modes qualitatively differ across regimes: blind’s persistent trajectory clusters around the previous-episode goal (literal “locked onto old position”); uniform and foveated wander diffusely (“locked onto old visual context” would fit). The visualisation cannot itself distinguish position-locked from context-locked failure but is consistent with the 2×2 dissociation. **[TODO: Both anomalies rest on single-seed runs; multi-seed replication is in progress and would promote the dissociation from candidate to robust. A direct “policy GPS reliance” test (ablating the GPS sensor mid-rollout) would distinguish “policy reads the GPS code” from “policy reads non-GPS memory” — the matched anomaly predicts a small SPL drop, the uniform anomaly a large one.]**

4.6 Boundaries and additional probes

Occupancy controls Two probes fail to discriminate between conditions and so define the boundary of our claims. Coarse occupancy (“have I visited this region?”) decodes above chance in every condition, so one-bit spatial working memory is ubiquitous and not part of the H1–H3 story. Metric 20×20 occupancy reconstruction fails everywhere, consistent with the SPACE benchmark’s finding that fine-grained layout requires non-linear probes or longer horizons [Ramakrishnan et al., 2025]: our claims are about representational *format* and *content*, not about high-resolution metric layout.

Goal vector No condition stores the ego-frame goal vector linearly: goal *distance* (DtG) decodes strongly across conditions ($R^2 \in [0.81, 0.90]$, validating the probe machinery), while goal *direction* and the full 2-D ego-to-goal vector are at chance or negative for every condition (Appendix A, Figure 10). This is unsurprising — PointGoal provides the goal vector as a per-step input, so the policy has no incentive to re-encode it — and is not part of the H1/H2/H3 story.

Per-unit spatial information and population structure A population coding analysis (Appendix A, Figure 9) sharpens the H1 picture in a direction we did not expect. Per-unit spatial information [Banino et al., 2018]: rich-encoder conditions have a few units with *higher* peak spatial information (≈ 2 bits) than blind (≈ 1 bit max). What separates conditions is what those units actually encode — a sparse-decoding curve (probe GPS using only the top- k spatial-info units) reveals that blind reaches $R^2 \approx 1.0$ at $k = 512$, matched climbs to $R^2 \approx 0.5$, and the three rich-encoder networks fail to decode GPS at any k . The peaked units in rich-encoder conditions therefore appear to encode features correlated with position (visual landmarks, trajectory-phase markers) rather than position itself, while blind’s flatter, lower-peak distribution is what actually carries the linear position code. Intrinsic dimension via PCA-95% confirms this: blind’s representation is the most compressed (180 dimensions), matched intermediate (271), and rich-encoder networks the highest (291–303). *Caveat*: the spatial-info computation depends on the per-scene grid resolution and we do not claim a sharp dissociation between “population code” and “few dedicated cells”; the qualitative direction (peak per-unit info $\neq$ population GPS readout) is what we report.

5 Discussion

5.1 Two orthogonal axes: sensor structure and gaze location

Our findings unify under a candidate account we call the *encoder–memory race*: top-layer linear GPS encoding in the LSTM is *maintained* when the encoder hands the LSTM minimal spatial-feature variety per step (blind, matched-compute), and *overwritten* when richer visual features substitute (uniform, foveated, foveated-learned). The encoder feature-map probe (§4.4) confirms that no sighted encoder linearly re-derives GPS, so the H1 race is not encoder-vs-LSTM at the position-information level; it is at whether the LSTM has visual variety enough to substitute for integrating its Layer-0 GPS-sensor signal across time. H2 (§4.3) cuts orthogonally: *what* the LSTM encodes lives in distinct linear subspaces across conditions, and the difference matters for behaviour (cross-condition memory transplants drop SPL beyond transplant-mechanics noise; the 5×5 matrix is asymmetric). H3 (§4.5) asks whether, within the rich-encoder pass-through regime, gaze location is a second content axis; the foveated-shifted causal control is in training. We present the unifying account as a candidate, not a proven mechanism: a direct causal test (e.g. ablating the visual stream mid-rollout) is left for follow-up.

5.2 Why this, not alternative accounts

Task competence. All conditions solve PointGoal at 93–99% success; the encoded-content gap is not a skill difference.

Probe capacity / fitting. DtG decodes robustly everywhere ($R^2 \geq 0.81$), demonstrating Ridge can fit high- R^2 scalars when the hidden state contains them. Hewitt–Liang label-permutation selectivity is high where GPS is high (blind, matched) and at chance where GPS is at chance — the latter follows from real and control probes facing the same near-zero target-information ceiling. Rich-encoder chance-level R^2 is therefore a hidden-state property, not a probe artefact.

Sample size / trajectory distribution. Identical 500-episode deterministic evaluation across conditions; the earlier 4-step stochastic-sampling artefact is documented and retired in §3.3.

Linear-probe artefact. Two complementary checks. (a) *Behavioural*: memory transplant goes outside the probe framework and reproduces the bottleneck-vs-pass-through ordering; the 5×5 matrix’s asymmetry adds a structure that no linear-similarity measure can read. (b) *Non-linear probe*: a 2-layer MLP recovers $R^2 \in [0.51, 0.73]$ from rich-encoder top-layer states where Ridge sees chance. The bottleneck-vs-pass-through divergence is therefore not about whether *any* probe can recover position; it is specifically about whether a *linear* top-layer probe can. Since the DD-PPO action head is a linear projection from $\mathbf{h}_2$, the linear metric is what the policy itself can use, so H1 maps onto policy-accessible encoding regardless of how the MLP interpretation resolves. We cannot from MLP recovery alone exclude exploitation of position-correlated features (wall distance, scene identifiers, episode-phase) rather than position itself; Appendix E tests this more directly.

5.3 Biological precedent as convergent evidence

Each of our three findings has a parallel in the animal-navigation literature, which we take as one piece of convergent (not conclusive) evidence that the effects are not specific to a single LSTM architecture. Compensatory memory under reduced sensory access is documented in congenitally blind humans, who develop enhanced hippocampal coupling and outperform sighted controls on non-visual spatial-working-memory tasks [Kupers and Ptito, 2014], and in rodents whose grid-cell periodicity collapses in darkness despite intact head-direction inputs [Chen et al., 2016]—both format-level, not just strength-level, effects. Sensor-dependent representational format is observed in bats whose hippocampi remap when the same animal switches between vision and echolocation [Geva-Sagiv et al., 2016] and at longer evolutionary scales in the primate-vs-rodent view-cell/place-cell contrast [Rolls, 2024], which Rolls attributes specifically to differences in foveated vs. near-panoramic sampling statistics. Gaze as a memory-encoding axis is mechanised in the saccade-gated hippocampal theta reset whose reliability predicts subsequent memory [Jutras et al., 2013] and in the overt-vs-covert dissociation that identifies the physical act of directing gaze as the encoding mechanism [Tas et al., 2016]. We do not claim our agents implement these biological mechanisms, nor that the parallel structure settles the generality question; we do claim that the parallels make a pure single-architecture artefact less likely, and that the controlled agent setting complements biological comparisons where sensor type, architecture, and training objective cannot be independently manipulated.

5.4 Implications for ML practice and architectural scope

Within the recurrent PointNav setting we study, three implications follow. (i) *Encoder capacity correlates inversely with linear, top-layer, policy-accessible LSTM spatial encoding*; if causal (Appendix E tests this), it is also a practical training-time lever for inducing cognitive-map-like internal representations. (ii) *“More pixels” is not the right lens*: the linear top-layer GPS code is non-monotonic in raw pixel count (uniform 256^2 lacks it; matched 48^2 with 1×1 feature map has a strong one). The encoder feature-map probe shows neither encoder linearly preserves GPS, so the mechanism is not encoder position retention; we attribute the pattern to encoder spatial-feature variety per step, with the scaling sweep and log-polar control needed to confirm. (iii) *Gaze (H3 pending)*: within the rich-encoder regime, we predict shifting the static gaze location alters what the LSTM encodes; the foveated-shifted causal control is in training (§4.5).

Beyond the recurrent setting (untested) The three implications are established only in our LSTM setting. Transformer-based embodied agents build their internal state by accumulating information via learned attention, structurally similar to our gaze location in the sense that it selects which observation

content is read into the representation; H3 therefore suggests a testable prediction for transformer navigators (strong constraints on attended tokens $\rightarrow$ content-level dissociations like ours) which we do not verify here. The sensor-structure axis (H1, H2) is a claim specifically about RL recurrent navigation agents under partial observability; whether equivalent effects arise in supervised visual learners, non-visual modalities, or multimodal systems is an open question.

5.5 Limitations

(i) *Intermediate checkpoint for foveated-fixed*: all foveated-fixed analyses use ckpt.36 at $\sim 174\text{M}$, the last checkpoint before the silent-corruption window (Appendix C); the other four conditions use the final converged checkpoint. (ii) *Minimal learned-gaze architecture*. Our learned-gaze module is a deterministic sigmoid MLP with slow-gaze rollout updates, chosen for computational compatibility with DD-PPO rollouts and simplicity as a controlled variable, not as a competitive active-gaze model. Its collapse under pure task supervision is specific to this architecture (and so far to a single seed; multi-seed replication is in progress); richer gaze mechanisms (stochastic gaze with entropy bonus, information-seeking intrinsic reward, attention-based token selection) would need separate investigation. We therefore do not claim that implicit task pressure cannot produce active gaze in general; we claim only that our minimal decoder collapsed, and we use the collapsed-gaze position as an empirical starting point for the gaze-location test (foveated-shifted control, §4.5) — a test whose logic does not depend on the gaze decoder working. (iii) *Cross-heading probe*: the proposal included a cross-heading generalisation probe (train on heading A , test on opposite heading). Our implementation returned “insufficient samples” sentinels at our probe sizes and does not appear in our results. (iv) *Foveation model*: Gaussian blur with quadratic falloff is a simplification; log-polar or multi-scale pyramid models may behave differently. A log-polar control (F3) and a $\sigma_{\max} \in \{2, 4, 12, 20\}$ strength sweep (F1–F4) are in training at submission and will populate §4.4 and Appendix D. (v) *Linear probes only*: non-linear probes might reveal structure Ridge regression misses, particularly for occupancy maps. (vi) *Architecture scope*: we use a recurrent (LSTM) backbone to remain comparable with the emergent-map literature we build on. The core claims (sensor structure and gaze location as representational-content axes) should hold for transformer-based navigation agents whose internal state is built by learned attention, but directly verifying this requires re-running the five-condition ablation on a transformer backbone and is left for future work; see §5 for the specific predictions.

6 Conclusion

We ran a controlled five-condition ablation isolating the effect of visual-input structure on the recurrent representations learned by an otherwise-identical DD-PPO navigation agent. Three findings emerge.

(1) *Encoder-memory race (H1)*. Top-layer LSTM GPS encoding is *temporally stable* when the encoder hands the LSTM minimal spatial-feature variety per step (blind, matched-compute), and *transient* when richer visual features substitute (uniform, foveated, foveated-learned). The encoder feature-map probe shows no sighted encoder linearly decodes GPS, so the H1 mechanism lives inside the LSTM, not at the encoder. H1 is most precisely a claim about *linear, top-layer, late-episode* GPS encoding — what the linear DD-PPO action head can read.

(2) *Condition-specific representational format (H2)*. Even where rich-encoder LSTMs are spatially pass-through, their hidden-state subspaces are linearly separable to single-nearest-neighbour granularity ($\text{CKA} < 10^{-4}$, probe transfer $\ll -800$, 1-NN purity 1.000). The 5×5 memory-transplant matrix is asymmetric: bottleneck donor states hurt rich-encoder recipients far more than the reverse. Whatever non-spatial task-relevant features each LSTM carries, it carries them in its own linear subspace.

(3) *Gaze location (H3, pending) and a having-vs.-using dissociation*. Whether gaze location is a second content axis *within* the rich-encoder regime is in training (foveated-shifted control). Independently of that pending test, plotting probe-readable GPS against behavioural memory reliance reveals a 2×2 dissociation: matched has linear GPS but ignores it; uniform has no linear GPS but heavily uses memory. Both anomalies rest on single-seed runs and we mark them as candidate dissociations pending multi-seed replication.

The pattern parallels three threads of animal-navigation research — enhanced hippocampal coupling in congenitally blind humans, the disruption of rodent grid-cell firing in darkness with navigation nonetheless preserved, and cross-modality hippocampal remapping in bats. We use these as analogy only: the convergent direction across systems is suggestive of a sensor–memory coupling principle, but our experiments do not test architectures other than LSTM, so we do not claim the principle is architecture-independent. Open questions: (i) whether the encoder-bottleneck ordering extends to finer-grained resolution sweeps (Appendix E); (ii) whether H3 stands on the clean causal test; (iii) whether transformer-based navigation agents reproduce the encoder–memory race; (iv) whether equivalent effects shape representations in non-navigation embodied tasks; (v) multi-seed replication of all five conditions to bound seed-to-seed variability of the H1 numbers.

References

- Andrei Atanov, Yicheng Fu, Gaurav Singh, Jianren Yu, Andrew Spielberg, and Amir Zamir. How far can a 1-pixel camera go? solving vision tasks using photoreceptors and computationally designed visual morphology. In *European Conference on Computer Vision (ECCV)*, 2024.
- Andrea Banino, Caswell Barry, Benigno Uria, Charles Blundell, Timothy Lillicrap, Piotr Mirowski, Alexander Pritzel, Martin J Chadwick, Thomas Degris, Joseph Modayil, et al. Vector-based navigation using grid-like representations in artificial agents. *Nature*, 557(7705):429–433, 2018.
- Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics*, 48(1):207–219, 2022.
- Guifen Chen, John A. King, Neil Burgess, and John O’Keefe. Absence of visual input results in the disruption of grid cell firing in the mouse. *Current Biology*, 26(17):2335–2342, 2016.
- Christopher J Cueva and Xue-Xin Wei. Emergence of grid-like representations by training recurrent neural networks to perform spatial localization. In *International Conference on Learning Representations (ICLR)*, 2018.
- Kshitij Dwivedi, Gemma Roig, Aniruddha Kembhavi, and Roozbeh Mottaghi. What do navigation agents learn about their environment? In *Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Maya Geva-Sagiv, Sandro Romani, Liora Las, and Nachum Ulanovsky. Hippocampal global remapping for different sensory modalities in flying bats. *Nature Neuroscience*, 19(7):952–958, 2016.
- James Gornet and Matt Thomson. Automated construction of cognitive maps with visual predictive coding. *Nature Machine Intelligence*, 6:820–833, 2024.
- Jacqueline Gottlieb, Pierre-Yves Oudeyer, Manuel Lopes, and Adrien Baranes. Information-seeking, curiosity, and attention: computational and neural mechanisms. *Trends in Cognitive Sciences*, 17(11):585–593, 2013.
- Mary M. Hayhoe and Jonathan S. Matthis. Control of gaze in natural environments: effects of rewards and costs, uncertainty and memory in target selection. *Philosophical Transactions of the Royal Society B*, 373(1718):20170368, 2018.
- Jay A Hennig, Sandra A Romero Pinto, Takahiro Yamaguchi, Scott W Linderman, Naoshige Uchida, and Samuel J Gershman. Emergence of belief-like representations through reinforcement learning. *PLOS Computational Biology*, 19(9):e1011067, 2023.
- John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2019.
- Michael J. Jutras, Pascal Fries, and Elizabeth A. Buffalo. Oscillatory activity in the monkey hippocampus during visual exploration and memory formation. *Proceedings of the National Academy of Sciences*, 110(32):13144–13149, 2013.
- Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *International Conference on Machine Learning (ICML)*, 2019.

- Ron Kupers and Maurice Ptito. Compensatory plasticity and cross-modal reorganization following early visual deprivation. *Neuroscience & Biobehavioral Reviews*, 41:36–52, 2014.
- Volodymyr Mnih, Nicolas Heess, Alex Graves, and Koray Kavukcuoglu. Recurrent models of visual attention. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2014.
- Jiri Najemnik and Wilson S. Geisler. Optimal eye movement strategies in visual search. *Nature*, 434(7031):387–391, 2005.
- Motahareh Pourrahimi and Pouya Bashivan. Emergent brain-like representations in a goal-directed neural network model of visual search. *bioRxiv preprint*, 2025.
- Santhosh Kumar Ramakrishnan, Erik Wijmans, Philipp Krähenbühl, and Vladlen Koltun. Does spatial cognition emerge in frontier models? In *International Conference on Learning Representations (ICLR)*, 2025.
- Edmund T. Rolls. Hippocampal spatial view cells for memory and navigation, and their underlying connectivity in humans. *Hippocampus*, 33(6):533–572, 2023.
- Edmund T. Rolls. Why do primates have view cells instead of place cells? *Trends in Cognitive Sciences*, 28(9):862–874, 2024.
- Ruth Rosenholtz. Capabilities and limitations of peripheral vision. *Annual Review of Vision Science*, 2:437–457, 2016.
- Manolis Savva, Abhishek Kadian, Oleksandr Maksymets, Yili Zhao, Erik Wijmans, Bhavana Jain, Julian Straub, Jia Liu, Vladlen Koltun, Jitendra Malik, Dhruv Batra, and Devi Parikh. Habitat: A platform for embodied AI research. In *International Conference on Computer Vision (ICCV)*, 2019.
- Jinghuan Shang and Michael S. Ryoo. Active vision reinforcement learning under limited visual observability. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Nathan Sprague and Dana H. Ballard. Eye movements for reward maximization. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2003.
- Hans Strasburger, Ingo Rentschler, and Martin Jüttner. Peripheral vision and pattern recognition: A review. *Journal of Vision*, 11(5):13, 2011.
- Amanda C. Tas, Steven J. Luck, and Andrew Hollingworth. The role of attention in the encoding of visual working memory. *Journal of Experimental Psychology: Human Perception and Performance*, 42(8):1121–1138, 2016.
- Erik Wijmans, Abhishek Kadian, Ari Morcos, Stefan Lee, Irfan Essa, Devi Parikh, Dhruv Batra, and Manolis Savva. DD-PPO: Learning near-perfect pointgoal navigators from 2.5 billion frames. In *International Conference on Learning Representations (ICLR)*, 2020.
- Erik Wijmans, Manolis Savva, Irfan Essa, Stefan Lee, Ari Morcos, and Dhruv Batra. Emergence of maps in the memories of blind navigation agents. In *International Conference on Learning Representations (ICLR)*, 2023.
- Sylvia Wirth, Pierre Baraduc, Angélique Planté, Serge Pinede, and Jean-René Duhamel. Gaze-informed, task-situated representation of space in primate hippocampus during virtual navigation. *PLOS Biology*, 15(2):e2001045, 2017.
- Fei Xia, Amir R Zamir, Zhiyang He, Alexander Sax, Jitendra Malik, and Silvio Savarese. Gibson Env: Real-world perception for embodied agents. In *Computer Vision and Pattern Recognition (CVPR)*, 2018.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer:
answerYes

Justification: The abstract and §1 state the three hypotheses H1/H2/H3 with explicit signatures (what pattern supports each), and preview the three findings with specific magnitudes (path-history $R^2=0.56$, CKA <0.004 , compass $R^2=0.94$ vs GPS $R^2=0.03$). These are expanded in §4 and §5 with no stronger claims than the data support.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer:

answerYes

Justification: §5.4 explicitly lists six limitations (i–vi) covering checkpoint staleness, single-seed H3, missing proposed analyses, foveation-model simplification, linear-probes-only, and LSTM-architecture scope.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer:

answerNA

Justification: The paper reports empirical results; it contains no theoretical claims or proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer:
answerYes

Justification: §3.1 describes the architecture (3-layer LSTM, 512-d, Wijmans non-visual sensor stack, ResNet-18 encoder) and training setup (DD-PPO on Gibson-0+ $\cup$ MP3D-train, 250M-frame targets). §3.2–§3.6 describe all probing and cross-condition protocols. Appendix A details the gaze-diversity ablation. Code will be released upon publication.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer:
answerYes

Justification: §3.7 commits to open-source release of the training and probing code. Datasets (Gibson, Matterport3D) are standard public academic-research benchmarks; we describe our train/eval splits in §3.1.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).

- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer:

answerYes

Justification: §3.1 specifies training pool and eval set; §3.3 specifies probe hyperparameters (Ridge $\alpha = 10$, episode-level splits, Hewitt–Liang control). Appendix A gives the ablation’s loss formulation and coefficient.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer:

answerNo

Justification: We report single-seed numbers per condition. This is documented as Limitation (ii) in §5.4. Multi-seed replication for tighter confidence intervals is in the immediate follow-up.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).

- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer:

answerYes

Justification: §3.7 reports per-condition training budget (5–7 days on 2×V100 for 250M frames) and probing budget (<1h per condition on one V100).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer:

answerYes

Justification: The research uses public academic-research benchmarks, publicly-released open-source RL frameworks, and involves no human subjects or deployment. We also report and remediate a numerical-stability bug upstream (Appendix B).

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer:

answerYes

Justification: §5.3 discusses implications — sensor design as a tunable axis for interpretable embodied agents. Negative societal impacts are minimal: the work is fundamental research on simulated indoor navigation with no direct deployment pathway.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer:

answerNA

Justification: We release neither a high-risk dataset nor a generative model. Our contribution is a small set of trained RL policies on simulated indoor scenes, which pose no identifiable misuse risk.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer:

answerYes

Justification: Gibson Database and Matterport3D are used under their released academic-research licenses; Habitat-Sim and habitat-lab are MIT-licensed. Our use complies with their terms; citations §3.1.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.

- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer:

answerYes

Justification: Code will be released with documentation upon publication. Figures are novel; they are released under the paper's license alongside the code. The gaze-diversity auxiliary loss and NaN-sanitisation patch are small standalone modules with docstrings.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer:

answerNA

Justification: The paper involves no crowdsourcing and no human-subjects data.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer:

answerNA

Justification: No human-subjects research; no IRB/equivalent review was required.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer:

answerNA

Justification: LLMs were used only for writing assistance on the manuscript. They are not part of the core research methodology; all empirical results come from standard RL training and linear probing.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.

A Supplementary visualisations




Figure 7: Supplementary visualisations referenced from the main text.




Figure 8: Memory-transplant midpoint sweep on three representative donor→recipient pairs (foveated→uniform, foveated→blind, foveated-learned→foveated) across midpoints $\{0, 15, 30, 60, 90\}$ time-steps ($n=100\text{--}150$ episodes/cell). The full 5×5 matrix in Figure 3 (right) is at midpoint $t=30$; this supplementary figure shows that the cross-self gap is roughly stable for $t \geq 15$ and that the $t=0$ self-/cross- baseline collapses to a protocol-noise floor (so the diagonal in Figure 3 is correctly read as “rollout-divergence noise” rather than “condition-mismatch noise”). Foveated-learned→foveated incurs -0.17 SPL on average across $t \geq 15$, foveated→uniform -0.10 , foveated→blind -0.02 , consistent with the asymmetry pattern in the full matrix.

B Gaze-Diversity Auxiliary Loss (H3 Ablation)

Motivation In the first foveated-learned training run the gaze decoder collapsed (per-dimension $\text{std} \approx 5 \times 10^{-4}$). We attribute this to two mechanisms: (i) once the sigmoid output saturates, its derivative $\sigma'(x) \rightarrow 0$ attenuates gradient signal; (ii) consistent-gaze views yield more predictable visual features, so the reward-gradient path has no explicit incentive for variety.

Target-variance hinge We add $\mathcal{L}_{\text{div}} = \lambda \cdot \text{relu}(\sigma_{\text{target}}^2 - \text{Var}_{\text{batch}}[g])$ where $g \in [0, 1]^2$, $\text{Var}_{\text{batch}}$ averaged over the two gaze dimensions, $\sigma_{\text{target}}^2 = 0.01$ ($\text{std} \approx 0.1$), $\lambda = 0.01$. The hinge disappears




Figure 9: Population coding analysis (referenced from §4.6). (a) Per-unit spatial-information distribution: each curve is one condition’s units sorted by spatial information (descending), x-axis log scale. Counter-intuitively, rich-encoder conditions (uniform / foveated / fov-learned) have a few more peaked spatially-tuned units (max ≈ 2 bits) than bottleneck conditions; blind has the flattest distribution (max ≈ 1 bit, no unit reaches the dashed 1-bit threshold). (b) Sparse-vs-distributed GPS decoding: probe R^2 when only the top- k spatial-info units are used as features. Blind reaches $R^2 \approx 1.0$ at $k = 512$; matched climbs to $R^2 \approx 0.5$; rich-encoder networks fail to decode GPS at any k even though their top units carry higher per-unit spatial information. The combined picture: *higher per-unit spatial information $\neq$ position readout*. Rich-encoder peaked units appear to encode features correlated with position (e.g. visual landmarks, trajectory phase) rather than position itself.




Figure 10: Goal-vector probe R^2 (referenced from §4.6). Goal *distance* (DtG) decodes strongly across all conditions ($R^2 \in [0.81, 0.90]$, validating probe machinery); goal *direction* and the full ego-frame goal vector are at chance or negative everywhere. PointGoal provides the goal vector as a per-step input, so the policy has no incentive to re-encode it.

once variance exceeds the target. We also evaluated an un-capped $-\lambda \cdot \text{Var}[g]$ variant as a negative control.

Pilot 1 (uncapped, job 2840256) After 10M frames, gaze variance exceeded that of a uniform distribution over $[0, 1]^2$ (std $[0.43, 0.42]$ vs. theoretical max 0.289): gaze was essentially random. Task performance collapsed: SPL plateaued at 0.05, running-mean metrics transitioned to NaN at ~ 8 M frames. The uncapped regulariser destroyed the consistent-foveation signal the policy needs.

Pilot 2 (hinge, job 2842288) Gaze converged to an extreme corner of image space: mean $(0.001, 0.998)$ with std $[0.020, 0.023]$ — well below the target std of 0.1 and therefore with the hinge active throughout. Task SPL plateaued at 0.05 as in pilot 1. The policy apparently found the minimum-task-cost gaze that partially satisfies the diversity pressure: a sigmoid saturated at one corner, producing a constant image-corner view.

Interpretation Both pilots argue that in our deterministic sigmoid-MLP gaze architecture with slow-gaze rollouts, gaze collapse is a pre-requisite for task learning rather than an accident. Task signal alone does not spontaneously produce informative fixation, and an exogenous anti-collapse regulariser moves gaze toward whatever region reduces regulariser loss fastest, which is generically *not* the task-useful region. A genuine H3 test likely requires (i) a stochastic gaze location with




Figure 11: Per-layer and MP3D companion figures (referenced from §4.2).

entropy bonus, so diversity is sampled at rollout time rather than enforced on a deterministic output; or (ii) a task-aligned intrinsic reward (e.g., information gain about goal-relative position); or (iii) an architectural change separating gaze generation from visual-feature extraction (e.g., attention-based gaze).

C Training Stability: Silent Weight Corruption in DD-PPO

During an initial foveated training run we observed a failure mode that did not manifest as a crash but as slow, silent corruption of the learned weights. After $\approx 170\text{M}$ frames of stable progress (reward ≈ 10 , SPL ≈ 0.83), one mini-batch on one DDP worker produced a NaN gradient. Because `torch.nn.utils.clip_grad_norm_` does not sanitise non-finite gradients—the global norm becomes NaN, the clip coefficient becomes NaN, and `grad *= NaN` remains NaN—the subsequent `optimizer.step()` wrote NaN into every parameter. DDP all-reduce propagated the corruption across workers. The training loop did not raise: it continued for 2.5 days with reward=NaN and SPL= 0 until `torch.multinomial` finally rejected the NaN probability tensor. All checkpoints saved during this window were unrecoverable (90 of 97 parameter tensors entirely NaN).

To isolate the origin of the NaN gradient we re-ran training from the last clean checkpoint with `torch.autograd.set_detect_anomaly(True)` and per-module forward-output hooks. The anomaly traceback pinned the failure to `MseLossBackward0`, i.e. the backward pass of the value-head loss $\frac{1}{2}(V(s) - R)^2$; rare return-vs-value discrepancies overflowed float32 in the squaring step, producing non-finite gradients downstream.

Our mitigation is a minimally invasive patch to `PPO.before_step` that zeroes any non-finite gradient element before clipping. A bad mini-batch becomes a no-op update rather than a catastrophic one. The patch is a no-op on clean training paths. Across the five conditions reported here, no NaN events were observed during full training after deployment, and all final checkpoints contain only finite weights. We reported the underlying issue upstream (habitat-lab issue #2226) with a minimal reproduction and a proposed native fix.

D Foveation experiments: status and design

This appendix documents the four controlled foveation experiments described in §4.4. All four were submitted to the cluster at the time of writing; the table below tracks their state. Numbers will populate the corresponding paragraphs in §4.4 in revision.

Table 4: Foveation experiment status (submission time).

Experiment	Tests	Status
F1–F4 strength sweep ($\sigma_{\max} \in \{2, 4, 12, 20\}$)	encoder–memory race as continuous lever	in training
F3 log-polar foveation	spatial-sampling vs. blur bottleneck	in training
Encoder feature-map probe (matched / uniform / foveated)	where position information is held	in probing
Foveated-shifted (gaze (0.49, 0.62))	gaze-location axis (links to H3)	in training

Why disclose this rather than wait. The four experiments are designed to falsify or refine specific predictions of the encoder–memory race framework, and each has a falsifiable signature. We disclose the design openly so that (i) the present submission is honest about what foveation-specific evidence already exists vs. what is in flight, and (ii) any reader can identify which result would change the paper’s interpretation. The current evidence already places foveation-fixed and foveation-learned within the framework; the experiments above sharpen the “foveation” part of the story specifically.

E Encoder-capacity scaling sweep

[TODO: The scaling sweep trains matched-compute variants at input resolutions {32, 48, 64, 96, 128, 192} alongside the existing blind (no encoder) and uniform 256×256 anchors. Probe results on the trained checkpoints will populate Figure 12 and this appendix section. Hypothesis: GPS probe R^2 decreases monotonically with input resolution, from $R^2 \approx 0.95$ at no-encoder to $R^2 \approx 0$ at 256×256, crossing into the pass-through regime between 48 and 128 pixels.]



Figure 12: Encoder-capacity scaling curve. **[TODO: Pending data.]** x-axis: input resolution in pixels; y-axis: GPS and compass probe R^2 (5-fold CV, mean $\pm\sigma$).